

Programmatic on Linux Kernel with eBPF

- DevClub Tech Meetup#6 eBPF -



JoJo
(Cloud Native Stylist)



JoJo
(Cloud Native Stylist)



THAI PROGRAMMER
สมาคมโปรแกรมเมอร์ไทย

กรรมการ และประธานคณะทำงาน

Jumpbox®

"เราเชื่อว่า การเรียนรู้ทำให้ชีวิตคุณดีขึ้น"

-Jumpbox Team-

Jumpbox®



Jumpbox®

Tech Passion | Sharing | Society

Why eBPF?



eBPF Foundation

Initial release: 2014; 10 years ago

Open source community, Meta, Google, Isovalent, Microsoft, Netflix



NETFLIX

Reference:

- [eBPF Foundation](#)
- [eBPF](#)

Everything You Need to Know in 5 Minutes

- DevClub Tech Meetup#6 eBPF -

eBPF has become the key technology for

eBPF has become the key technology for
infrastructure software

Accurate Definition

" **eBPF** is a **programming language** & **runtime** to extend operating system"

-Thomas Graf, CTO, Isovalent -



Reference:

- [Keynote: eBPF](#)
- [KubeCon2022 eBPF Day North America](#)

Practical Comparison

" **eBPF** is like **JavaScript/Lua** but for Kernel Developers"

-Thomas Graf, CTO, Isovalent -

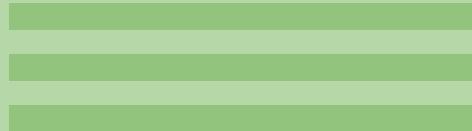


Reference:

- [Keynote: eBPF](#)
- [KubeCon2022 eBPF Day North America](#)



Website



Submit

```
function hello() {
| alert('hello')
}
</script>
<form onsubmit="return false;">
| <input type="submit"
|   name="hello"
|   onclick="hello()">
</form>
```



</> Process

execve()

Operating System



System Calls

```
int syscall__ret_execve(struct pt_regs *ctx)
{
    struct comm_event event = {
        .pid = bpf_get_current_pid_tgid() >> 32,
        .type = TYPE_RETURN,
    };

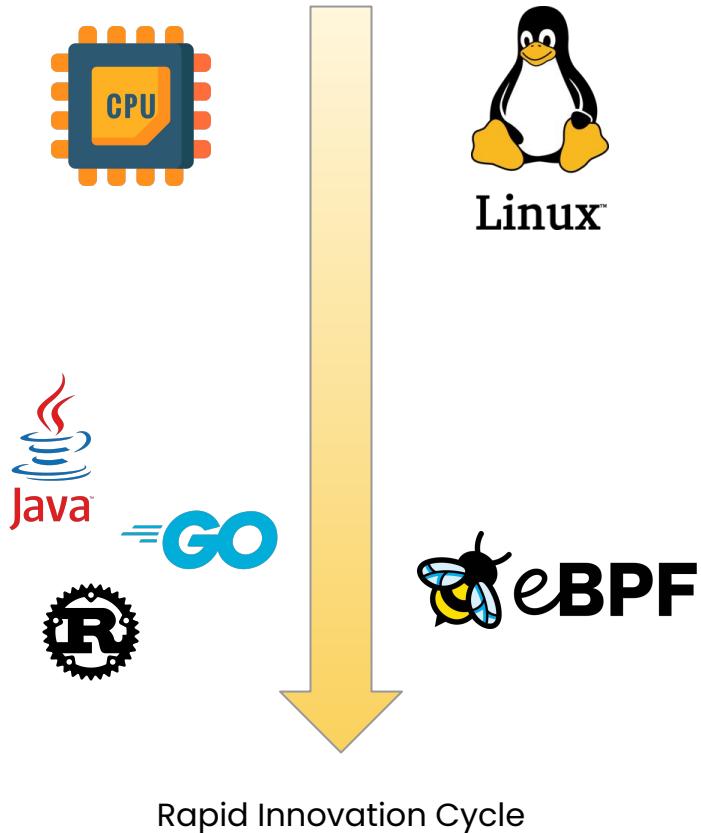
    bpf_get_current_comm(&event.comm, sizeof(event.comm));
    comm_events.perf_submit(ctx, &event, sizeof(event));

    return 0;
}
```

Jumpbox®

**Operating Systems are like hardware,
hard to change and with long innovation cycles**

Long Innovation Cycle



Long innovation cycles result in the need to predict use cases or stick to providing building blocks

Programmability allows you to continuously adapt to changing requirements and innovate quickly

How is  eBPF different
from  or  ?



- Designed specifically to be embedded into operating systems (Linux, Windows)
- Can interface with internal operating system APIs not available to user space
- Restricted to run safely in kernel context
- Aimed at kernel developers, hard to learn



- Designed to be embedded into applications
- General Purpose
- Aimed at application developers, easy to learn
- ... **but** cannot run in an operating system context

How does eBPF work?



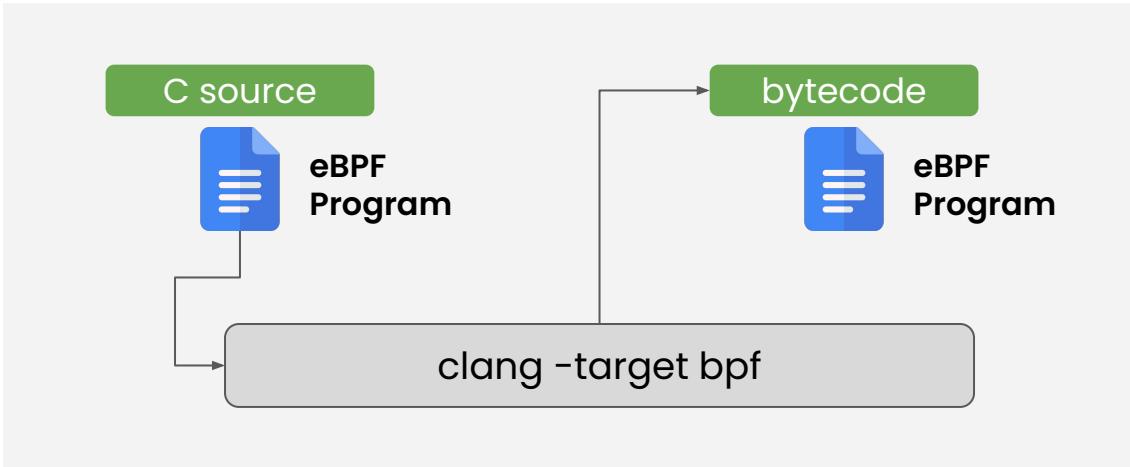
eBPF Language



eBPF Runtime

Jumpbox®

eBPF Language



cilium/ebpf



libbpf-rs | redbpf | Aya
Rust



libbpf
C++



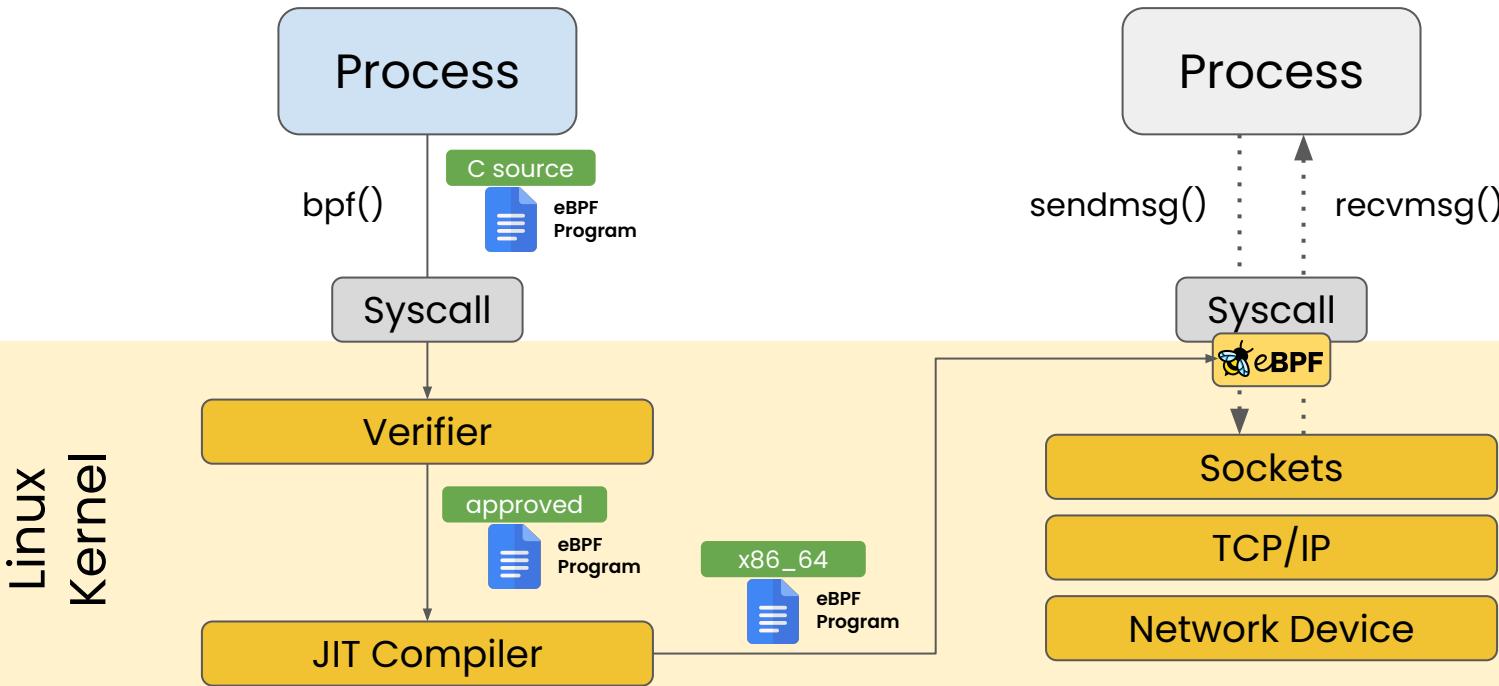
bcc

Multiple SDKs and Compilers exist to get to eBPF bytecode

Jumpbox®



eBPF Runtime



The runtime accepts bytecode, verifies it, just-in-time compiles it, and runs it at the requested hook point.



secure

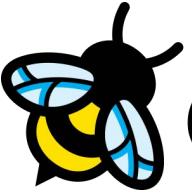
Runtime Verification
Sandbox Concept
Program Signing

efficient

JIT Compiler
Embedded in OS
Per-CPU data Structures

portable

Generic Bytecode
Data Type Discovery
Stable API to OS

Who controls  eBPF ?



eBPF Foundation

Platinum



TIGERA



HUAWEI



CROWDSTRIKE



Silver

NETFLIX



DaoCloud

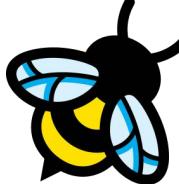


数悦铭金
DATOMS



FUTUREWEI

Jumpbox®

Where is  eBPF used
today?

Cloud Native eBPF Landscape

Application Observability



Networking & Service Mesh



Calico



Falco

CLOUD NATIVE COMPUTING FOUNDATION

Jumpbox®

Hyperscalers



<https://github.com/facebookincubator/katran>

Blazing fast,  eBPF-based
L4 load-balancer used at
Facebook/Meta

Jumpbox®

All major providers have picked
 eBPF-based **Networking** & **Security**
for their Kubernetes platforms



Google Cloud

Jumpbox®

Smartphones



Android BPF Loader

During Android boot, all eBPF programs located at `/system/etc/bpf/` are loaded. These programs are binary objects built by the Android build system from C programs and are accompanied by `Android.bp` files in the Android source tree. The build system stores the generated objects at `/system/etc/bpf/`, and those objects become part of the system image.

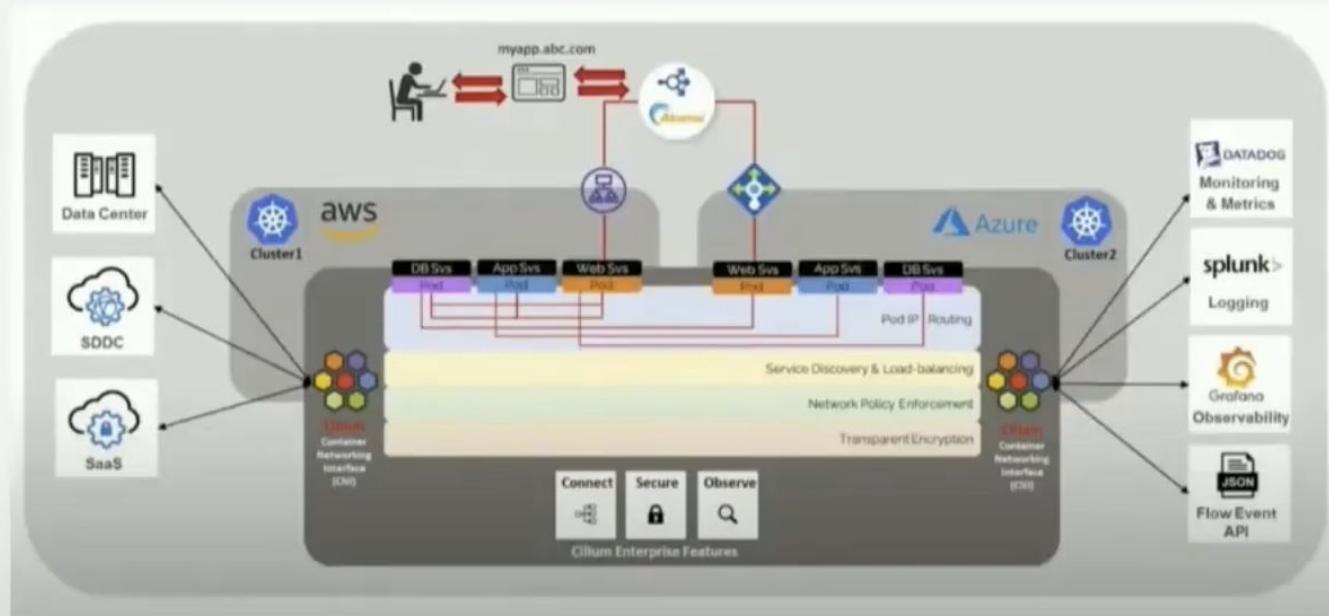
Examples of eBPF in Android

The following programs in AOSP provide additional examples of using eBPF:

- The `netd` [eBPF C Program](#) is used by the networking daemon (`netd`) in Android for various such as socket filtering and statistics gathering. To see how this program is used, check the [eBPF traffic monitor](#) sources.
- The `time_in_state` [eBPF C program](#) calculates the amount of time an Android app spends at different CPU frequencies, which is used to calculate power.
- In Android 12, the `gpu_mem` [eBPF C program](#) tracks total GPU memory usage for each process and for the entire system. The program is used for GPU memory profiling.

Enterprises Multi-Cloud Networking

S&P Global



eBPF, a road to invisible network: S&P Global's Network Transformation Journey
eBPF Summit 2022

One more thing...



Grafana



ISOVALENT

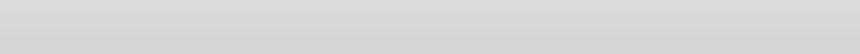
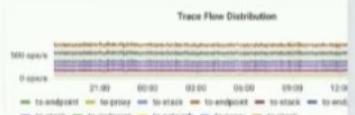
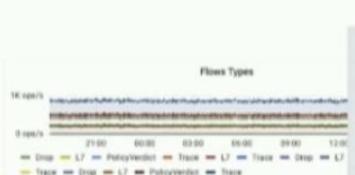
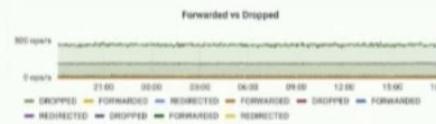
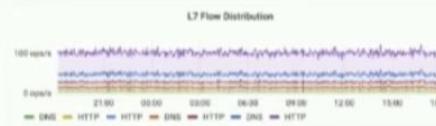
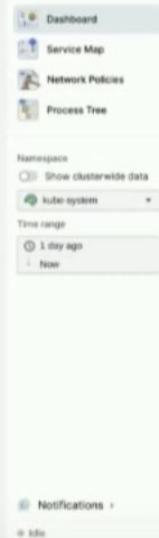


New Strategic Partnership to provide
eBPF -based Observability & Monitoring

jumpbox®



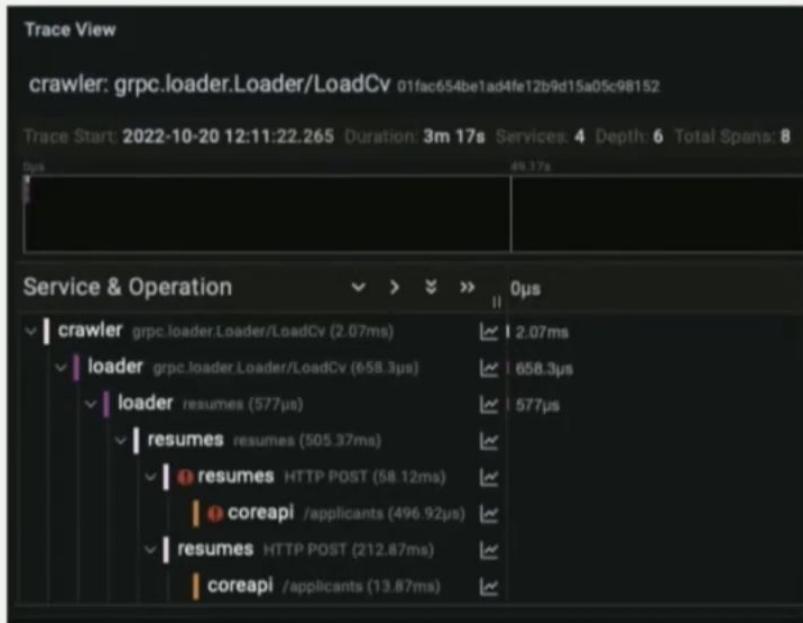
eBPF-based Network Observability



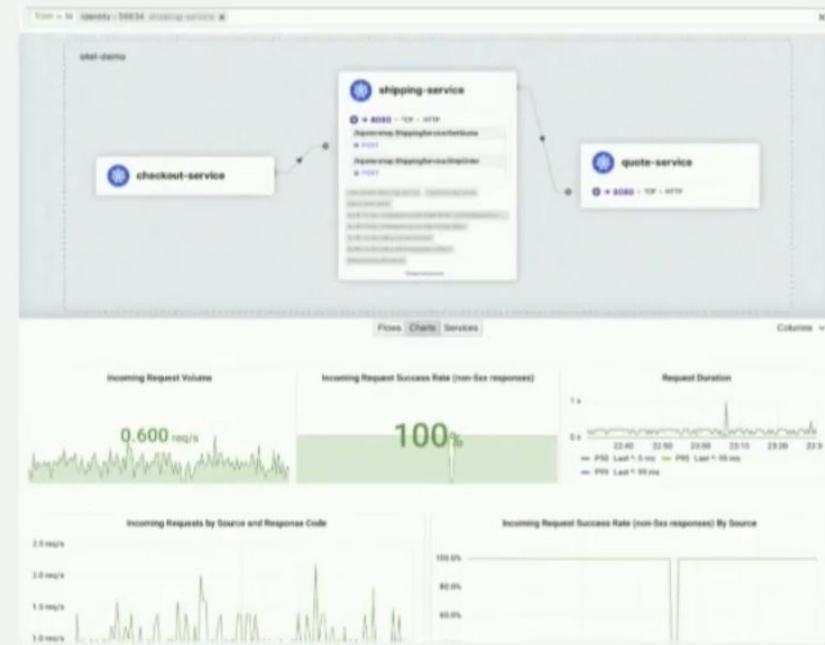
Embedded dashboards
in hubble-ui

Jumpbox®

eBPF-based Tracing & HTTP Observability



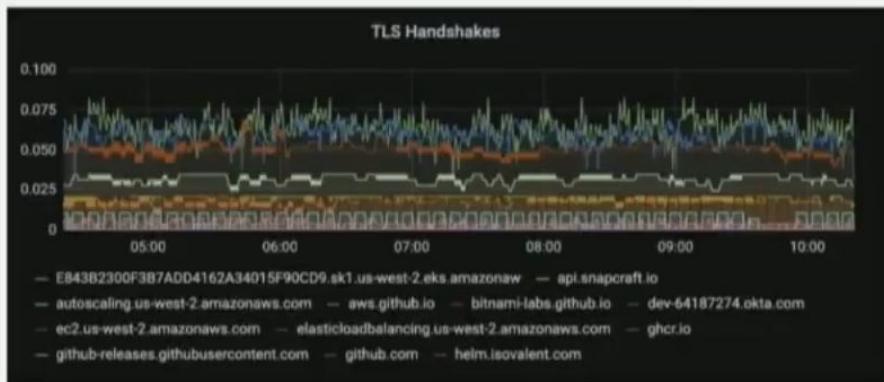
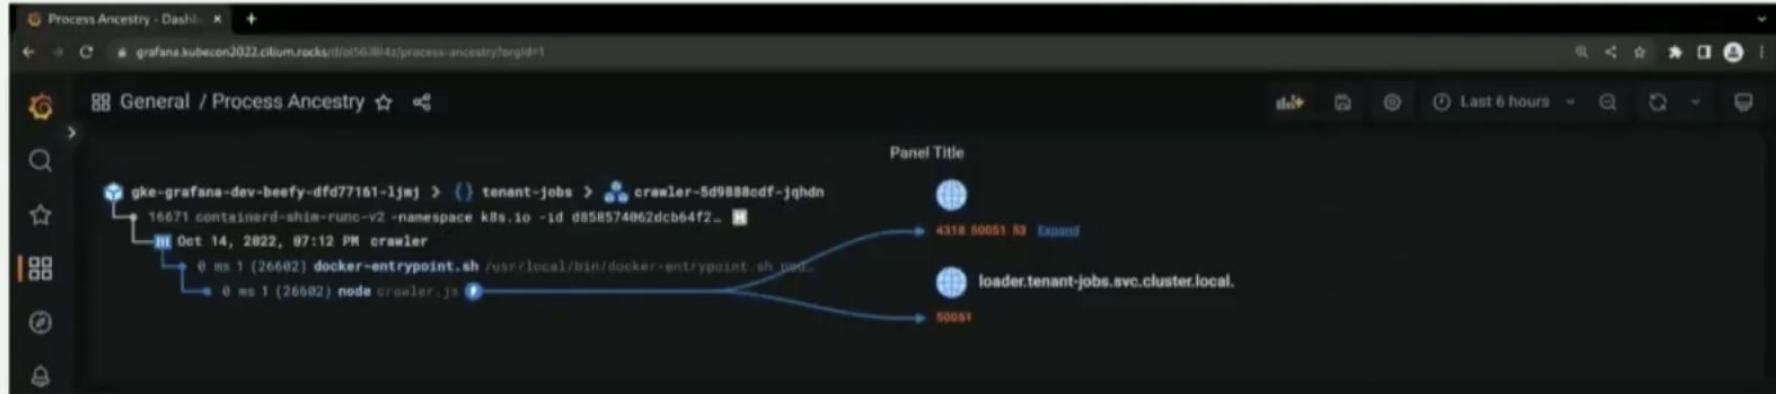
 **Grafana Tempo**



Jumpbox®



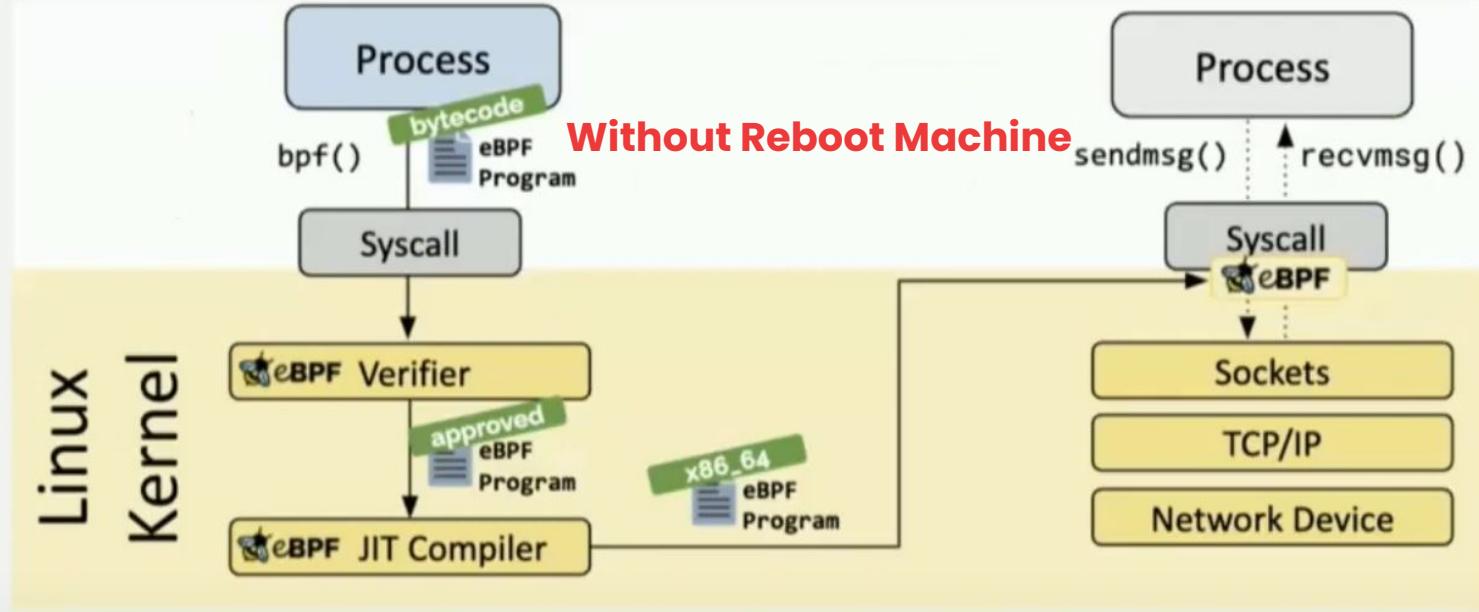
eBPF-based Security Observability



Run custom program of your choice

Run custom program of your choice
in the **Kernel**

eBPF Runtime



The runtime accepts bytecode, verifies it, just-in-time compiles it, and runs it at the requested hook point.

Let's get started with BPF

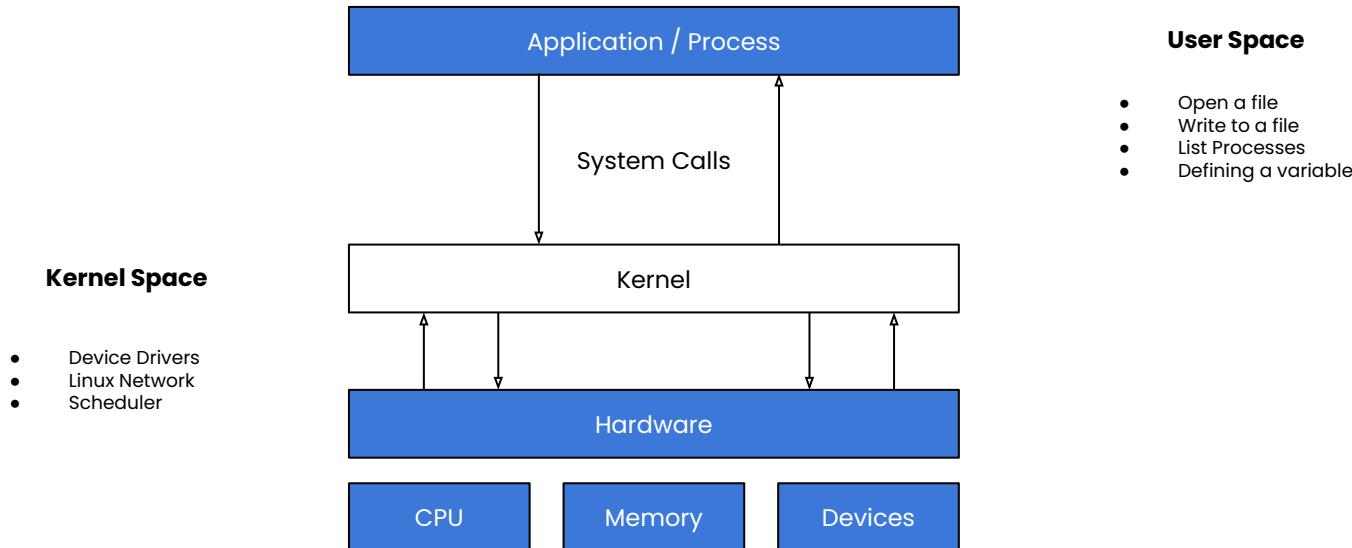
Before we go

Before Let's talk about

Before Let's talk about
User Space and **Kernel Space**

Linux Kernel (Cont.)

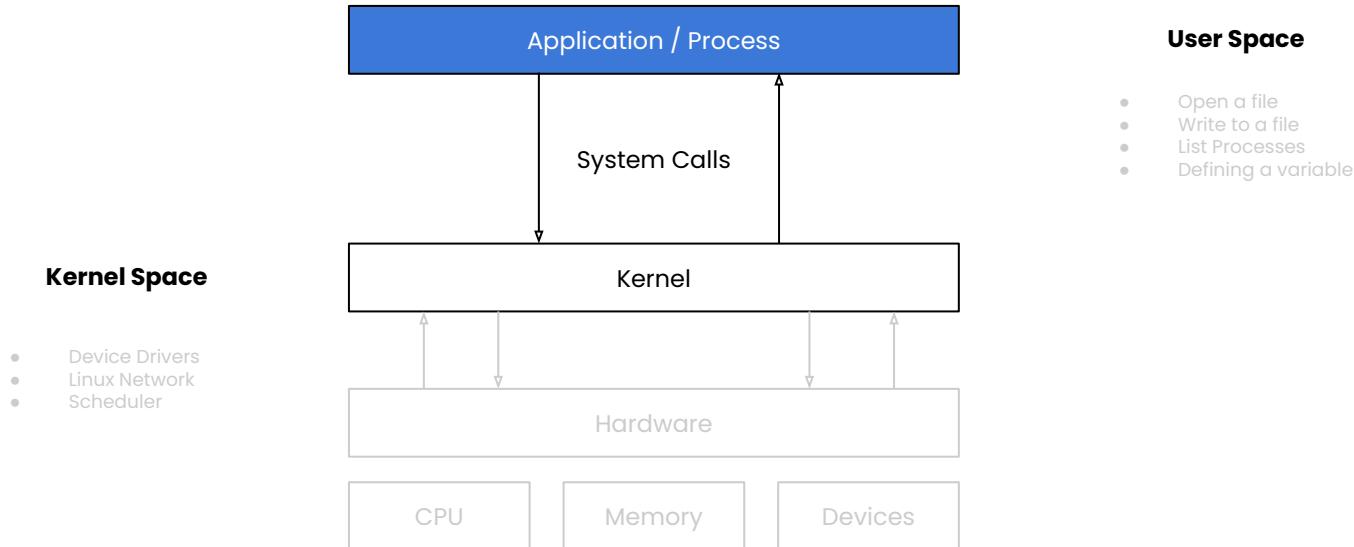
- Kernel Space and User Space



Today,
We will focus on this area

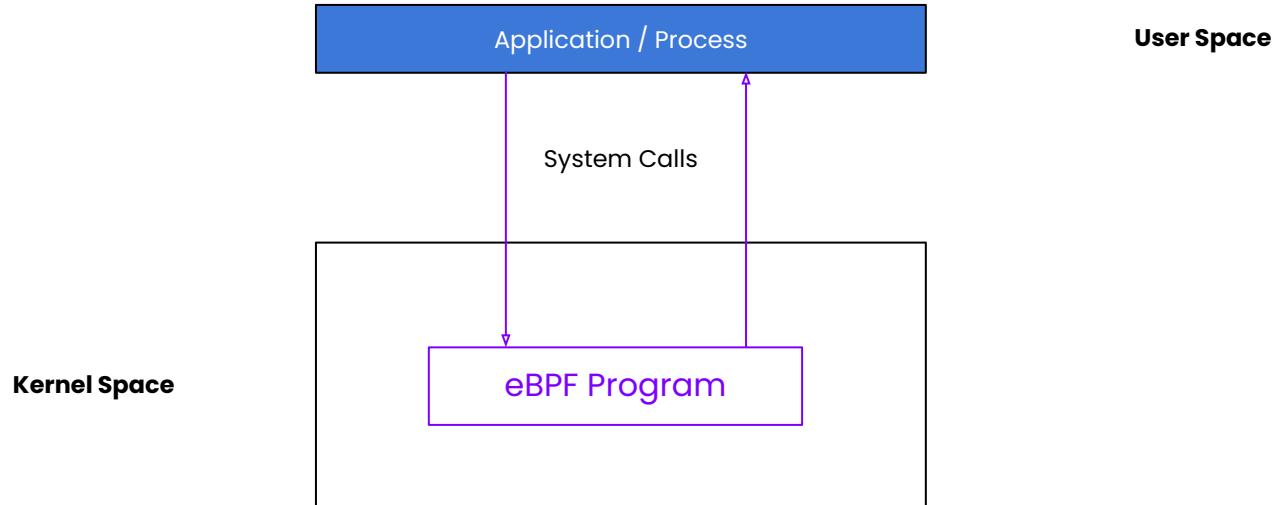
Linux Kernel (Cont.)

- Kernel Space and User Space



Map to eBPF

eBPF with logical structure



What is BPF? (Cont.)

- The bpf() system call performs a range of operations related to extended Berkeley Packet Filters. Extended BPF (eBPF) is similar to the original (“classic”) BPF (cBPF) used to filter network packets.
- For both cBPF and eBPF programs, the kernel statically analyzes the programs (eBPF verifier) before loading them in order to ensure that they cannot harm the running system.

bpftrace

 CI passing  IRC bpftrace  igtm alerts 7  discourse 18 topics

bpftrace is a high-level tracing language for Linux enhanced Berkeley Packet Filter (eBPF) available in recent Linux kernels (4.x). bpftrace uses LLVM as a backend to compile scripts to BPF-bytecode and makes use of BCC for interacting with the Linux BPF system, as well as existing Linux tracing capabilities: kernel dynamic tracing (kprobes), user-level dynamic tracing (uprobes), and tracepoints. The bpftrace language is inspired by awk and C, and predecessor tracers such as DTrace and SystemTap. bpftrace was created by [Alastair Robertson](#).

To learn more about bpftrace, see the [Reference Guide](#) and [One-Liner Tutorial](#).

One-Liners

The following one-liners demonstrate different capabilities:

```
# Files opened by process  
bpftrace -e 'tracepoint:syscalls:sys_enter_open { printf("%s %s\n", comm, str(args->filename)); }'  
  
# Syscall count by program  
bpftrace -e 'tracepoint:raw_syscalls:sys_enter { @[comm] = count(); }'  
  
# Read bytes by process:  
bpftrace -e 'tracepoint:syscalls:sys_exit_read /args->ret/ { @[comm] = sum(args->ret); }'  
  
# Read size distribution by process:  
bpftrace -e 'tracepoint:syscalls:sys_exit_read { @[comm] = hist(args->ret); }'  
  
# Show per-second syscall rates:  
bpftrace -e 'tracepoint:raw_syscalls:sys_enter { @ = count(); } interval:s:1 { print(@); clear(@); }
```

Explore bpf syscalls in bpfttrace

```
vagrant:vagrant:~/libbpfgobeginners (ssh vagrant@127.0.0.1 -p 2222 -o LogLevel=FATAL -o Compression=yes -o DSAAuthentication=yes -o IdentitiesOnly=yes -o StrictHostKeyChecking=no -o UserKnownHostsFile=/dev/null -i /Users/liz/vagrant/ubuntu-)
```

```
vagrant:vagrant:~/libbpfgobeginners$ sudo bpftrace -e 'tracepoint:raw_syscalls:sys_enter { @[comm] = count(); }'
```

```
Attaching 1 probe ...
```

```
^C
```

```
@[vmstats]: 2  
@[memballoon]: 2  
@[timesync]: 4  
@[sudo]: 7  
@[multipathd]: 22  
@[true]: 29  
@[cpptools-srv]: 35  
@[irqbalance]: 38  
@[sshd]: 44  
@[gmain]: 44  
@[bpftrace]: 47  
@[gopls]: 65  
@[docker-init]: 96  
@[cpptools]: 97  
@[dbus-daemon]: 97  
@[vminfo]: 100  
@[systemd-logind]: 108  
@[containerd]: 220  
@[kubelet]: 235  
@[rs:main Q:Reg]: 252  
@[grep]: 324  
@[in:imuxsock]: 465  
@[systemd-run]: 477  
@[systemd-journal]: 710  
@[sh]: 903  
@[systemd]: 1052  
@[node]: 1956  
@[ls]: 2271  
@[dockerd]: 2639  
@[kubelet]: 2881  
@[runc]: 4610
```



```
vagrant@vagrant:~/libbpfgo-beginners$ sudo strace -e bpf bpfftrace -e 'tracepoint:raw_syscalls:sys_enter { @[comm] = count(); }'  
bpf(BPF_MAP_CREATE, {map_type=BPF_MAP_TYPE_ARRAY, key_size=4, value_size=4, max_entries=1, map_flags=0, inner_map_fd=0, map_name="  
map_ifindex=0, btf_fd=0, btf_key_type_id=0, btf_value_type_id=0}, 120) = 3  
bpf(BPF_MAP_CREATE, {map_type=BPF_MAP_TYPE_PERCPU_HASH, key_size=16, value_size=8, max_entries=4096, map_flags=0, inner_map_fd=0,  
map_name="@", map_ifindex=0, btf_fd=0, btf_key_type_id=0, btf_value_type_id=0}, 120) = -1 EINVAL (Invalid argument)  
bpf(BPF_MAP_CREATE, {map_type=BPF_MAP_TYPE_PERCPU_HASH, key_size=16, value_size=8, max_entries=4096, map_flags=0, inner_map_fd=0,  
map_name="", map_ifindex=0, btf_fd=0, btf_key_type_id=0, btf_value_type_id=0}, 120) = 3  
bpf(BPF_MAP_CREATE, {map_type=BPF_MAP_TYPE_PERF_EVENT_ARRAY, key_size=4, value_size=4, max_entries=2, map_flags=0, inner_map_fd=0,  
map_name="printf", map_ifindex=0, btf_fd=0, btf_key_type_id=0, btf_value_type_id=0}, 120) = 4  
Attaching 1 probe ...  
bpf(BPF_MAP_UPDATE_ELEM, {map_fd=4, key=0x7ffe26e88d4c, value=0x7ffe26e88d50, flags=BPF_ANY}, 120) = 0  
bpf(BPF_MAP_UPDATE_ELEM, {map_fd=4, key=0x7ffe26e88d4c, value=0x7ffe26e88d50, flags=BPF_ANY}, 120) = 0  
bpf(BPF_PROG_LOAD, {prog_type=BPF_PROG_TYPE_TRACEPOINT, insn_cnt=27, insns=0x7efeb74ce000, license="GPL", log_level=0, log_size=0,  
log_buf=NULL, kern_version=KERNEL_VERSION(5, 8, 17), prog_flags=0, prog_name="sys_enter", prog_ifindex=0, expected_attach_type=BP  
F_CGROUP_INET_INGRESS, prog_btf_fd=0, func_info_rec_size=0, func_info=NULL, func_info_cnt=0, line_info_rec_size=0, line_info=NULL,  
line_info_cnt=0, attach_btf_id=0, attach_prog_fd=0}, 120) = 9  
^Cstrace: Process 1688560 detached
```

```
vagrant@vagrant:~/libbpfgo-beginners$
```

```
@[vmstats]: 2  
@[memballoon]: 2  
@[gmain]: 5  
@[systemd-journal]: 5  
@[sudo]: 7  
@[multipathd]: 8  
@[cppools-srv]: 16  
@[dockerd]: 24  
@[containerd]: 28  
@[bpfftrace]: 28  
@[gopls]: 35  
@[irqbalance]: 38  
@[cppools]: 48  
@[sshd]: 53  
@[grep]: 108  
@[strace]: 191
```

```
Attaching 1 probe ...
bpf(BPF_MAP_UPDATE_ELEM, {map_fd=4, key=0x7ffe26e88d4c, value=0x7ffe26e88d50, flags=BPF_ANY}, 120) = 0
bpf(BPF_MAP_UPDATE_ELEM, {map_fd=4, key=0x7ffe26e88d4c, value=0x7ffe26e88d50, flags=BPF_ANY}, 120) = 0
bpf(BPF_PROG_LOAD, {prog_type=BPF_PROG_TYPE_TRACEPOINT, insn_cnt=27, insns=0x7efeb74ce000, license="GPL", log_level=0, log_size=0,
log_buf=NULL, kern_version=KERNEL_VERSION(5, 8, 17), prog_flags=0, prog_name="sys_enter", prog_ifindex=0, expected_attach_type=BP
F_CGROUP_INET_INGRESS, prog_btf_fd=0, func_info_rec_size=0, func_info=NULL, func_info_cnt=0, line_info_rec_size=0, line_info=NULL,
line_info_cnt=0, attach_btf_id=0, attach_prog_fd=0}, 120) = 9
^Cstrace: Process 1688560 detached
```

```
map_name pointer, map_index 0, bpf_id 0, bpf_key_type_id 0, bpf_value_type_id 0, 120) .
Attaching 1 probe ...
bpf(BPF_MAP_UPDATE_ELEM, {map_fd=4, key=0x7ffe26e88d4c, value=0x7ffe26e88d50, flags=BPF_ANY}, 120) = 0
bpf(BPF_MAP_UPDATE_ELEM, {map_fd=4, key=0x7ffe26e88d4c, value=0x7ffe26e88d50, flags=BPF_ANY}, 120) = 0
bpf(BPF_PROG_LOAD, {prog_type=BPF_PROG_TYPE_TRACEPOINT, insn_cnt=27, insns=0x7efeb74ce000, license="GPL", log_level=0, log_size=0,
log_buf=NULL, kern_version=KERNEL_VERSION(5, 8, 17), prog_flags=0, prog_name="sys_enter", prog_ifindex=0, expected_attach_type=BP
F_CGROUP_INET_INGRESS, prog_btf_fd=0, func_info_rec_size=0, func_info=NULL, func_info_cnt=0, line_info_rec_size=0, line_info=NULL,
line_info_cnt=0, attach_btf_id=0, attach_prog_fd=0}, 120) = 9
^Cstrace: Process 1688560 detached
```

eBPF programs & maps

bpf(BPF_PROG_LOAD, ...)

bpf(BPF_MAP_CREATE, ...)

eBPF Programs (Cont.)

- eBPF programs can be written in a restricted C that is compiled (using the clang compiler) into eBPF bytecode.
- Various features are omitted from this restricted C, such as loops, global variables, variadic functions, floating-point numbers, and passing structures as function arguments.



Reference:

- [restricted C](#)

eBPF Helper function (Cont.)

[eBPF Helper functions] are used by eBPF programs to interact with the system, or with the context in which they work.

For instance, they can be used to print debugging messages:

```
bpf_trace_printk()  
bpf_get_current_comm()  
bpf_perf_event_output()  
...
```

Reference:

- [eBPF Helper function](#)

eBPF Map (Cont.)

Maps are a generic data structure for the storage of different types of data.

They allow **sharing of data** between eBPF kernel programs and also between **kernel and user-space applications**.



Reference:

- [eBPF Helper function](#)

Attaching eBPF to events

eBPF programs are event-driven and are run when the kernel or an application passes a certain hook point.

Pre-defined hooks include system calls, function entry/exit, kernel tracepoints, network events, and several others.

If a predefined hook does not exist for a particular need, it is possible to create a kernel probe (kprobe) or user probe (uprobe) to attach eBPF program almost anywhere in kernel or user applications.

Reference:

- [Program type BPF_PROG_TYPE_KPROBE](#)

Attach custom code to an event

```
bpf(BPF_PROG_LOAD, ...) = x  
perf_event_open(...) = y  
ioctl(y, PERF_EVENT_IOC_SET_BPF, x)
```

vagrant@vagrant:~/libbpfgobeginners (ssh vagrant@127.0.0.1 -p 2222 -o LogLevel=FATAL -o Compression=yes -o DSAAuthentication=yes -o IdentitiesOnly=yes -o StrictHostKeyChecking=no -o UserKnownHostsFile=/dev/null -i /Users/liz/vagrant/ubuntu-20.10.v...

```
vagrant@vagrant:~/libbpfgobeginners$ sudo strace -e bpf,perf_event_open,ioctl bpftrace -e 'tracepoint:raw_syscalls:sys_enter { @[comm] = count(); }'  
bpf(BPF_MAP_CREATE, {map_type=BPF_MAP_TYPE_ARRAY, key_size=4, value_size=4, max_entries=1, map_flags=0, inner_map_fd=0, map_name="", map_ifindex=0, btf_fd=0, btf_key_type_id=0, btf_value_type_id=0}, 120) = 3  
bpf(BPF_MAP_CREATE, {map_type=BPF_MAP_TYPE_PERCPU_HASH, key_size=16, value_size=8, max_entries=4096, map_flags=0, inner_map_fd=0, map_name="@", map_ifindex=0, btf_fd=0, btf_key_type_id=0, btf_value_type_id=0}, 120) = -1 EINVAL (Invalid argument)  
bpf(BPF_MAP_CREATE, {map_type=BPF_MAP_TYPE_PERCPU_HASH, key_size=16, value_size=8, max_entries=4096, map_flags=0, inner_map_fd=0, map_name="", map_ifindex=0, btf_fd=0, btf_key_type_id=0, btf_value_type_id=0}, 120) = 3  
bpf(BPF_MAP_CREATE, {map_type=BPF_MAP_TYPE_PERF_EVENT_ARRAY, key_size=4, value_size=4, max_entries=2, map_flags=0, inner_map_fd=0, map_name="printf", map_ifindex=0, btf_fd=0, btf_key_type_id=0, btf_value_type_id=0}, 120) = 4  
Attaching 1 probe ...  
perf_event_open({type=PERF_TYPE_SOFTWARE, size=0 /* PERF_ATTR_SIZE_??? */, config=PERF_COUNT_SW_BPF_OUTPUT, ... }, -1, 0, -1, PERF_
```

```
map_name="printf", map_ifindex=0, btf_fd=0, btf_key_type_id=0, btf_value_type_id=0}, 120) = 4
Attaching 1 probe ...
perf_event_open({type=PERF_TYPE_SOFTWARE, size=0 /* PERF_ATTR_SIZE_??? */, config=PERF_COUNT_SW_BPF_OUTPUT, ...}, -1, 0, -1, PERF_FLAG_FD_CLOEXEC) = 6
ioctl(6, PERF_EVENT_IOC_ENABLE, 0)      = 0
bpf(BPF_MAP_UPDATE_ELEM, {map_fd=4, key=0x7ffdf9ccbc0c, value=0x7ffdf9ccbc10, flags=BPF_ANY}, 120) = 0
perf_event_open({type=PERF_TYPE_SOFTWARE, size=0 /* PERF_ATTR_SIZE_??? */, config=PERF_COUNT_SW_BPF_OUTPUT, ...}, -1, 1, -1, PERF_FLAG_FD_CLOEXEC) = 7
ioctl(7, PERF_EVENT_IOC_ENABLE, 0)      = 0
bpf(BPF_MAP_UPDATE_ELEM, {map_fd=4, key=0x7ffdf9ccbc0c, value=0x7ffdf9ccbc10, flags=BPF_ANY}, 120) = 0
bpf(BPF_PROG_LOAD, {prog_type=BPF_PROG_TYPE_TRACEPOINT, insn_cnt=27, insns=0x7f88a6521000, license="GPL", log_level=0, log_size=1024, log_buf=NULL, kern_version=KERNEL_VERSION(5, 8, 17), prog_flags=0, prog_name="sys_enter", prog_ifindex=0, expected_attach_type=F_CGROUP_INET_INGRESS, prog_btf_fd=0, func_info_rec_size=0, func_info=NULL, func_info_cnt=0, line_info_rec_size=0, line_info=NULL, line_info_cnt=0, attach_btf_id=0, attach_prog_fd=0}, 120) = 9
perf_event_open({type=PERF_TYPE_TRACEPOINT, size=0 /* PERF_ATTR_SIZE_??? */, config=22, ...}, -1, 0, -1, PERF_FLAG_FD_CLOEXEC) = 10
ioctl(8, PERF_EVENT_IOC_SET_BPF, 9)      = 0
ioctl(8, PERF_EVENT_IOC_ENABLE, 0)      = 0
^Cstrace: Process 1690732 detached
```

vagrant@vagrant:~/libbpfgo-beginners\$

```
@[multipathd]: 4
@[systemd-journal]: 5
@[sudo]: 7
@[cpptools-srv]: 10
@[gopls]: 11
@[bpftrace]: 16
@[containerd]: 20
```

```
map_name="printf", map_ifindex=0, btf_fd=0, btf_key_type_id=0, btf_value_type_id=0}, 120) = 4
Attaching 1 probe ...
perf_event_open({type=PERF_TYPE_SOFTWARE, size=0 /* PERF_ATTR_SIZE_??? */, config=PERF_COUNT_SW_BPF_OUTPUT, ... }, -1, 0, -1, PERF_FLAG_FD_CLOEXEC) = 6
ioctl(6, PERF_EVENT_IOC_ENABLE, 0)      = 0
bpf(BPF_MAP_UPDATE_ELEM, {map_fd=4, key=0x7ffdf9ccbc0c, value=0x7ffdf9ccbc10, flags=BPF_ANY}, 120) = 0
perf_event_open({type=PERF_TYPE_SOFTWARE, size=0 /* PERF_ATTR_SIZE_??? */, config=PERF_COUNT_SW_BPF_OUTPUT, ... }, -1, 1, -1, PERF_FLAG_FD_CLOEXEC) = 7
ioctl(7, PERF_EVENT_IOC_ENABLE, 0)      = 0
bpf(BPF_MAP_UPDATE_ELEM, {map_fd=4, key=0x7ffdf9ccbc0c, value=0x7ffdf9ccbc10, flags=BPF_ANY}, 120) = 0
bpf(BPF_PROG_LOAD, {prog_type=BPF_PROG_TYPE_TRACEPOINT, insn_cnt=27, insns=0x7f88a6521000, license="GPL", log_level=0, log_size=0,
log_buf=NULL, kern_version=KERNEL_VERSION(5, 8, 17), prog_flags=0, prog_name="sys_enter", prog_ifindex=0, expected_attach_type=BP
F_CGROUP_INET_INGRESS, prog_btf_fd=0, func_info_rec_size=0, func_info=NULL, func_info_cnt=0, line_info_rec_size=0, line_info=NULL,
line_info_cnt=0, attach_btf_id=0, attach_prog_fd=0}, 120) = 9
perf_event_open({type=PERF_TYPE_TRACEPOINT, size=0 /* PERF_ATTR_SIZE_??? */, config=22, ... }, -1, 0, -1, PERF_FLAG_FD_CLOEXEC) = 8
ioctl(8, PERF_EVENT_IOC_SET_BPF, 9)      = 0
ioctl(8, PERF_EVENT_IOC_ENABLE, 0)      = 0
^Cstrace: Process 1690732 detached
```



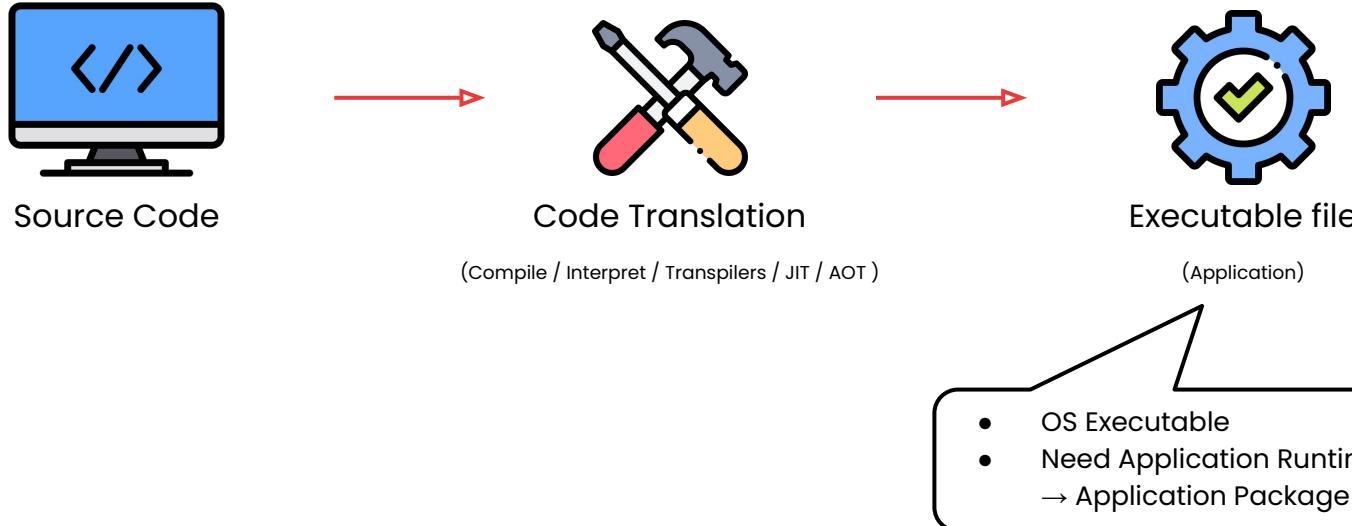
```
PERF_IOC_ENABLE),
ioctl(7, PERF_EVENT_IOC_ENABLE, 0)      = 0
bpf(BPF_MAP_UPDATE_ELEM, {map_fd=4, key=0x7ffdf9ccbc0c, value=0x7ffdf9ccbc10, flags=BPF_ANY}, 120) = 0
bpf(BPF_PROG_LOAD, {prog_type=BPF_PROG_TYPE_TRACEPOINT, insn_cnt=27, insns=0x7f88a6521000, license="GPL", log_level=0, log_size=0,
log_buf=NULL, kern_version=KERNEL_VERSION(5, 8, 17), prog_flags=0, prog_name="sys_enter", prog_ifindex=0, expected_attach_type=BP
F_CGROUP_INET_INGRESS, prog_btf_fd=0, func_info_rec_size=0, func_info=NULL, func_info_cnt=0, line_info_rec_size=0, line_info=NULL,
line_info_cnt=0, attach_btf_id=0, attach_prog_fd=0}, 120) = 9
perf_event_open({type=PERF_TYPE_TRACEPOINT, size=0 /* PERF_ATTR_SIZE_??? */, config=22, ... }, -1, 0, -1, PERF_FLAG_FD_CLOEXEC) = 8
ioctl(8, PERF_EVENT_IOC_SET_BPF, 9)      = 0
ioctl(8, PERF_EVENT_IOC_ENABLE, 0)      = 0
^Cstrace: Process 1690732 detached
```

How to write eBPF hello world?

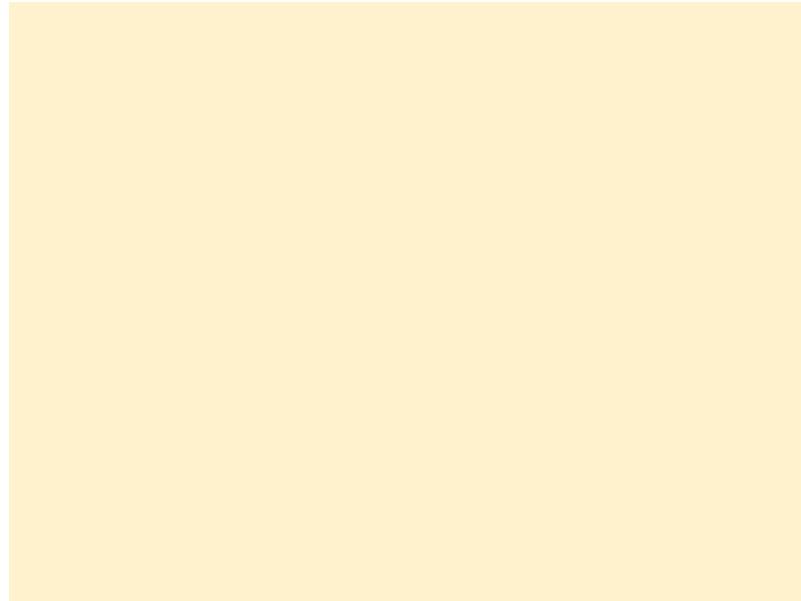
Back to Basics

Back to Basics
with Application Building

Application Building (Cont.)



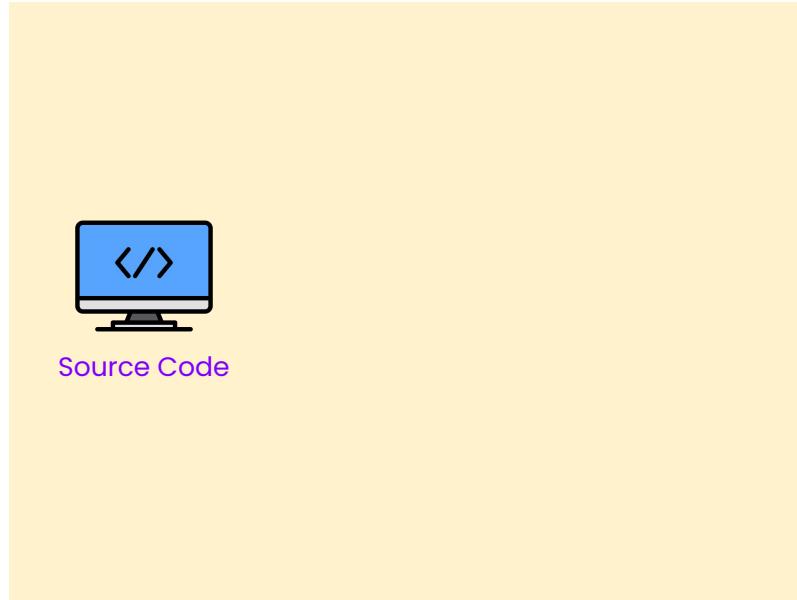
eBPF with logical structure flow (Cont.)



User Space

eBPF with logical structure flow (Cont.)

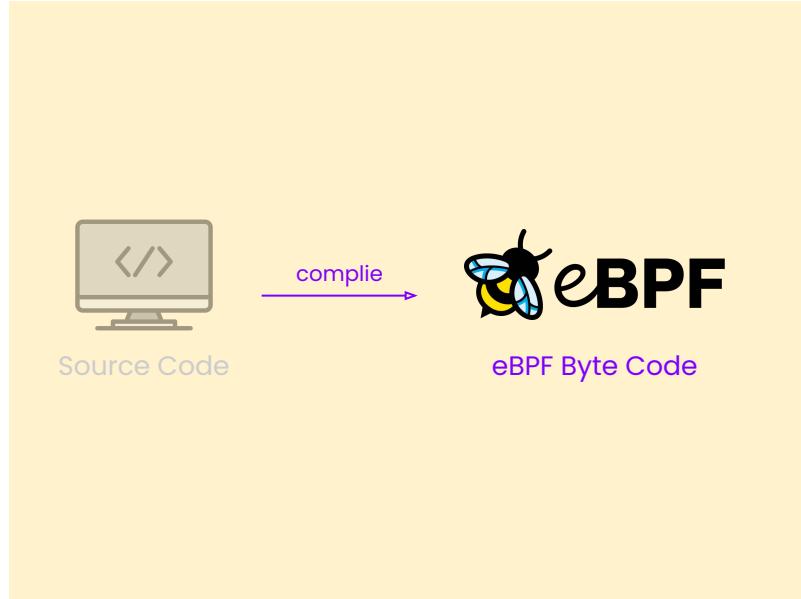
Written in our choice language



User Space

eBPF with logical structure flow (Cont.)

Written in our choice language

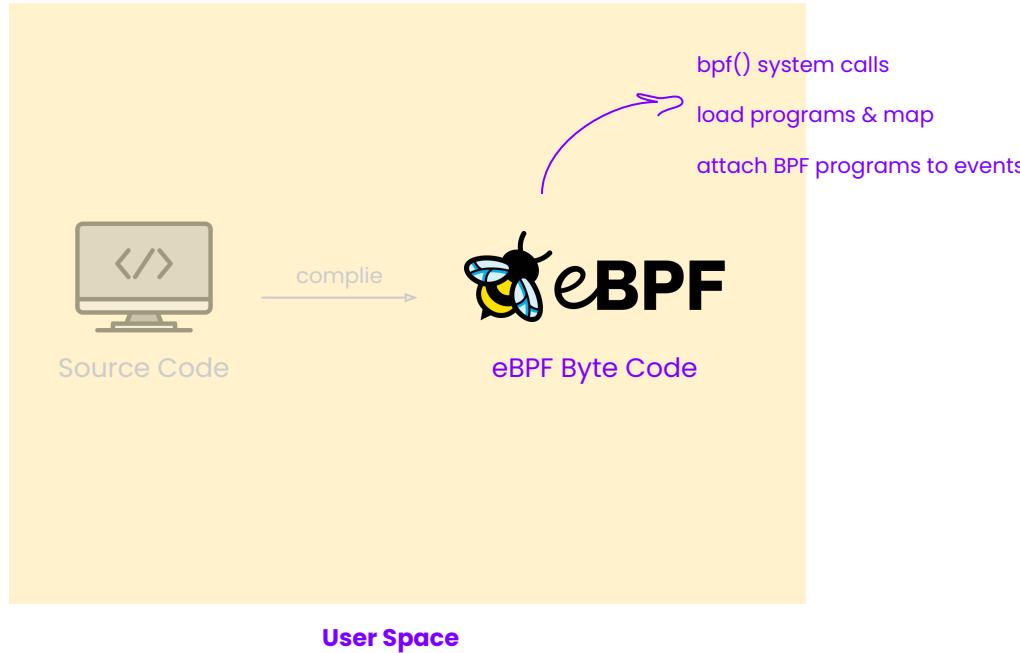


User Space

Jumpbox®

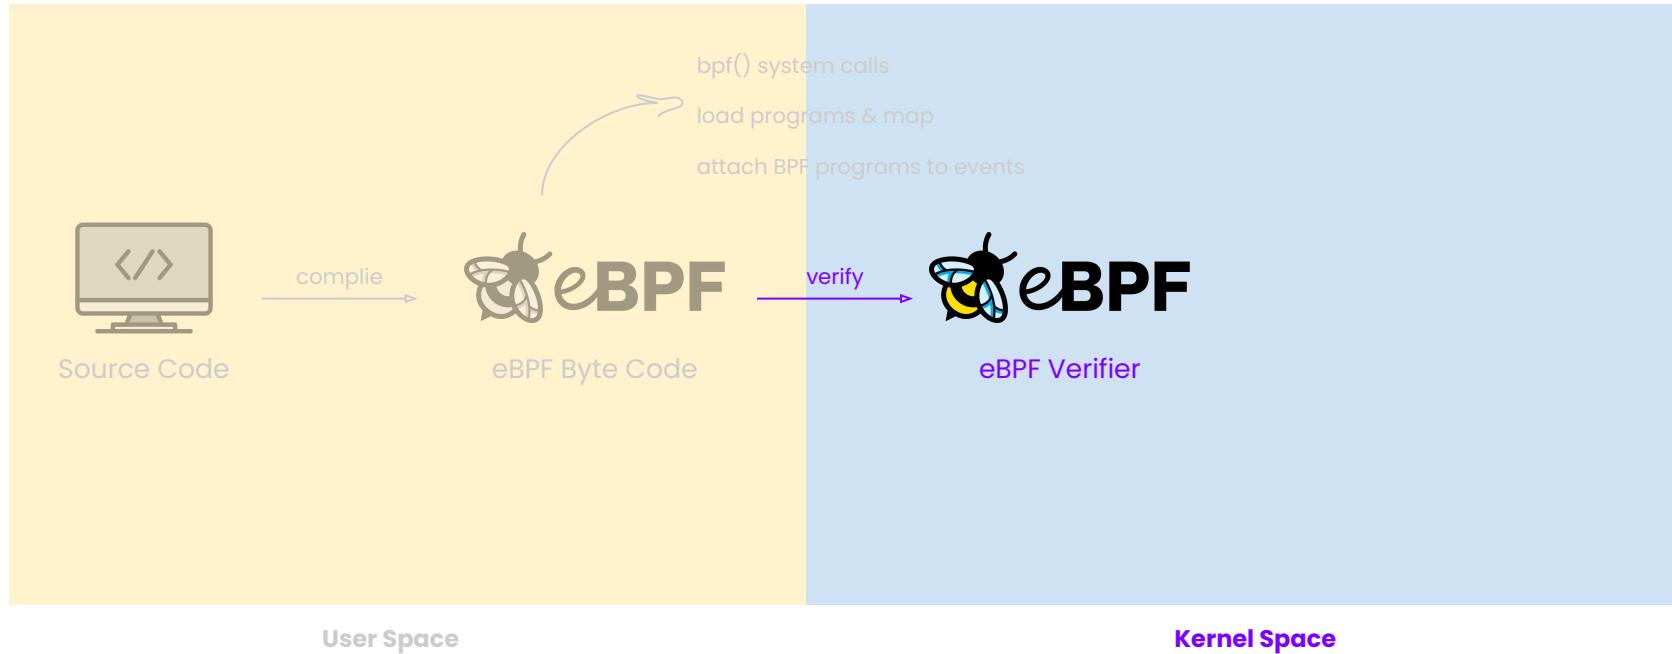
eBPF with logical structure flow (Cont.)

Written in our choice language



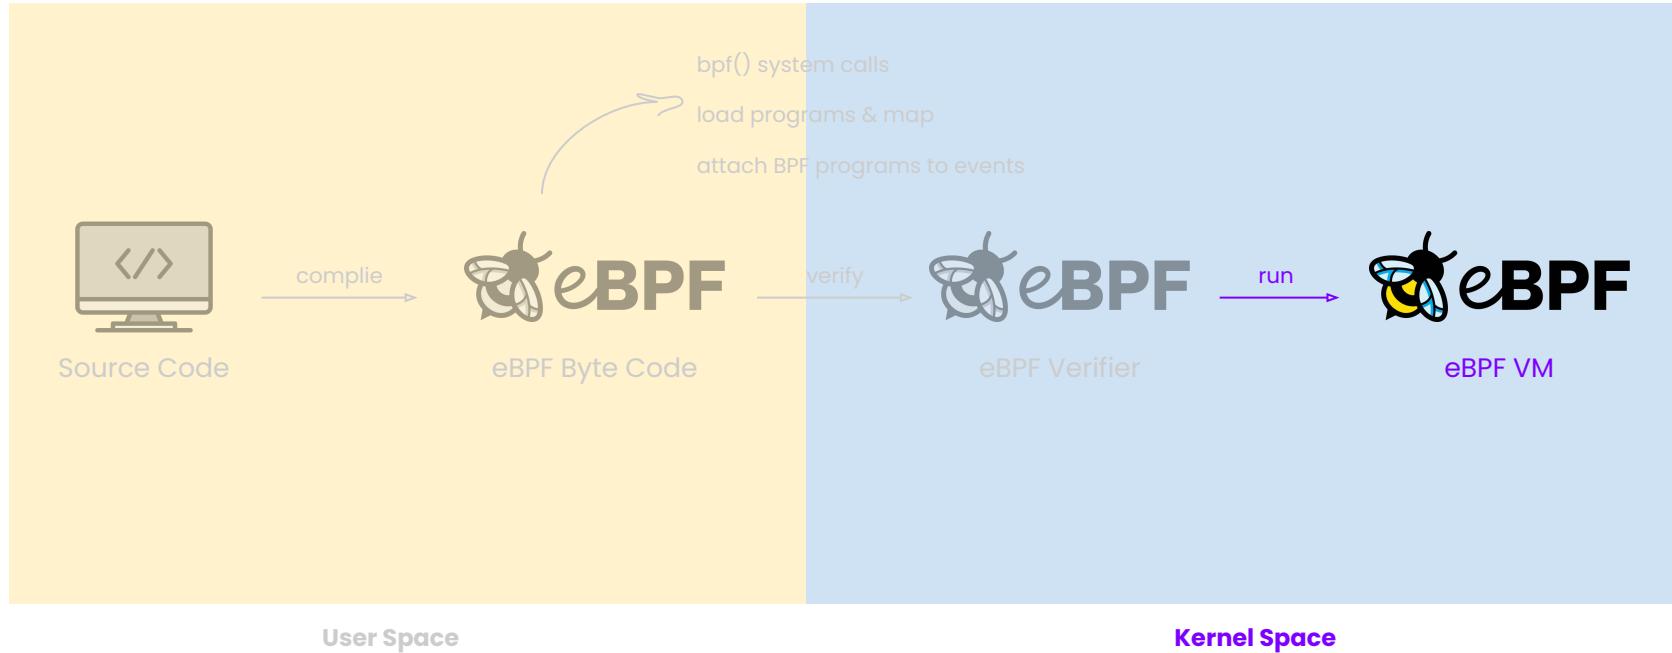
eBPF with logical structure flow (Cont.)

Written in our choice language

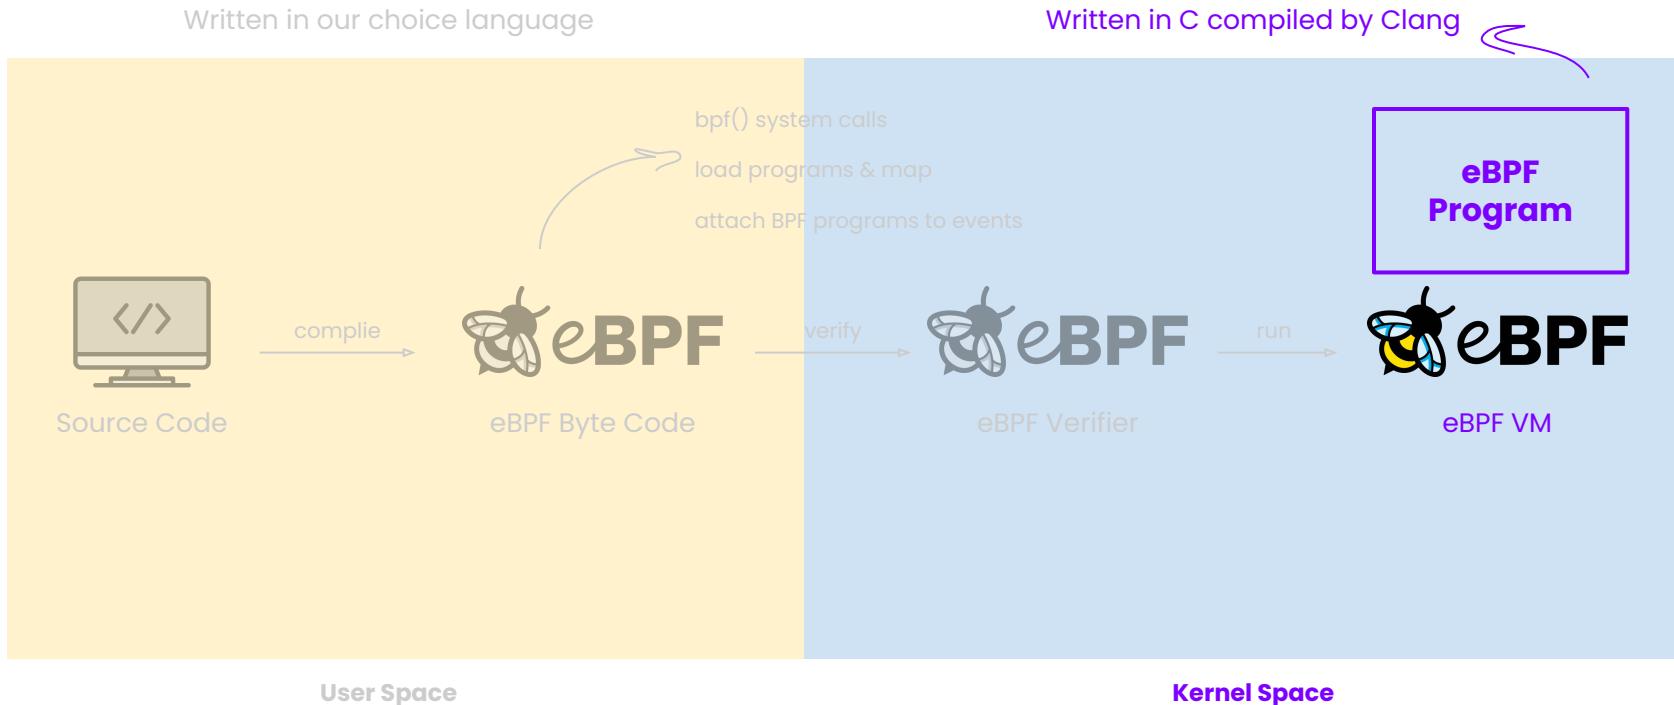


eBPF with logical structure flow (Cont.)

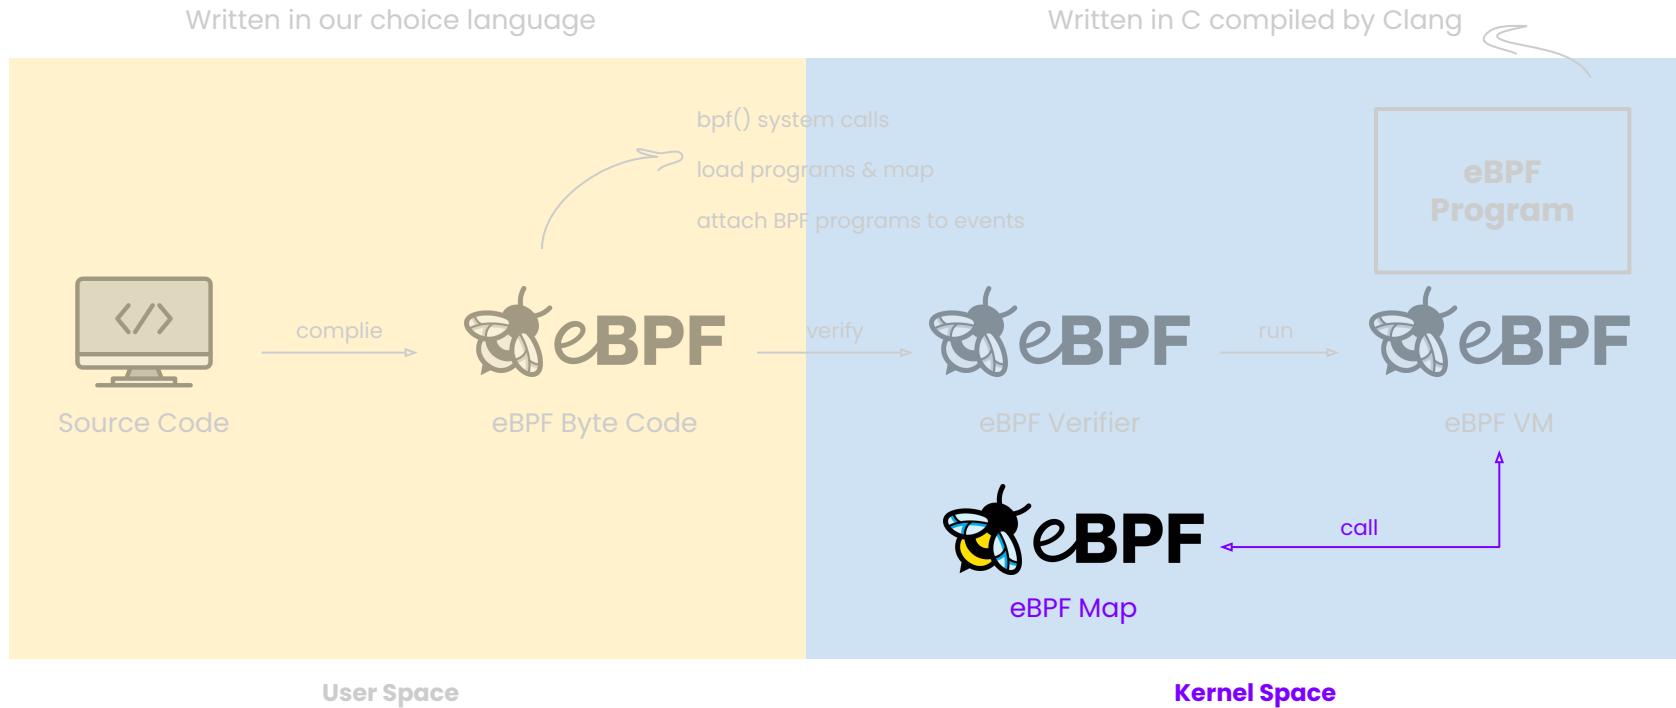
Written in our choice language



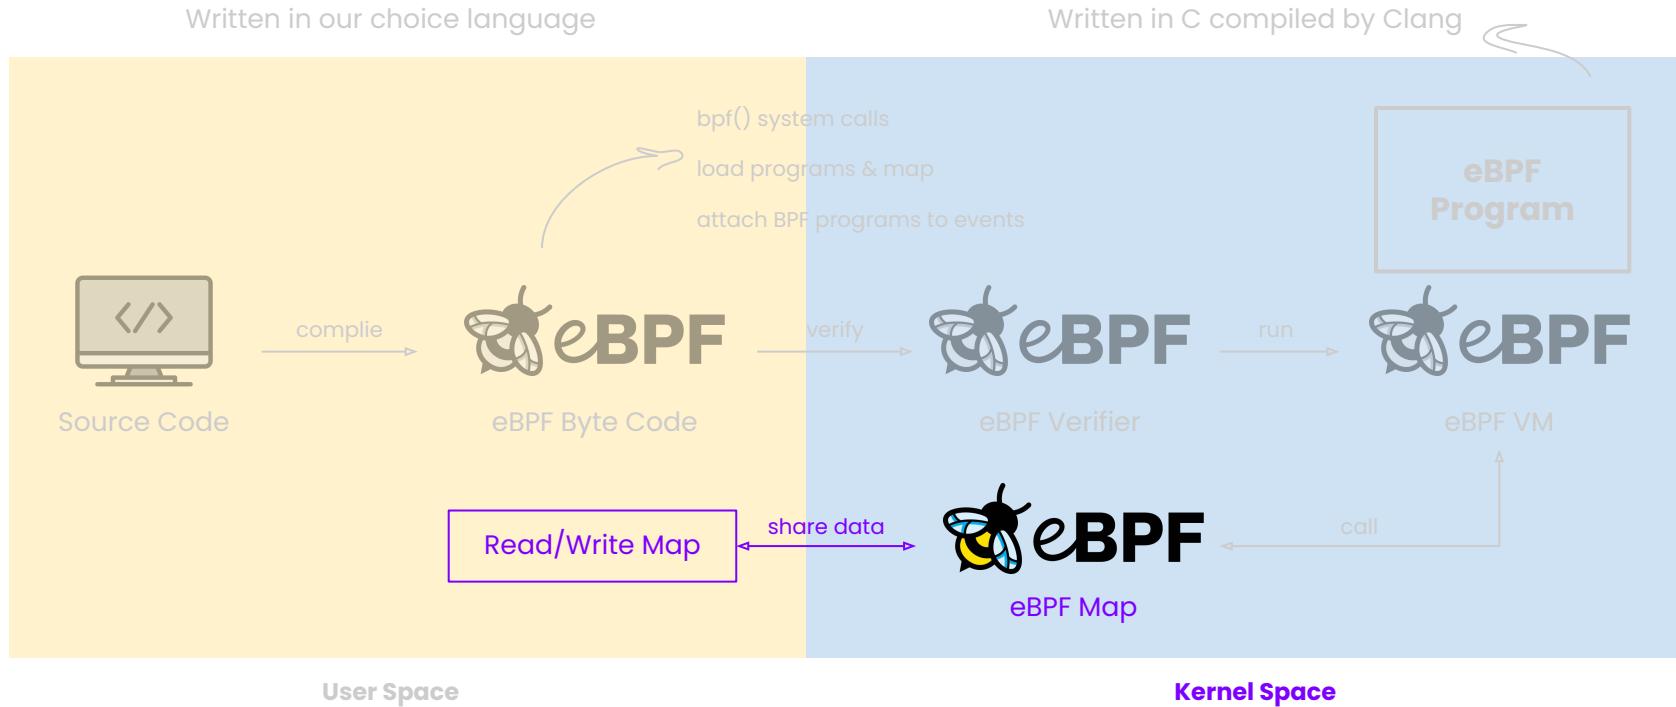
eBPF with logical structure flow (Cont.)



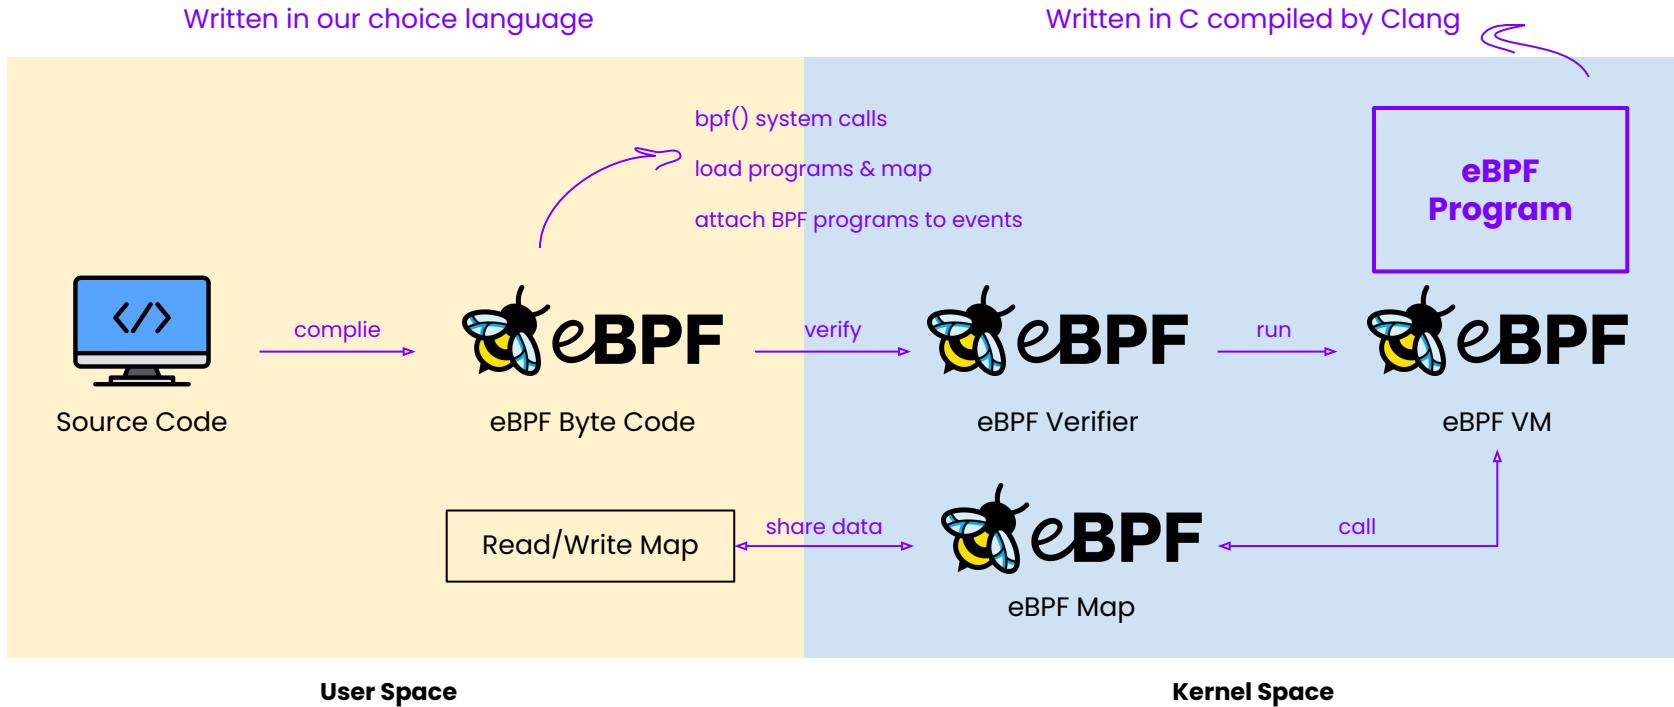
eBPF with logical structure flow (Cont.)



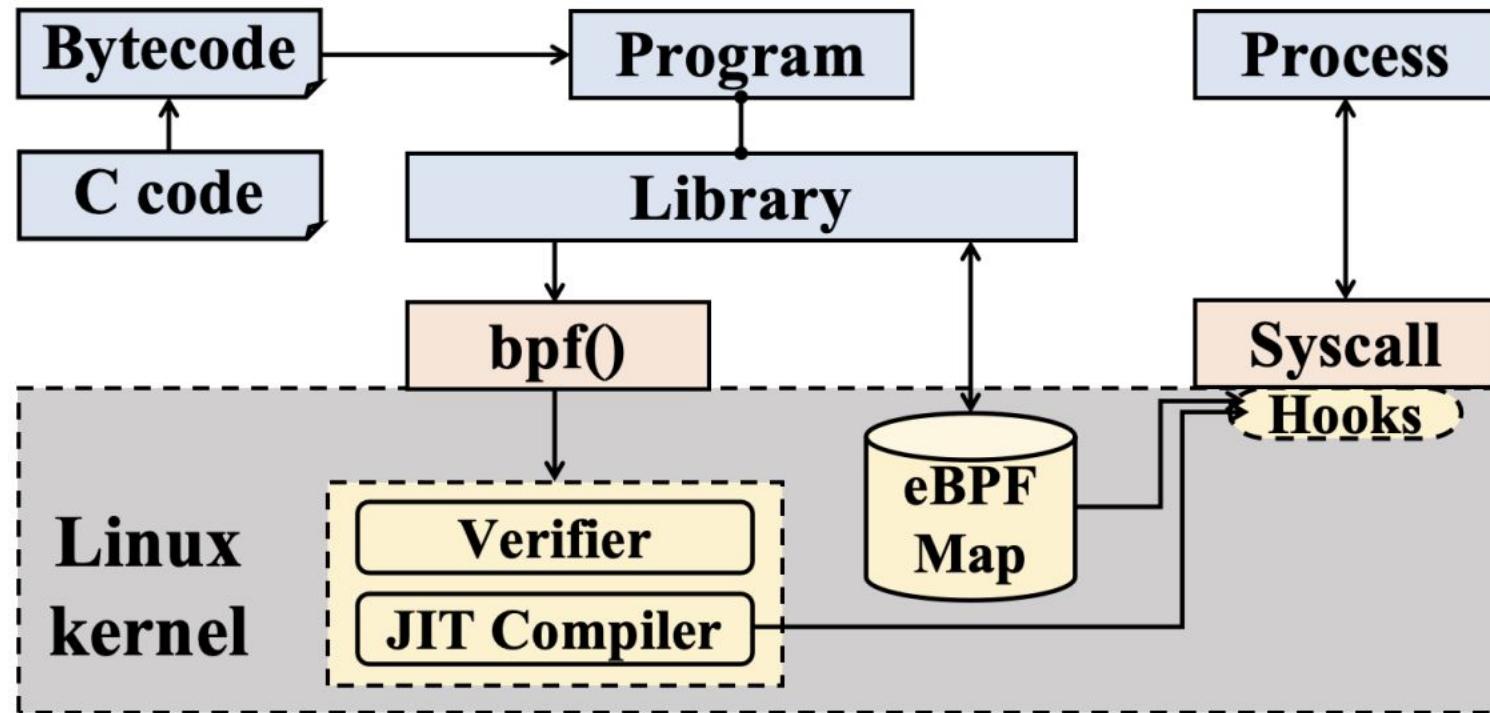
eBPF with logical structure flow (Cont.)



eBPF with logical structure flow (Cont.)



eBPF with logical structure flow (Cont.)



Reference:

- [Full Logical Structure](#)

Review Keywords

- **tracepoint:** Trace Syscalls
- **uprobe:** Userspace Probe
- **kprobe:** Kernel Function Probe
- **xdp:** Packet Filtering
- **lsm:** Linux Security Module

Can MacOS users develop eBPF?



Lima: Linux Machines

openSSF best practices passing

[Lima](#) launches Linux virtual machines with automatic file sharing and port forwarding (similar to WSL2).

The original goal of Lima was to promote [containerd](#) including [nerdctl \(contaiNERD ctl\)](#) to Mac users, but Lima can be used for non-container applications as well.

Lima also supports other container engines (Docker, Podman, Kubernetes, etc.) and non-macOS hosts (Linux, NetBSD, etc.).

Reference:

- [Lima](#)

QEMU README

QEMU is a generic and open source machine & userspace emulator and virtualizer.

QEMU is capable of emulating a complete machine in software without any need for hardware virtualization support. By using dynamic translation, it achieves very good performance. QEMU can also integrate with the Xen and KVM hypervisors to provide emulated hardware while allowing the hypervisor to manage the CPU. With hypervisor support, QEMU can achieve near native performance for CPUs. When QEMU emulates CPUs directly it is capable of running operating systems made for one machine (e.g. an ARMv7 board) on a different machine (e.g. an x86_64 PC board).

QEMU is also capable of providing userspace API virtualization for Linux and BSD kernel interfaces. This allows binaries compiled against one architecture ABI (e.g. the Linux PPC64 ABI) to be run on a host using a different architecture ABI (e.g. the Linux x86_64 ABI). This does not involve any hardware emulation, simply CPU and syscall emulation.

QEMU aims to fit into a variety of use cases. It can be invoked directly by users wishing to have full control over its behaviour and settings. It also aims to facilitate integration into higher level management layers, by providing a stable command line interface and monitor API. It is commonly invoked indirectly via the libvirt library when using open source applications such as oVirt, OpenStack and virt-manager.

QEMU as a whole is released under the GNU General Public License, version 2. For full licensing details, consult the LICENSE file.

Reference:

- [Qemu](#)

Lima & Qemu Stack (Cont.)

Lima

Linux Virtual Machine



Lima & Qemu Stack (Cont.)

Lima

Linux Virtual Machine

Container Runtime Engine (e.g. containerD)

Linux Infrastructure

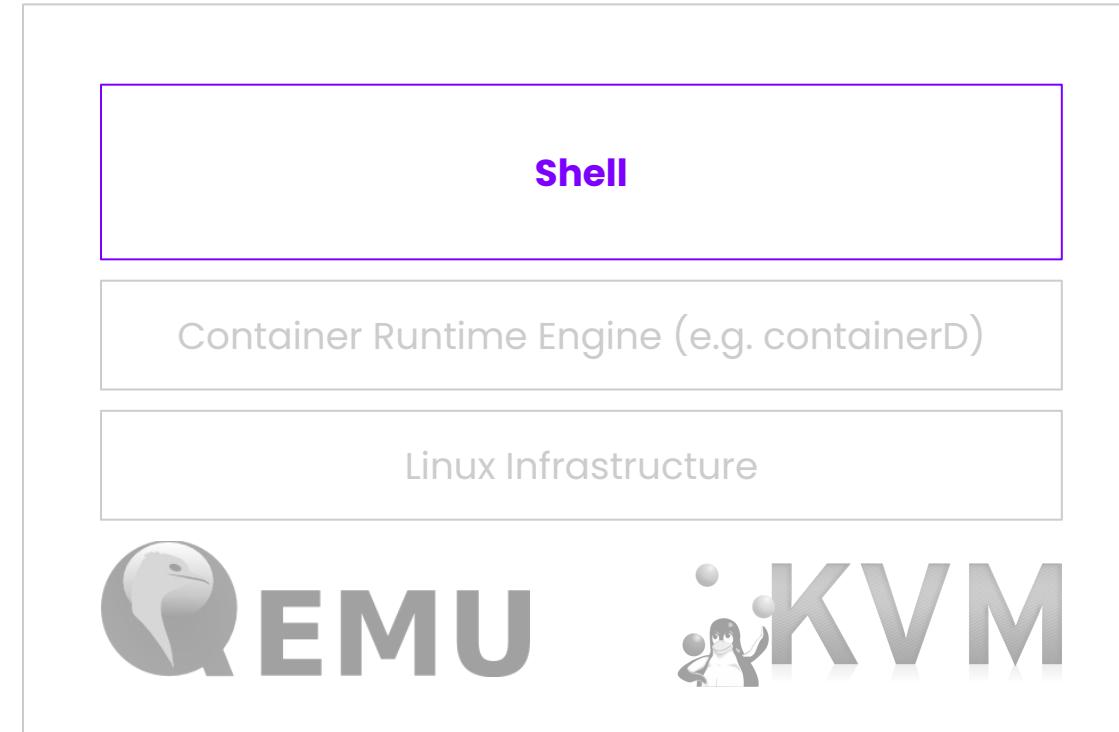


Jumpbox®

Lima & Qemu Stack (Cont.)

Lima

Linux Virtual Machine

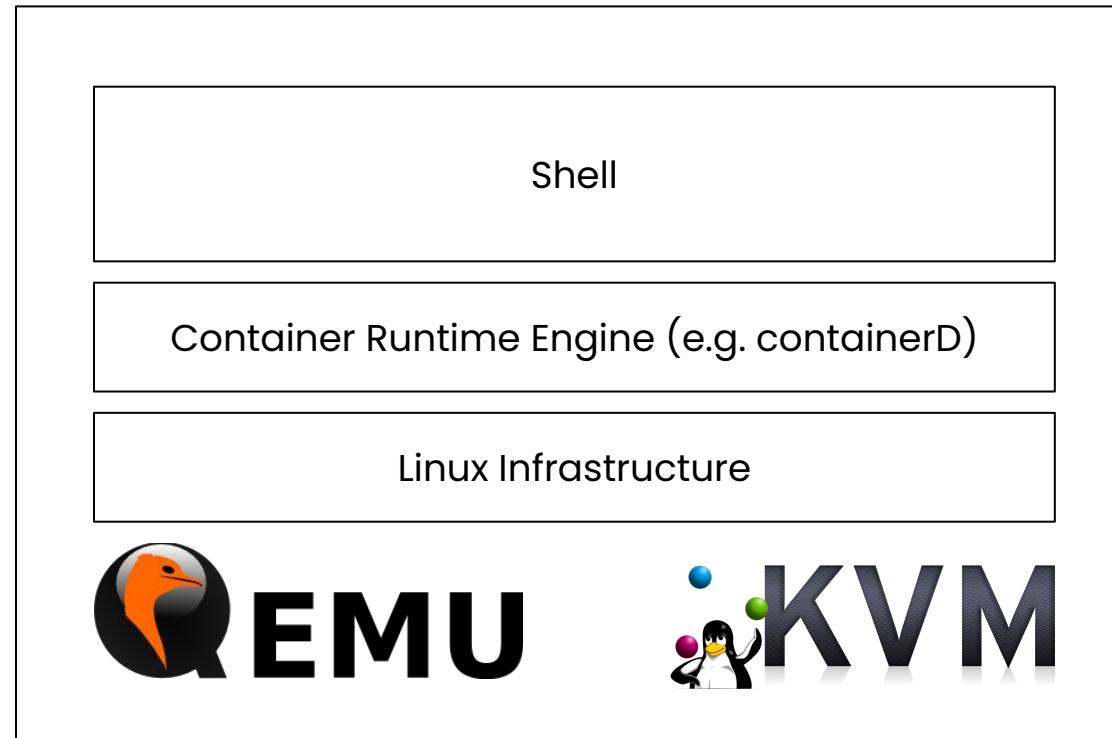


Jumpbox®

Lima & Qemu Stack (Cont.)

Lima

Linux Virtual Machine



Jumpbox®



Demo

Demo with Python (Cont.)

```
#!/usr/bin/python
from bcc import BPF

program = """
int hello_world(void *ctx) {
    bpf_trace_printk("Hello World! NongKai is here. Enjoy!\n");
    return 0;
}
"""

b = BPF(text=program)
syscall = b.get_syscall_fnname("execve")
b.attach_kprobe(event=syscall, fn_name="hello_world")

b.trace_print()
```

Reference:

- [Python: bcc Reference Guide](#)

Demo with Python (Cont.)

```
#!/usr/bin/python
from bcc import BPF
                bcc: python framework easy to write eBPF Program
program = """
int hello_world(void *ctx) {
    bpf_trace_printk("Hello World! NongKai is here. Enjoy!\n");
    return 0;
}
"""

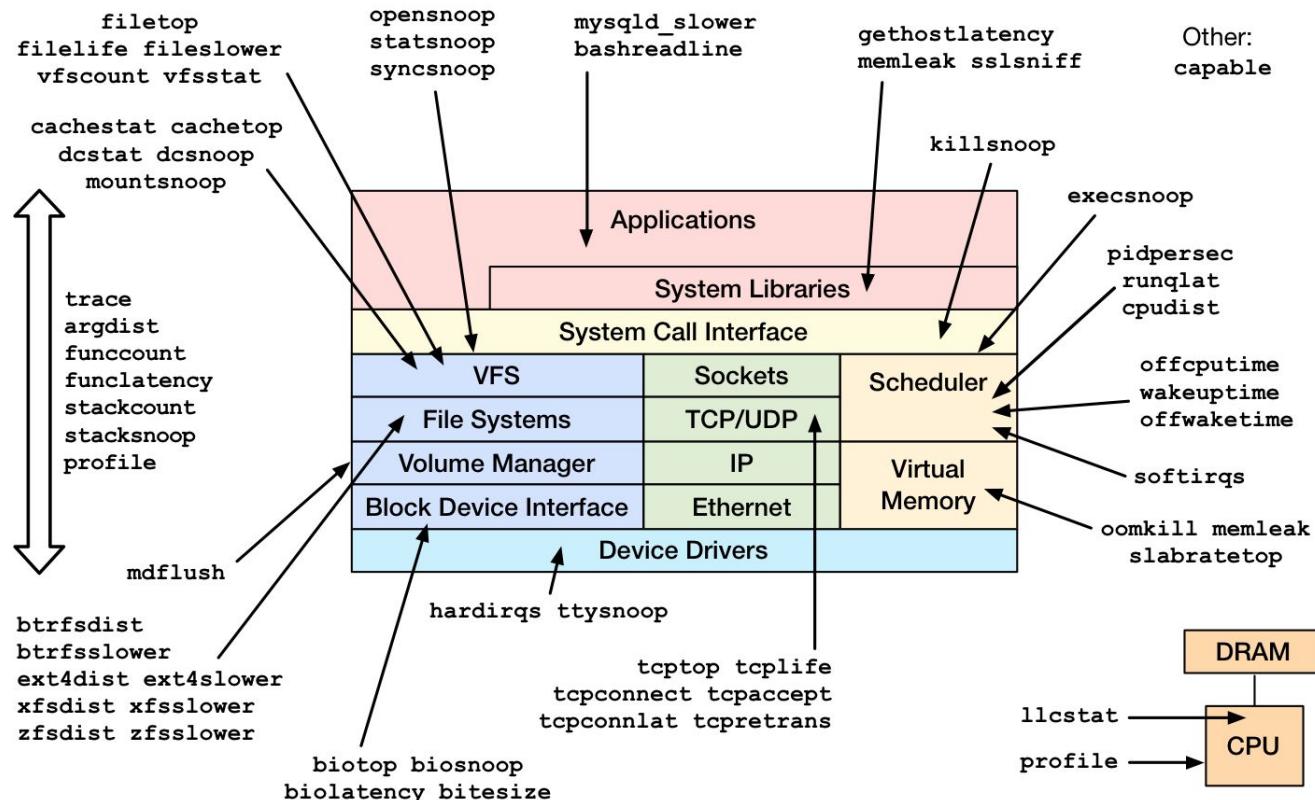
b = BPF(text=program)
syscall = b.get_syscall_fnname("execve")
b.attach_kprobe(event=syscall, fn_name="hello_world")

b.trace_print()
```

Reference:

- [Python: bcc Reference Guide](#)

Linux bcc/BPF Tracing Tools



Reference:

- [Linux BCC](#)

<https://github.com/iovisor/bcc#tools 2016>

Jumpbox®

Demo with Python (Cont.)

```
#!/usr/bin/python
from bcc import BPF
program = """
C function (hello_world):
when event invoked by system call

int hello_world(void *ctx) {
    bpf_trace_printk("Hello World! NongKai is here. Enjoy!\n");
    return 0;
}
"""

b = BPF(text=program)
syscall = b.get_syscall_fnname("execve")
b.attach_kprobe(event=syscall, fn_name="hello_world")

b.trace_print()
```

Reference:

- [Python: bcc Reference Guide](#)

Demo with Python (Cont.)

```
#!/usr/bin/python
from bcc import BPF

program = """
int hello_world(void *ctx) {
    bpf_trace_printk("Hello World! NongKai is here. Enjoy!\n");
    return 0
}
"""

b = BPF(text=program)
syscall = b.get_syscall_fnname("execve")
b.attach_kprobe(event=syscall, fn_name="hello_world")

b.trace_print()
```

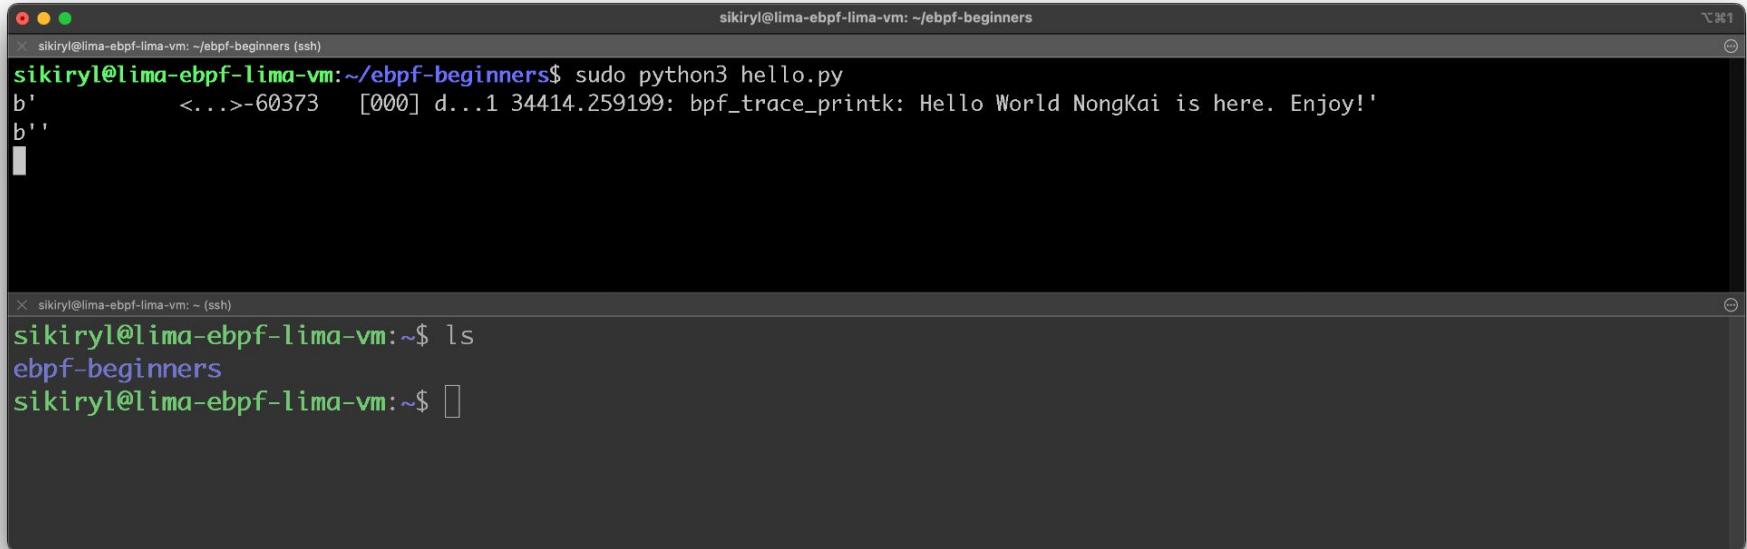
Catch event:
execute program (execve)

Reference:

- [Linux System Call: execve](#)

Demo with Python (Cont.)

Output:



```
sikiryl@lima-ebpf-lima-vm: ~/ebpf-beginners$ sudo python3 hello.py
b'          <...>-60373  [000] d...1 34414.259199: bpf_trace_printk: Hello World NongKai is here. Enjoy!'
b'
b'

sikiryl@lima-ebpf-lima-vm: ~$ ls
ebpf-beginners
sikiryl@lima-ebpf-lima-vm: ~$ 
```

Reference:

- [Linux System Call: execve](#)

Demo with Python (Cont.)

Output:

The screenshot shows two terminal sessions on a Linux system named 'lima-ebpf-lima-vm'.
Session 1 (Top): The user runs the command `sudo python3 hello.py`. The output is:
`sikiryl@lima-ebpf-lima-vm:~/ebpf-beginners$ sudo python3 hello.py
b'<...>-60373 [000] d...1 34414.259199: bpf_trace_printk: Hello World NongKai is here. Enjoy!'
b''
[]`
Session 2 (Bottom): The user runs the command `ls` to list files in the current directory. The output is:
`sikiryl@lima-ebpf-lima-vm:~$ ls
ebpf-beginners
sikiryl@lima-ebpf-lima-vm:~$`

Reference:

- [Linux System Call: execve](#)

Catch UserID execution event

```
sikiryl@lima-ebpf-lima-vm: ~/ebpf-beginners$ sudo python3 ebpf.py
No entries yet
ID 501: 1
ID 501: 2
ID 501: 2
[

sikiryl@lima-ebpf-lima-vm: ~$ ls
ebpf-beginners
sikiryl@lima-ebpf-lima-vm: ~$ id
uid=501(sikiryl) gid=1000(sikiryl) groups=1000(sikiryl)
sikiryl@lima-ebpf-lima-vm: ~$ ]
```

The screenshot shows two terminal windows side-by-side. The left window has a title bar "sikiryl@lima-ebpf-lima-vm: ~" and contains the command "id" followed by its output: "uid=501(sikiryl) gid=1000(sikiryl) groups=1000(sikiryl)". The right window has a title bar "sikiryl@lima-ebpf-lima-vm: ~/ebpf-beginners (ssh)" and contains the command "ls" followed by its output: "ebpf-beginners". Both windows have standard OS X-style window controls (red, yellow, green buttons) and a scroll bar on the right.

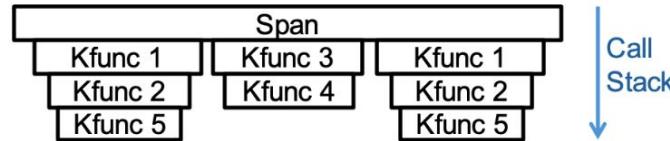
```
sikiryl@lima-ebpf-lima-vm: ~$ id
uid=501(sikiryl) gid=1000(sikiryl) groups=1000(sikiryl)
sikiryl@lima-ebpf-lima-vm: ~$ ls
ebpf-beginners
```

eBPF mentioned
at the latest KubeCon2024

OpenTelemetry Amplified- Full Observability with EBPF-Enabled Distributed Tracing

Fine-grained Traces

- Fine-grained Traces
 - Collect telemetry data in Kernel
 - Context Propagation between kernel functions
 - Merge kernel spans with request-level spans



Example Demo

■ Trace Visualization

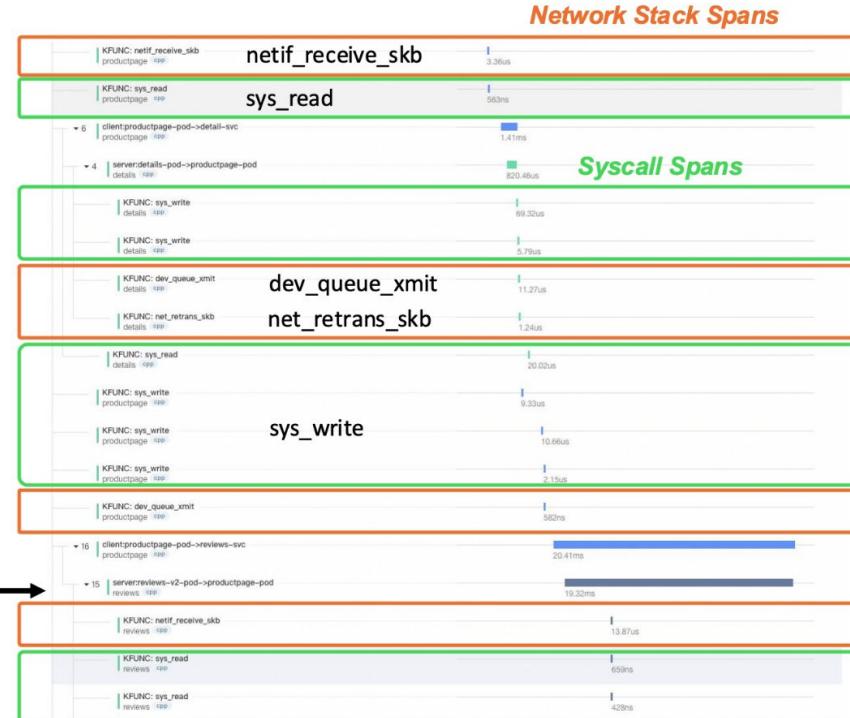
Request-level insights

Network insights

Security insights



Before



After

eBPF Topic in KubeCon2024

- [OpenTelemetry Amplified- Full Observability with EBPF-Enabled Distributed Tracing](#)
- [Redefining Service Mesh: Leveraging EBPF to Optimize Istio Ambient Architecture and Performance](#)
- [KubeSkoop- Deal with the Complexity of Network Issues and Monitoring with eBPF](#)
- [Unlocking LLM Performance with EBPF: Optimizing Training and Inference Pipelines](#)

Incoming Event



NOVEMBER 12, 2024

SALT LAKE CITY, UTAH

#CiliumeBPFDay

Reference:

- [Incoming event: Cilium+eBPF Day North America November 12, 2024 Salt Lake City, Utah](#)

Jumpbox®

Learn more about eBPF

- [eBPF Everything You Need to Know in 5 Minutes](#)
- [Getting Started with eBPF](#)
- [Beginner's Guide to eBPF Programming](#)
- [Beginner's Guide to eBPF Programming with Go](#)
- [Run Fast! Catch Performance Regressions in eBPF with Rust](#)
- [What is eBPF? Brightboard Lesson](#)
- [Brendan Gregg • BPF Performance Tools](#)



ISOVALENT

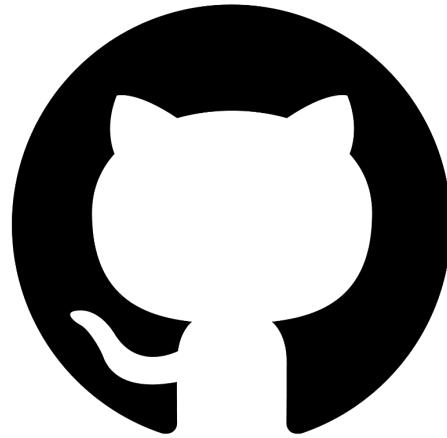
x



Jumpbox®

We are now partners

Jumpbox®



GitHub

<https://github.com/jumpbox-academy/learn-ebpf>

Jumpbox®

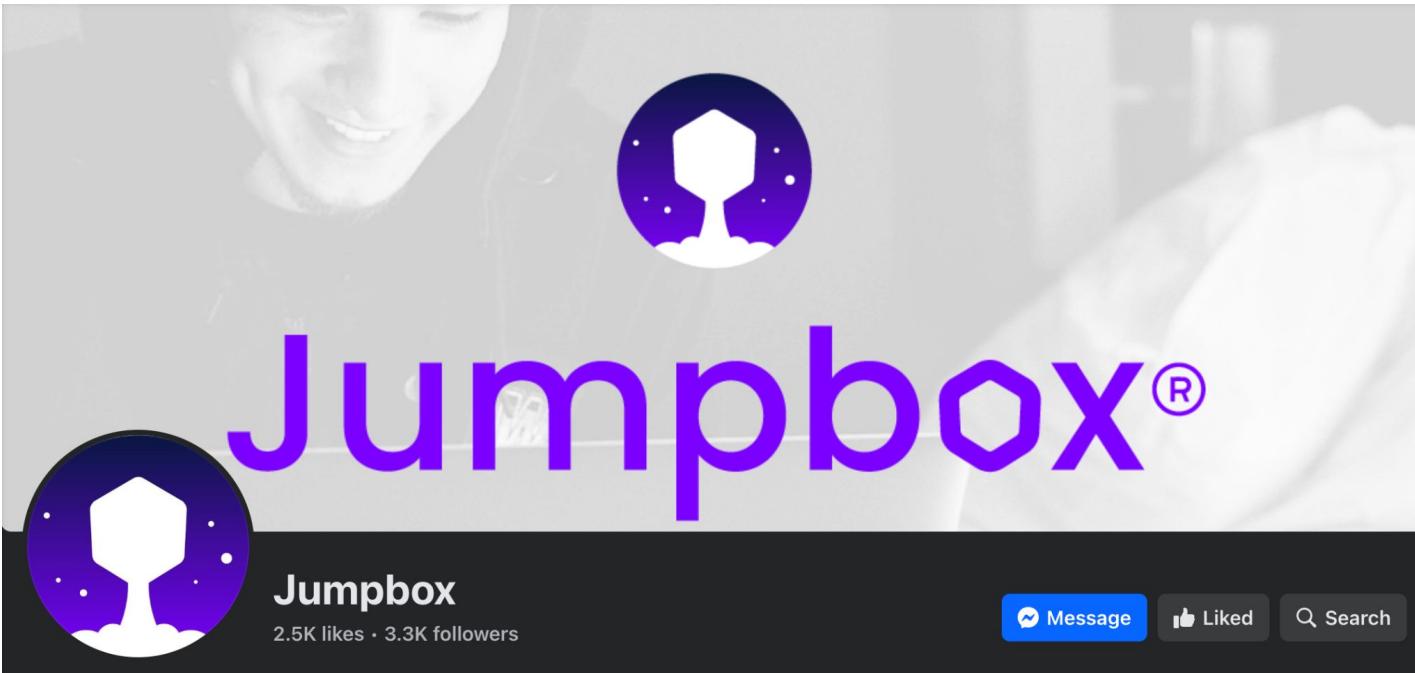


Jumpbox®

THE CLOUD CAMP

2024

facebook



<https://www.facebook.com/jumpbox.academy>

Jumpbox®



Jumpbox

@jumpbox.academy • 1.92K subscribers • 30 videos

More about this channel ...[more](#)

<https://www.youtube.com/@jumpbox.academy>

Jumpbox®

Contact Us



Jumpbox



@jumpbox



admin@jumpbox.co



063-245-2168 (JoJo)

062-796-1559 (Beau)



Jumpbox®

"เราเชื่อว่า การเรียนรู้ทำให้ชีวิตคุณดีขึ้น"

-Jumpbox Team-

Jumpbox®



Jumpbox®