

Kubernetes Security with *e*BPF

Mongkol Thongkraikaew

Head of Platform Engineering

Ascend Money



About Me



Mongkol Thongkraikaew
Head of Platform Engineering
Ascend Money

Open-Source Lovers, Cloud-Native Lovers, Tech Blogger



: Mongkol Thongkraikaew



Medium : @mongkol.ttm




: Mongkol Thongkraikaew

Agenda

- 🟡 **Kubernetes is Just Linux**
- 🟡 **Security Practices for Kubernetes**
- 🟡 **Overview of eBPF**
- 🟡 **Using eBPF to Secure Linux Systems**
- 🟡 **Kubernetes Security with eBPF**



Game Changer



CLOUD NATIVE
COMPUTING FOUNDATION

Game Changer



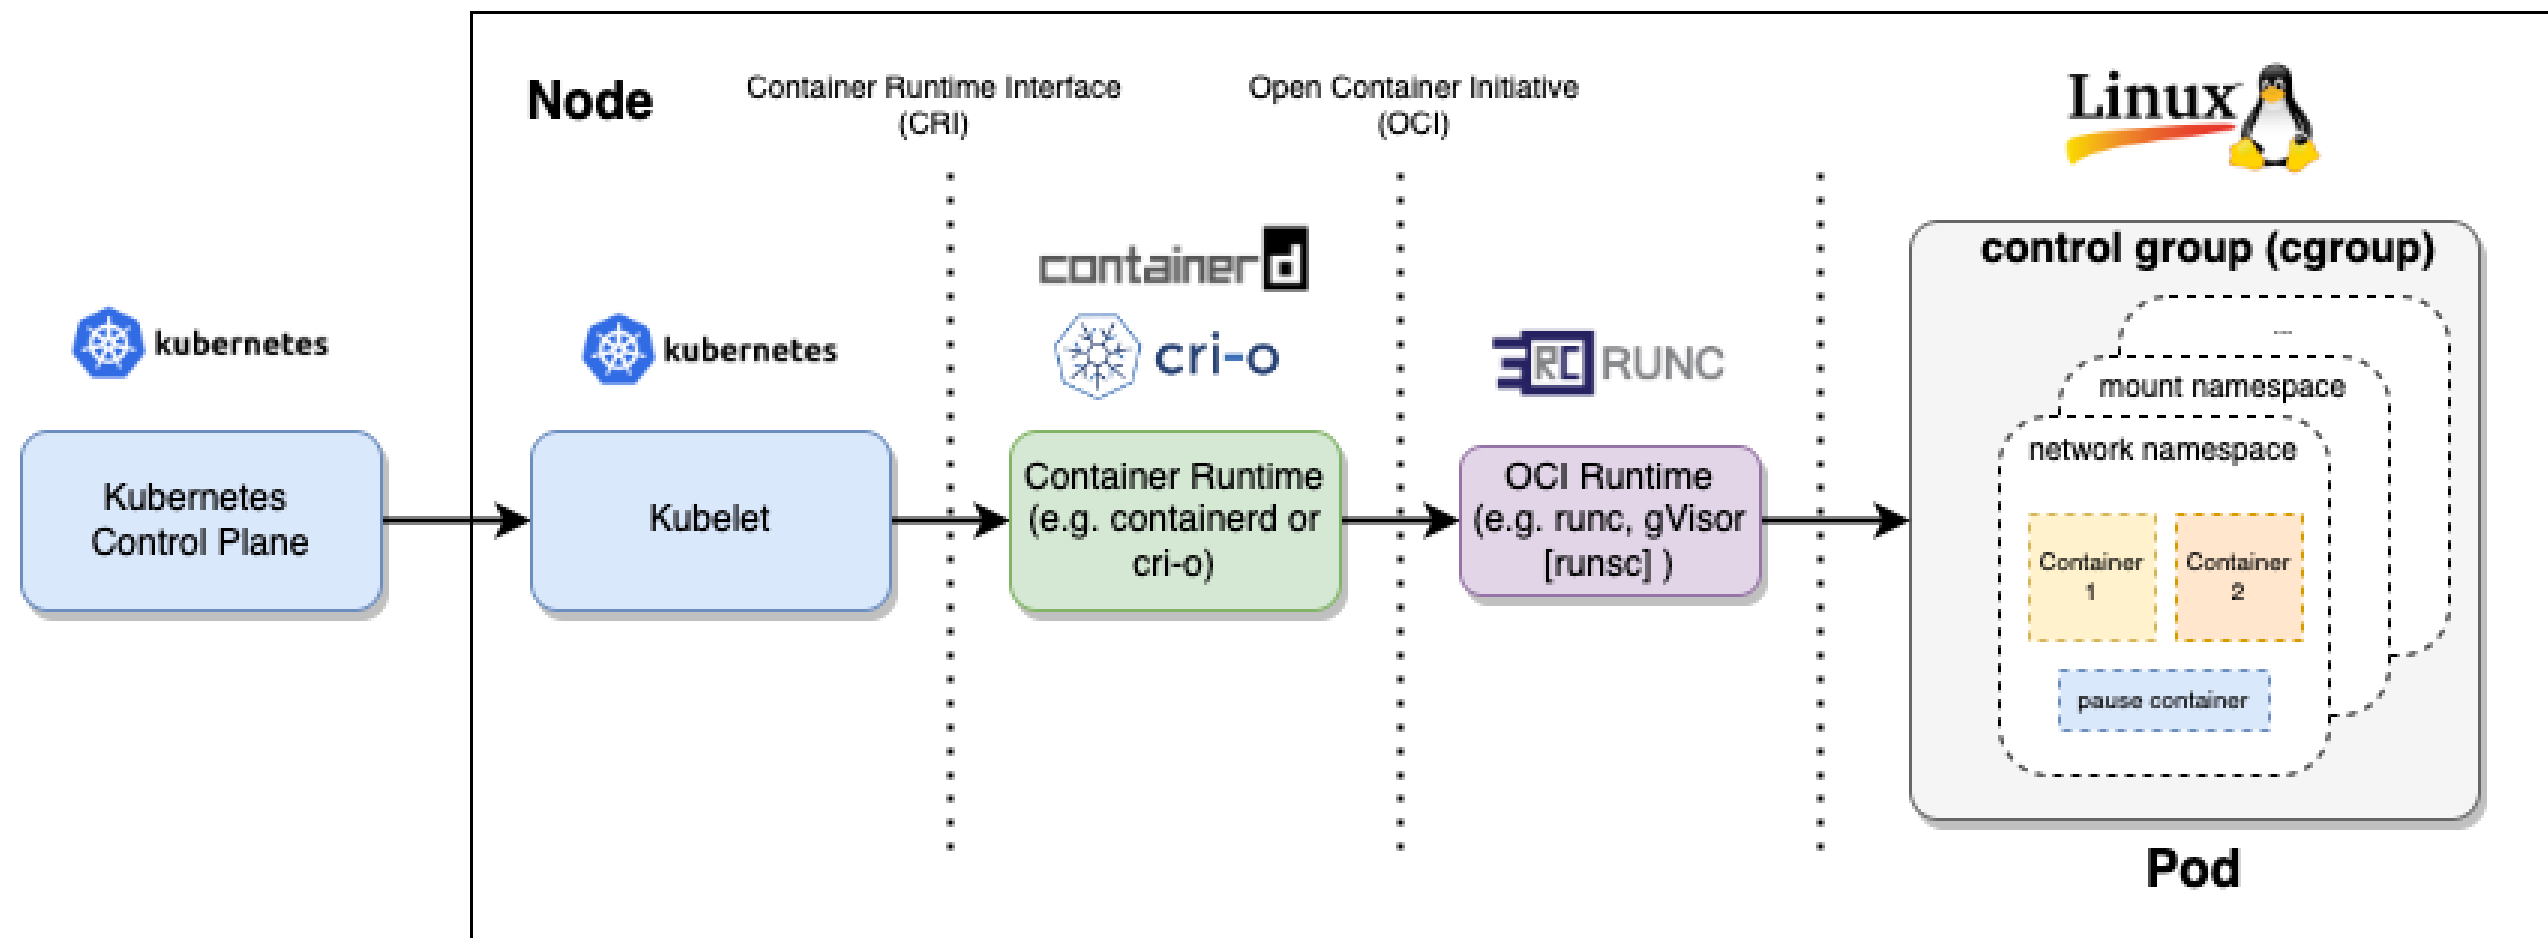


POD

Container Engine



Behind the scenes of Container



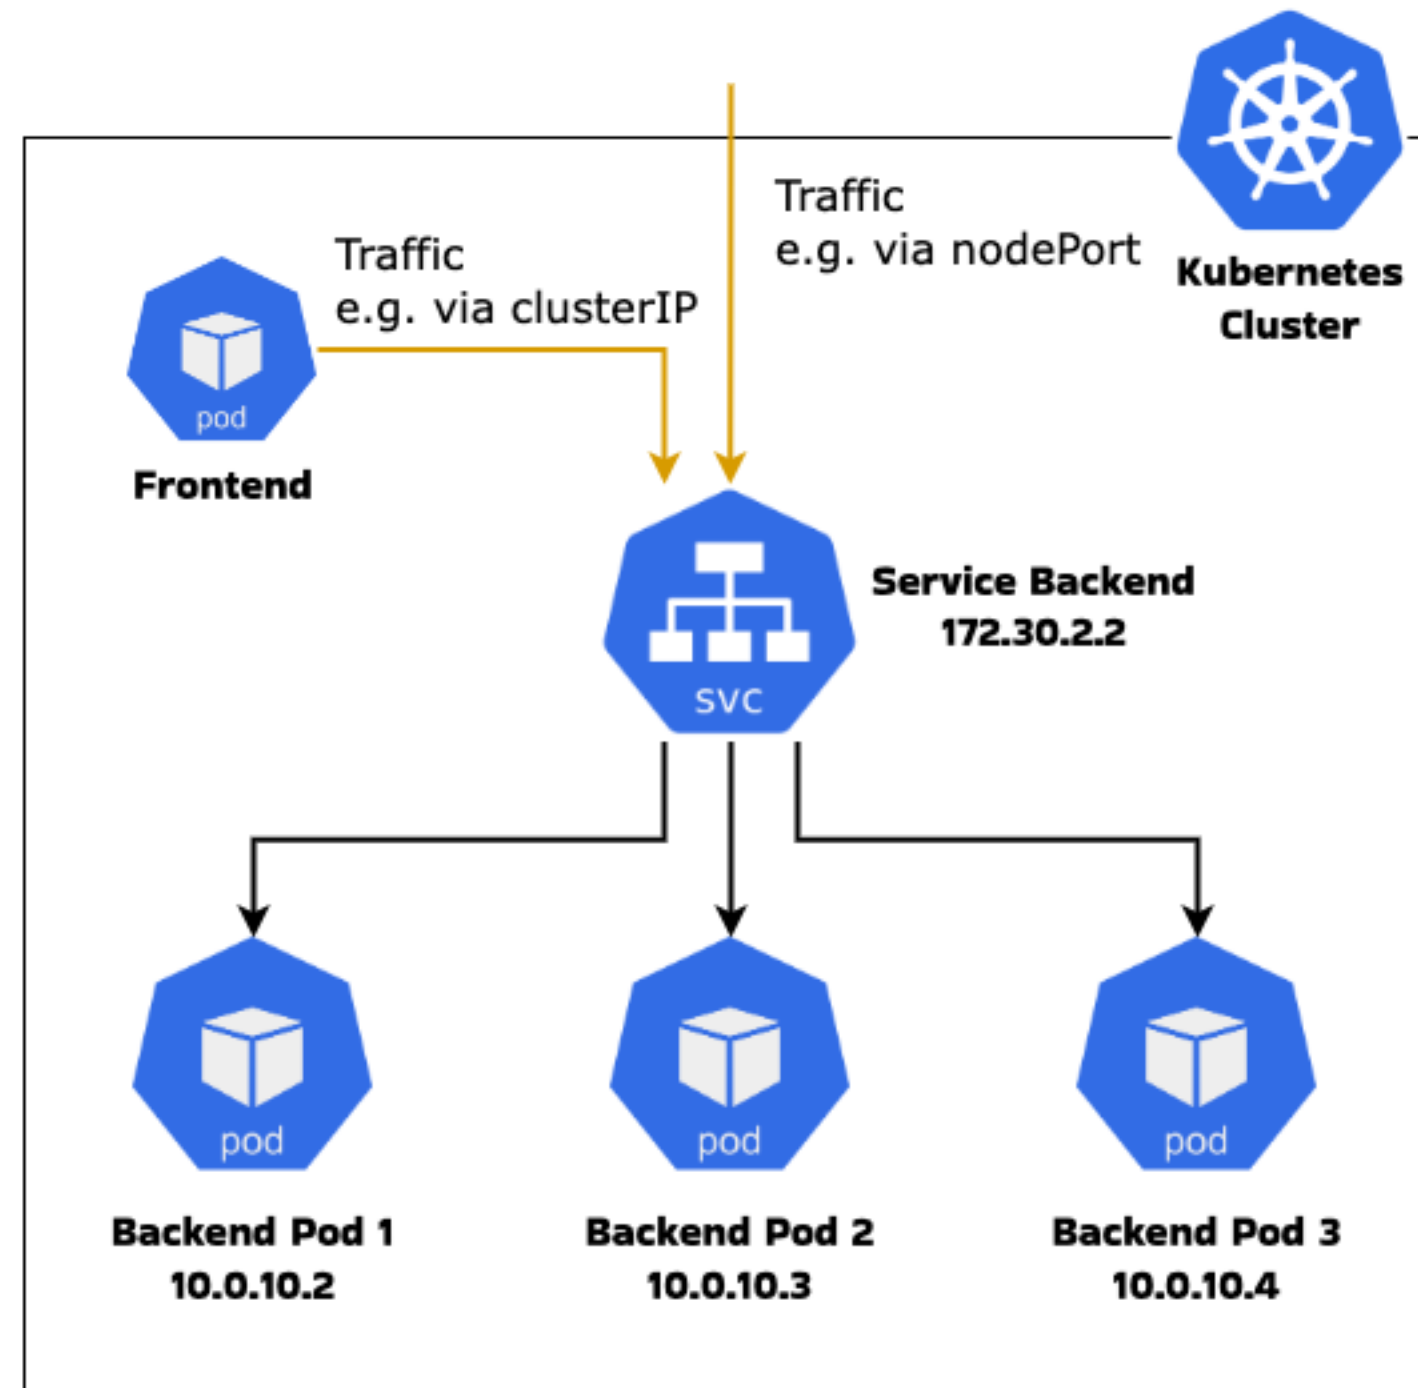
Container built from two main Linux kernels:

Linux Namespaces : The main feature that ensures the process is **completely separate** from other processes.

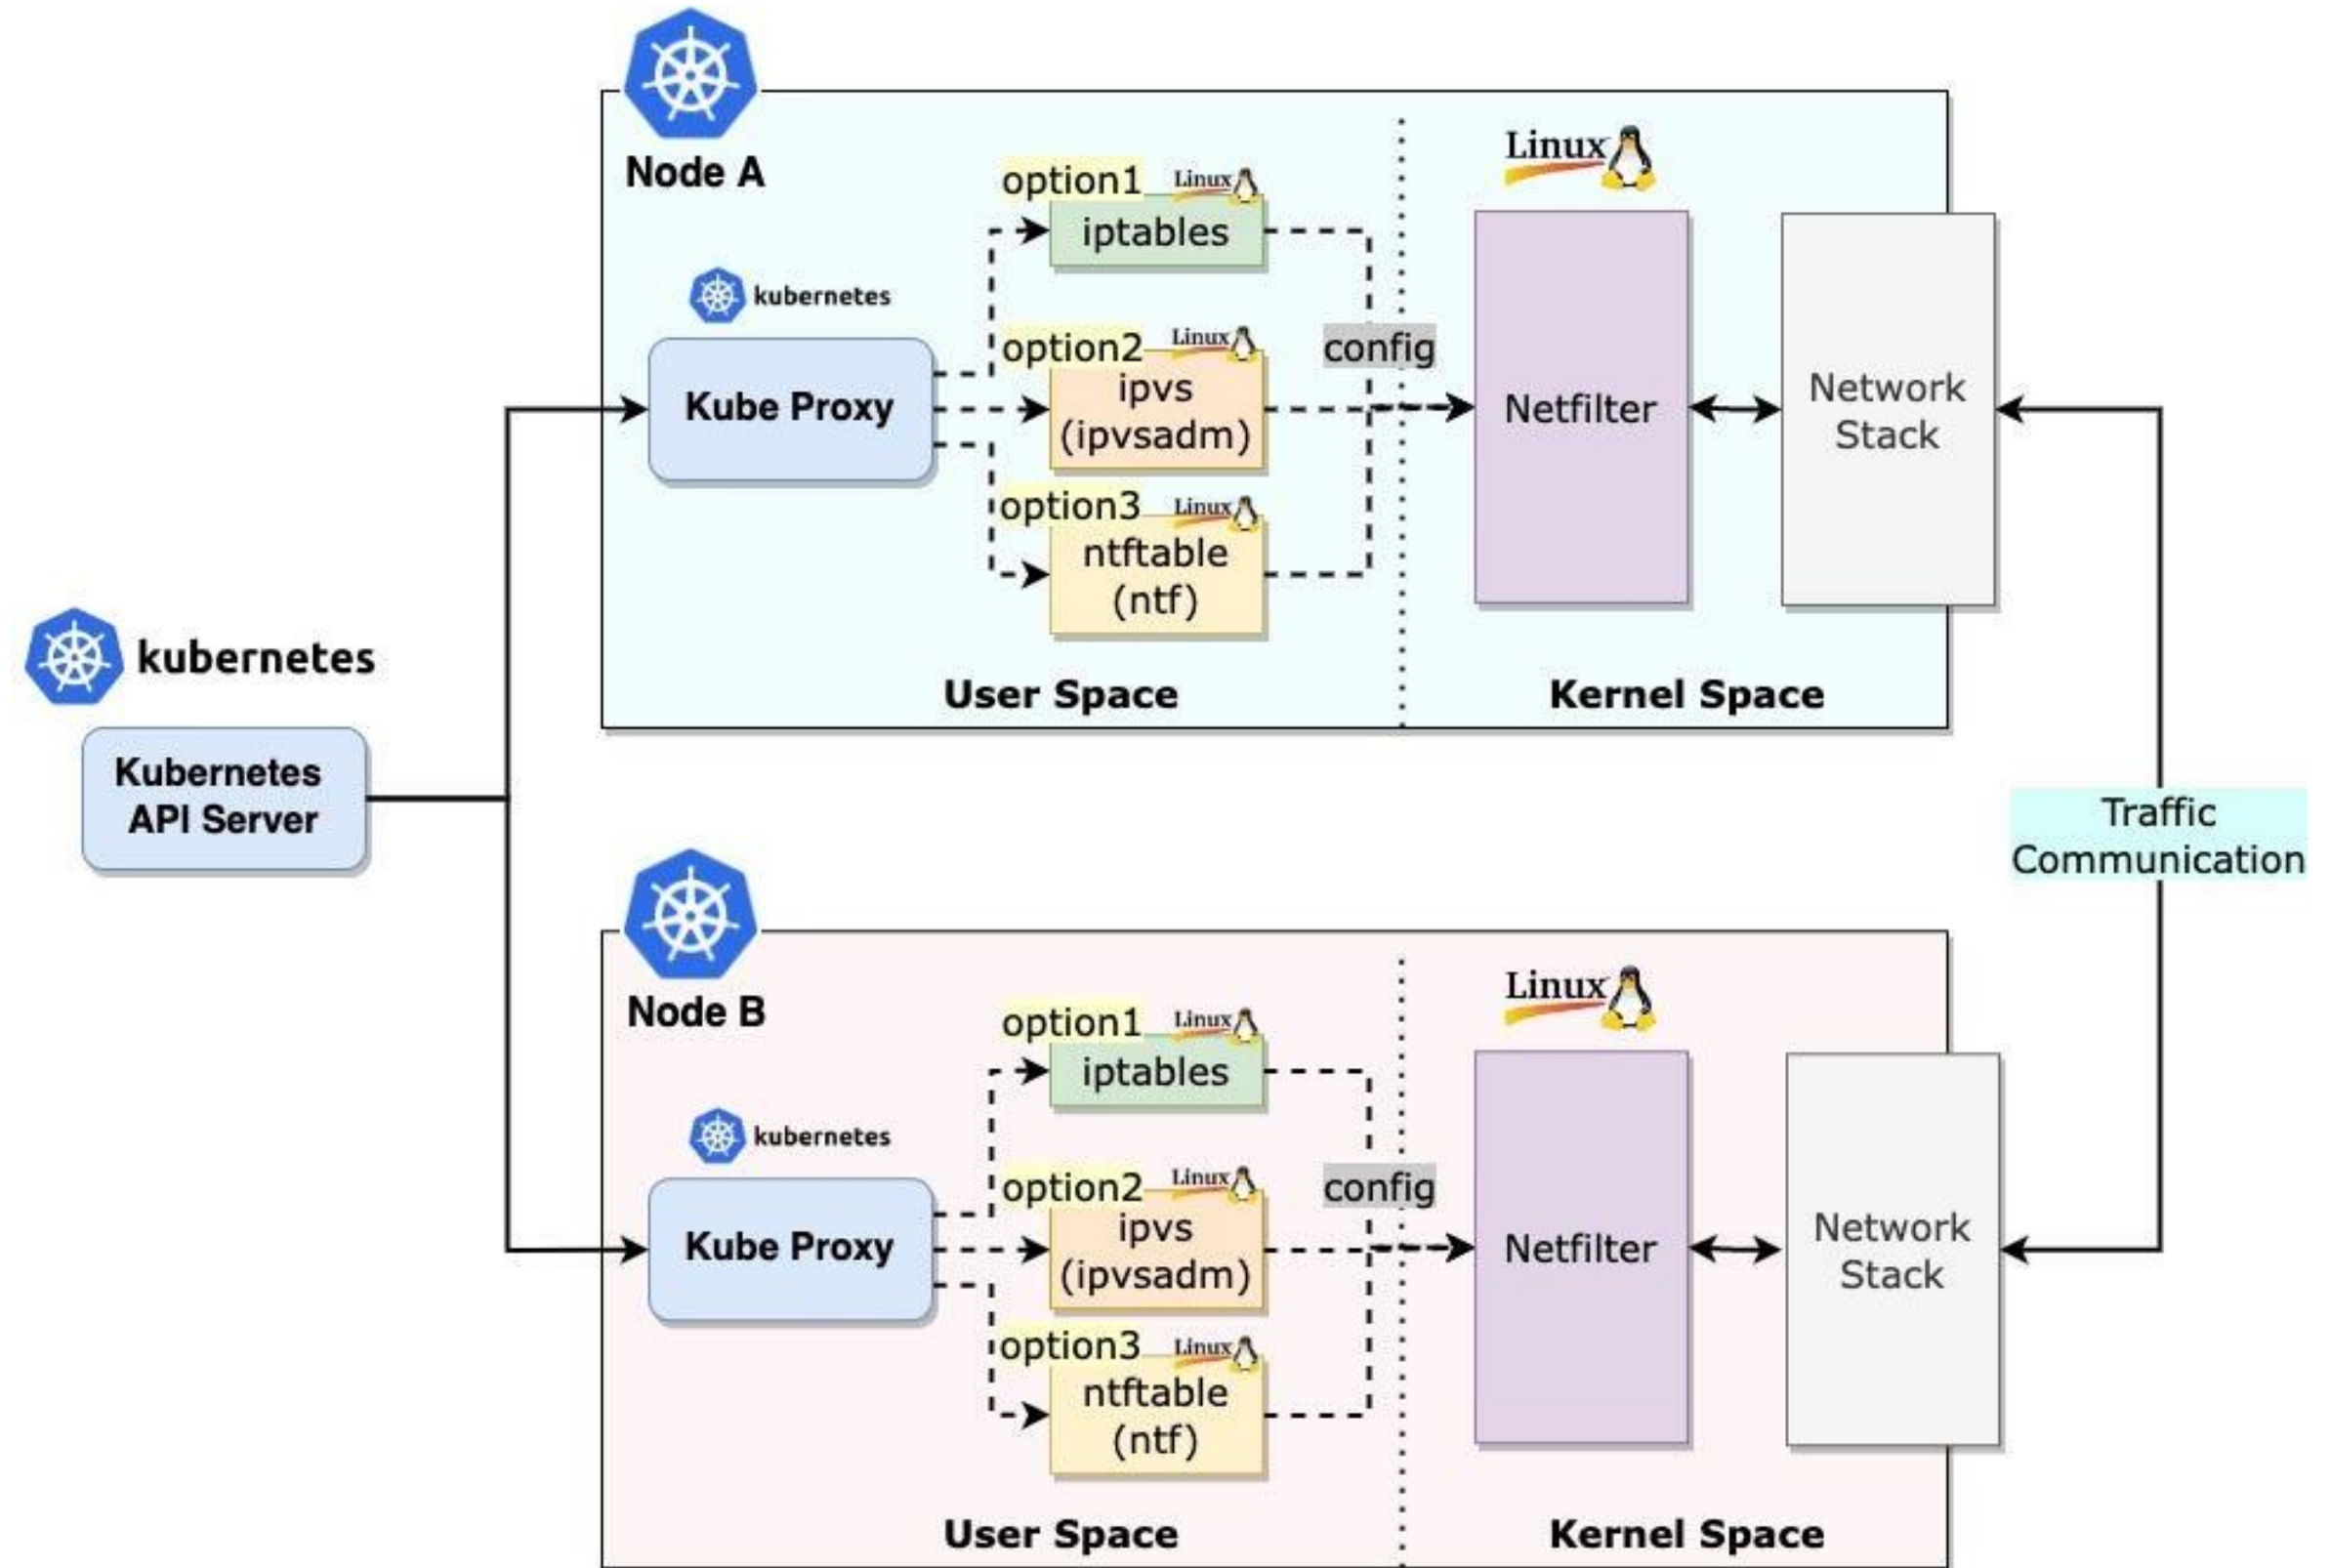

Control Group (Cgroups) : Allows you to **limit the system resources** (CPU, memory, disk I/O, network, etc) of a process.

“Container are just collection of processes”

Kubernetes Networking - Service





Kubernetes Networking - Service



Behind the scenes


Kubernetes Networking - Service



Behind the scenes


Kubernetes capabilities

Resource Limit



cgroup

Network Policy



iptables

Persistent Volume

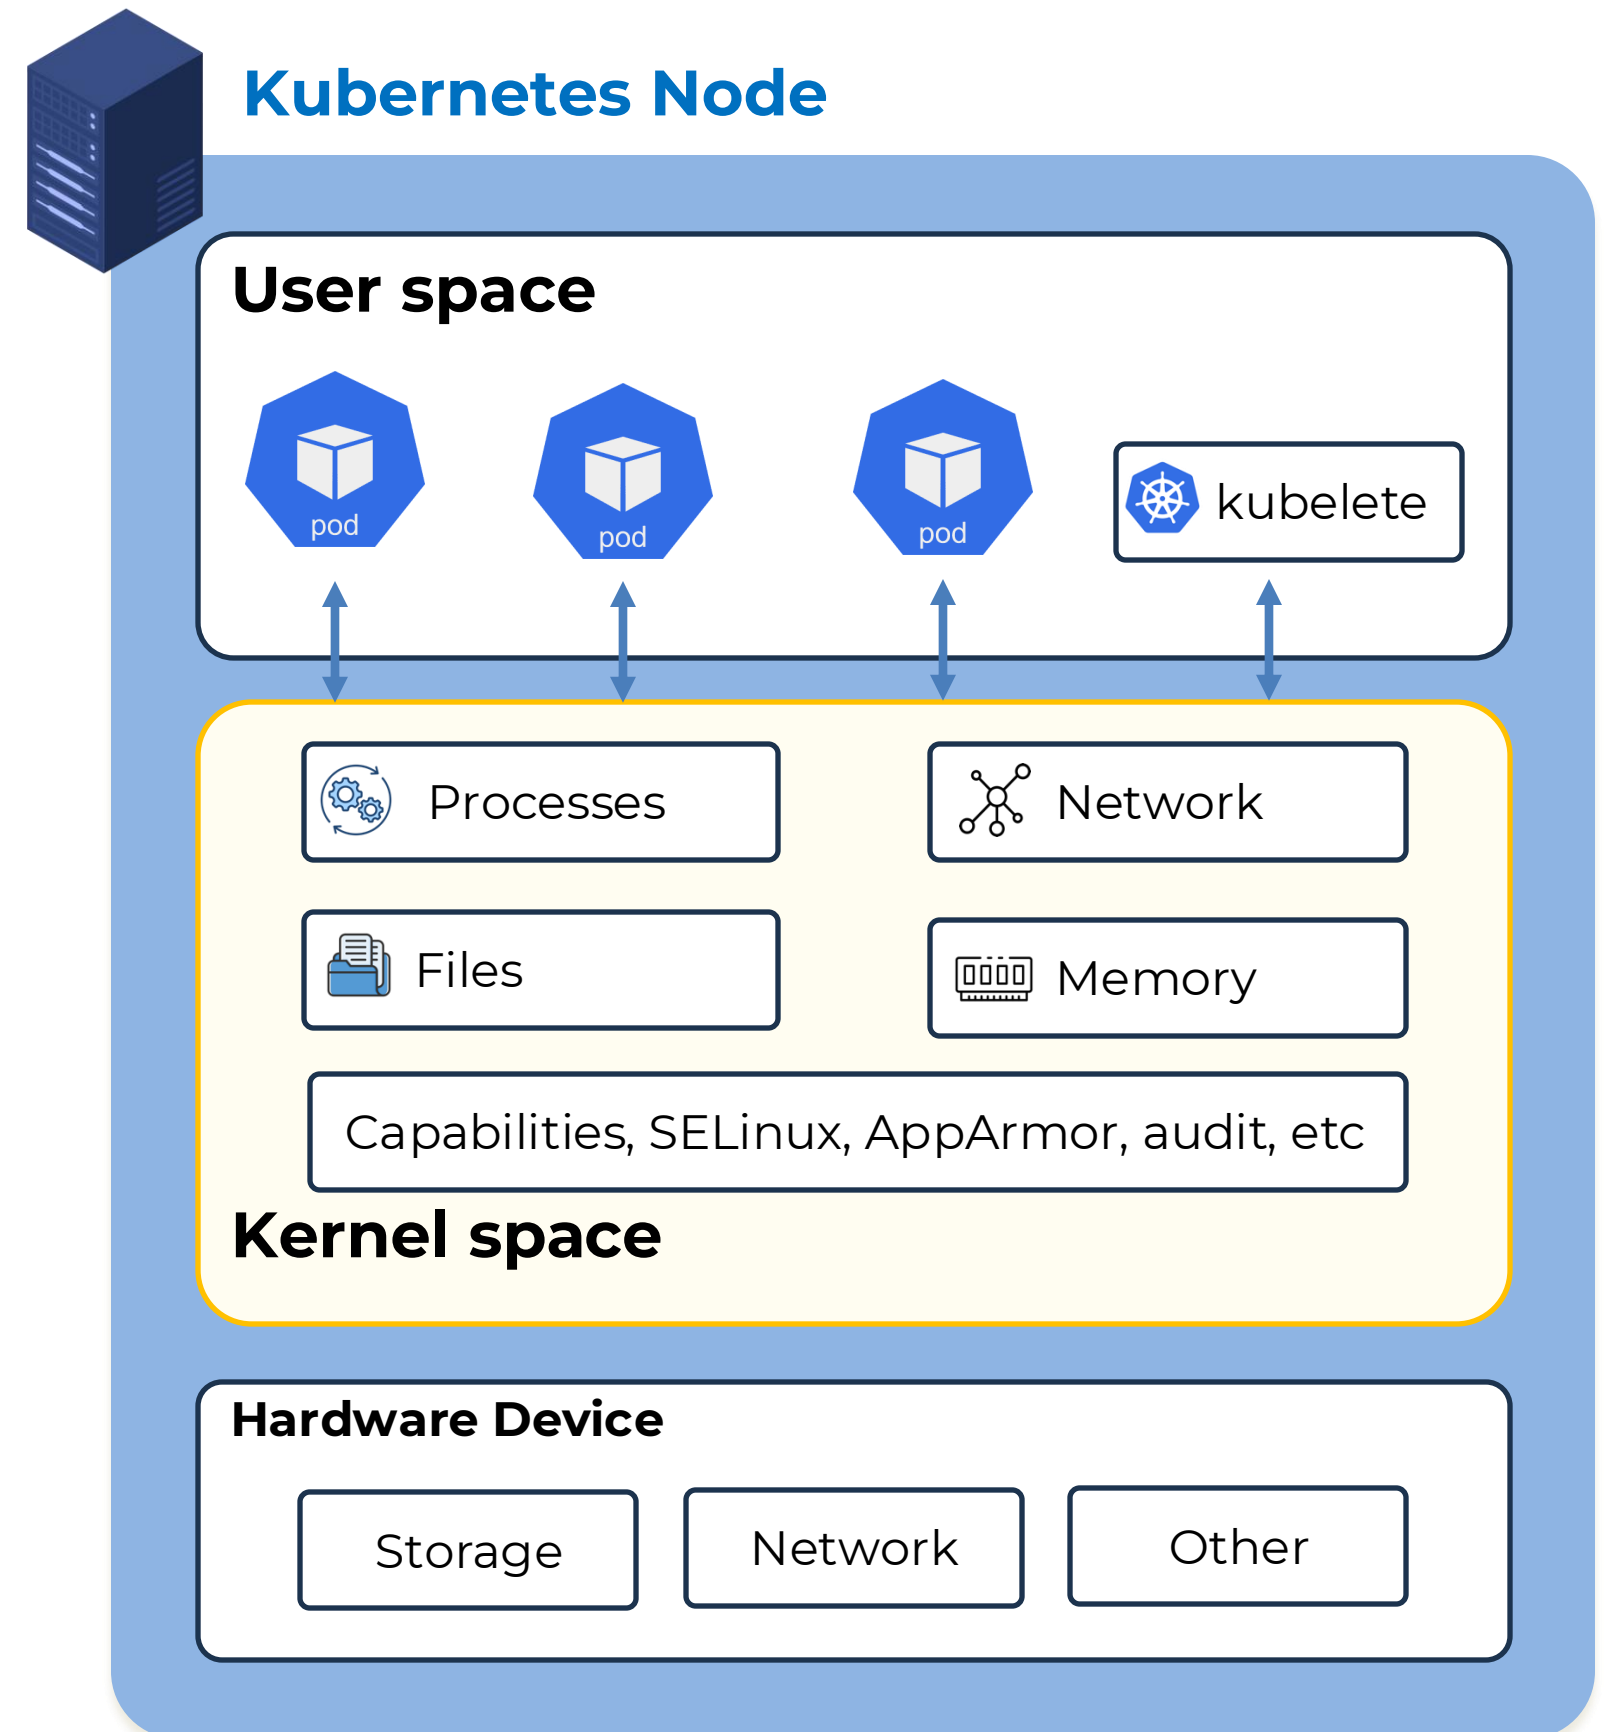


Virtual file system
(VFS)

Kubernetes is essentially a leveraged feature of ...




The Linux Kernel



Read More




●● Medium



Security Practice for Kubernetes




Application



Container



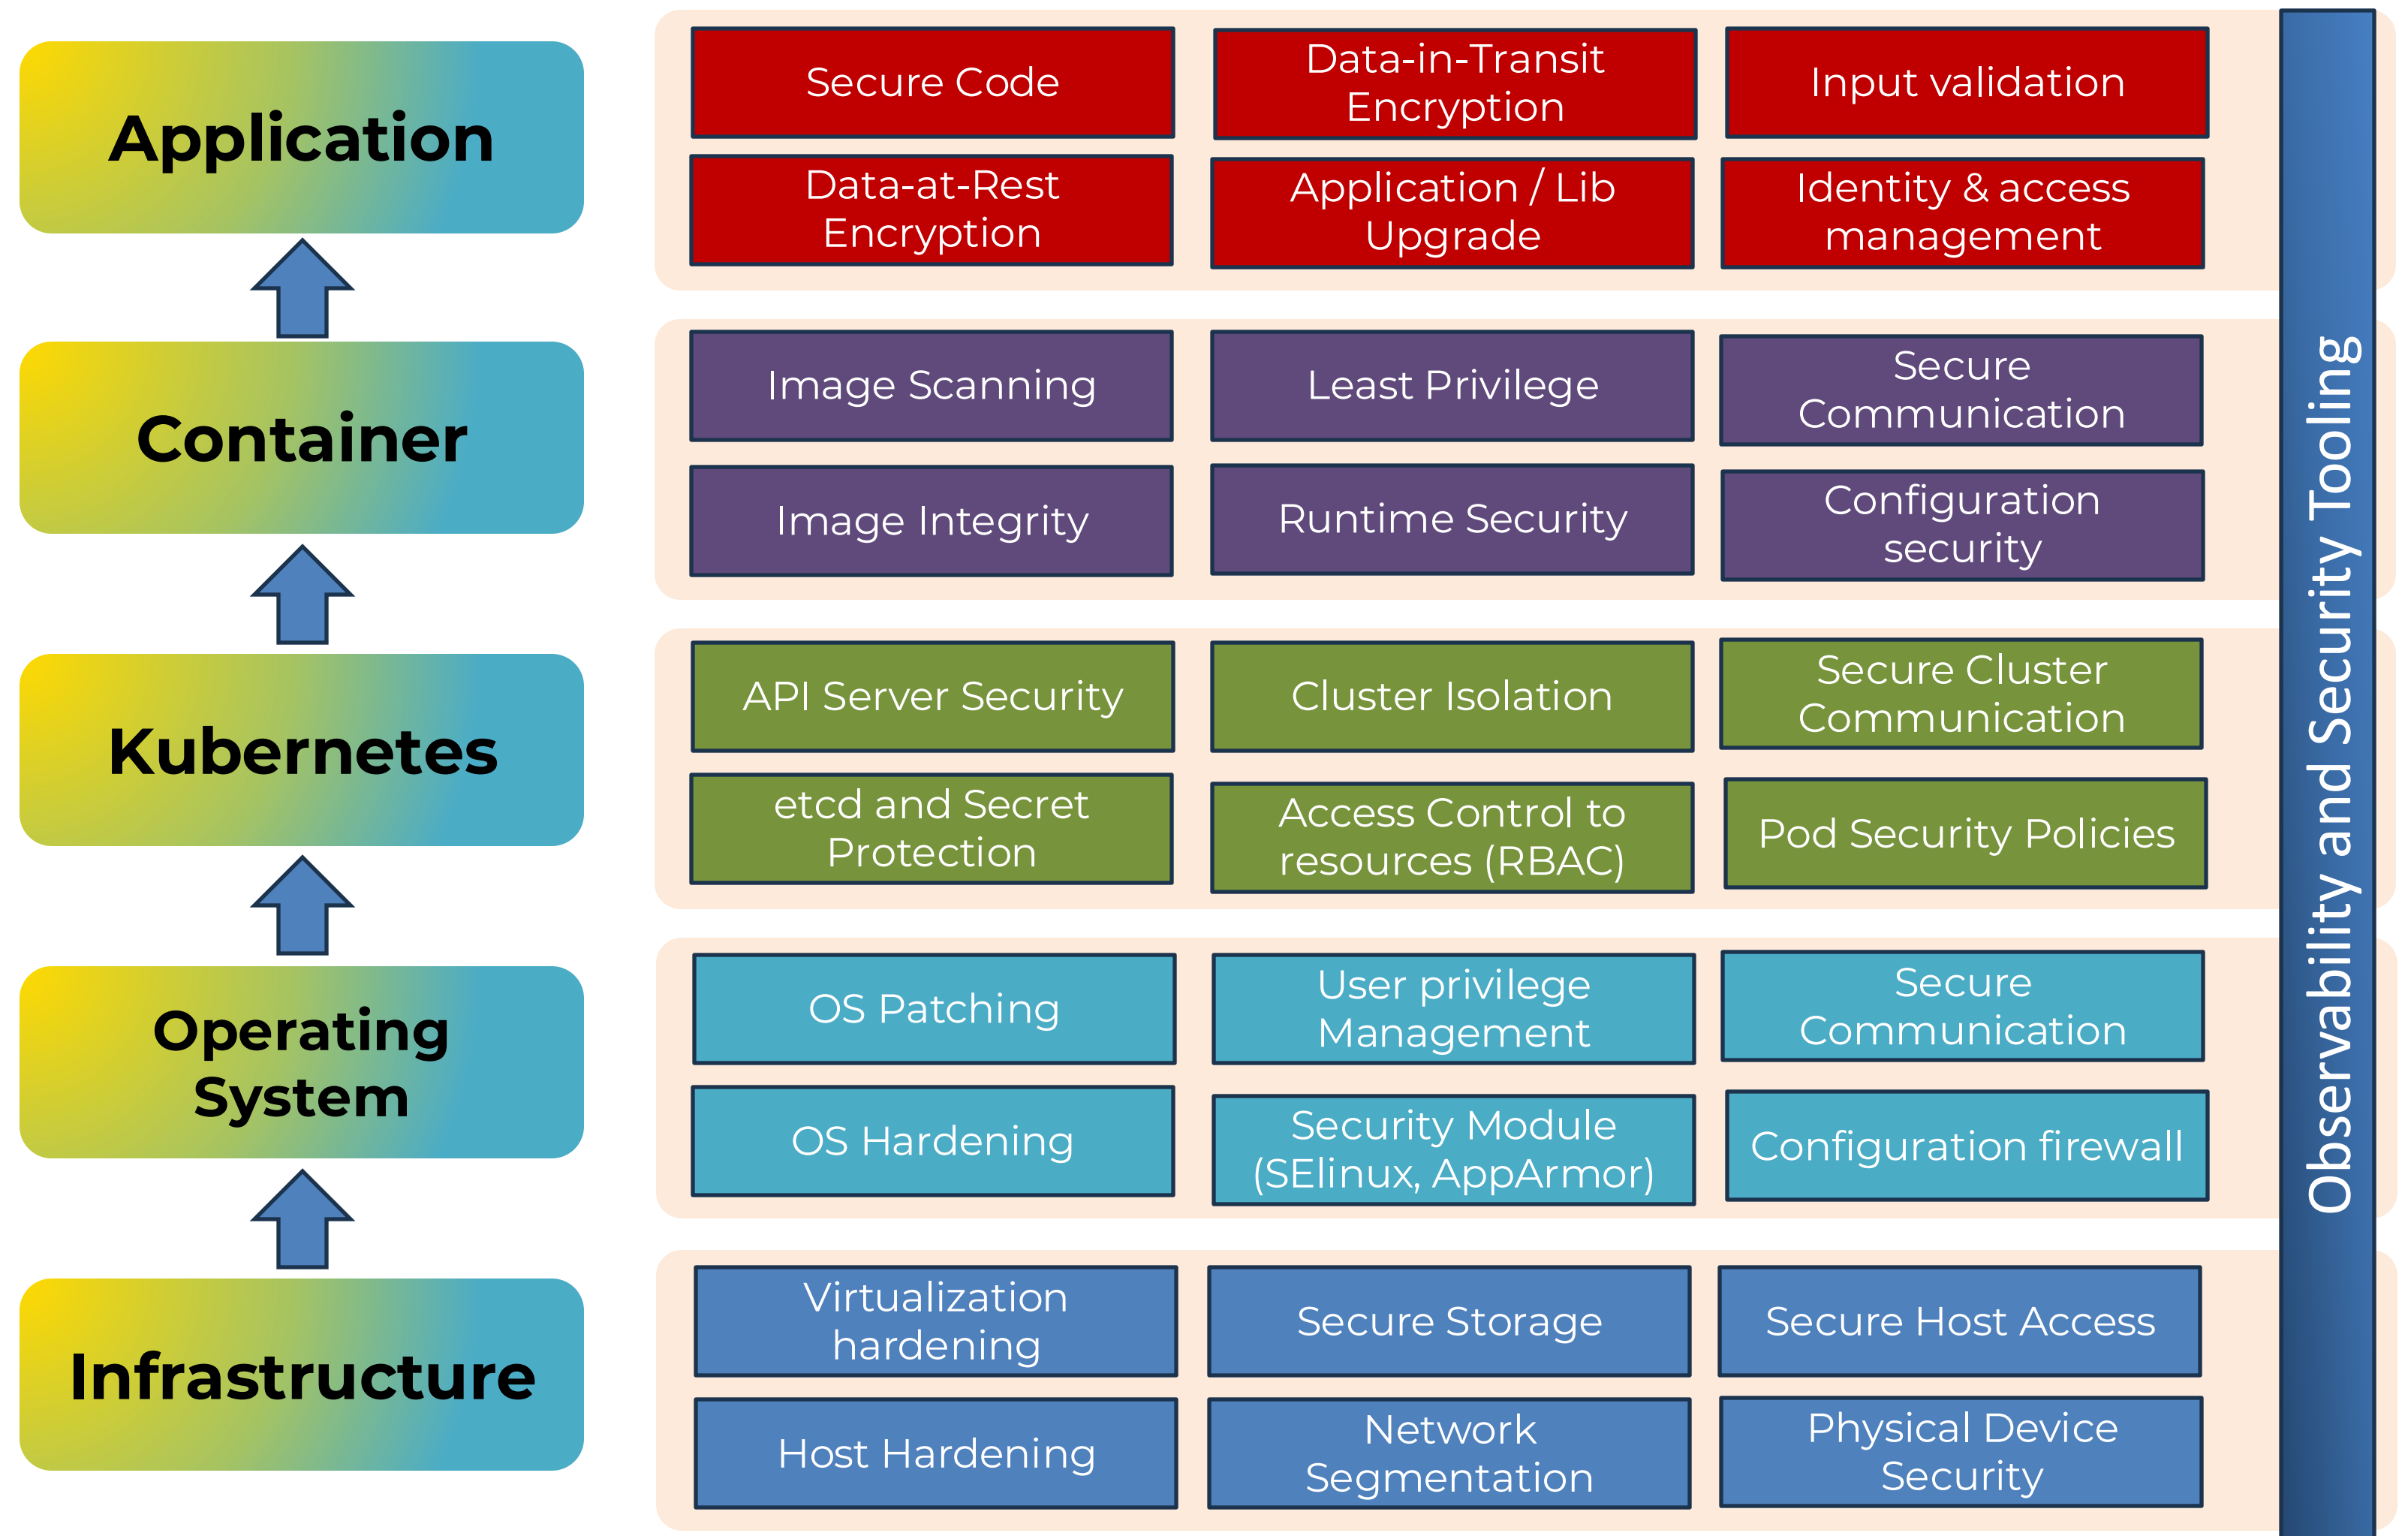
Kubernetes



**Operating
System**



Infrastructure



Recommend tools for security validation in Kubernetes



Tool that validates telco application's adherence to cloud native principles and best practices.

<https://github.com/cnti-testcatalog/testsuite>

- [Category: Security Tests](#)

[\[Container socket mounts\]](#) | [\[Privileged Containers\]](#) | [\[External IPs\]](#) | [\[SELinux Options\]](#) | [\[Sysctls\]](#) | [\[Privilege escalation\]](#) | [\[Symlink file system\]](#) | [\[Application credentials\]](#) | [\[Host network\]](#) | [\[Service account mapping\]](#) | [\[Ingress and Egress blocked\]](#) | [\[Insecure capabilities\]](#) | [\[Non-root containers\]](#) | [\[Host PID/IPC privileges\]](#) | [\[Linux hardening\]](#) | [\[CPU limits\]](#) | [\[Memory limits\]](#) | [\[Immutable File Systems\]](#) | [\[HostPath Mounts\]](#)



kube-bench is a tool that checks whether Kubernetes is deployed securely by running the checks documented in the CIS Kubernetes Benchmark.

<https://github.com/aquasecurity/kube-bench>



Kubescape is an open-source Kubernetes security platform that ensures comprehensive security throughout the development and deployment lifecycle, offering hardening, posture management, and runtime security for robust protection of Kubernetes environments.

<https://github.com/kubescape/kubescape>

Overview of eBPF

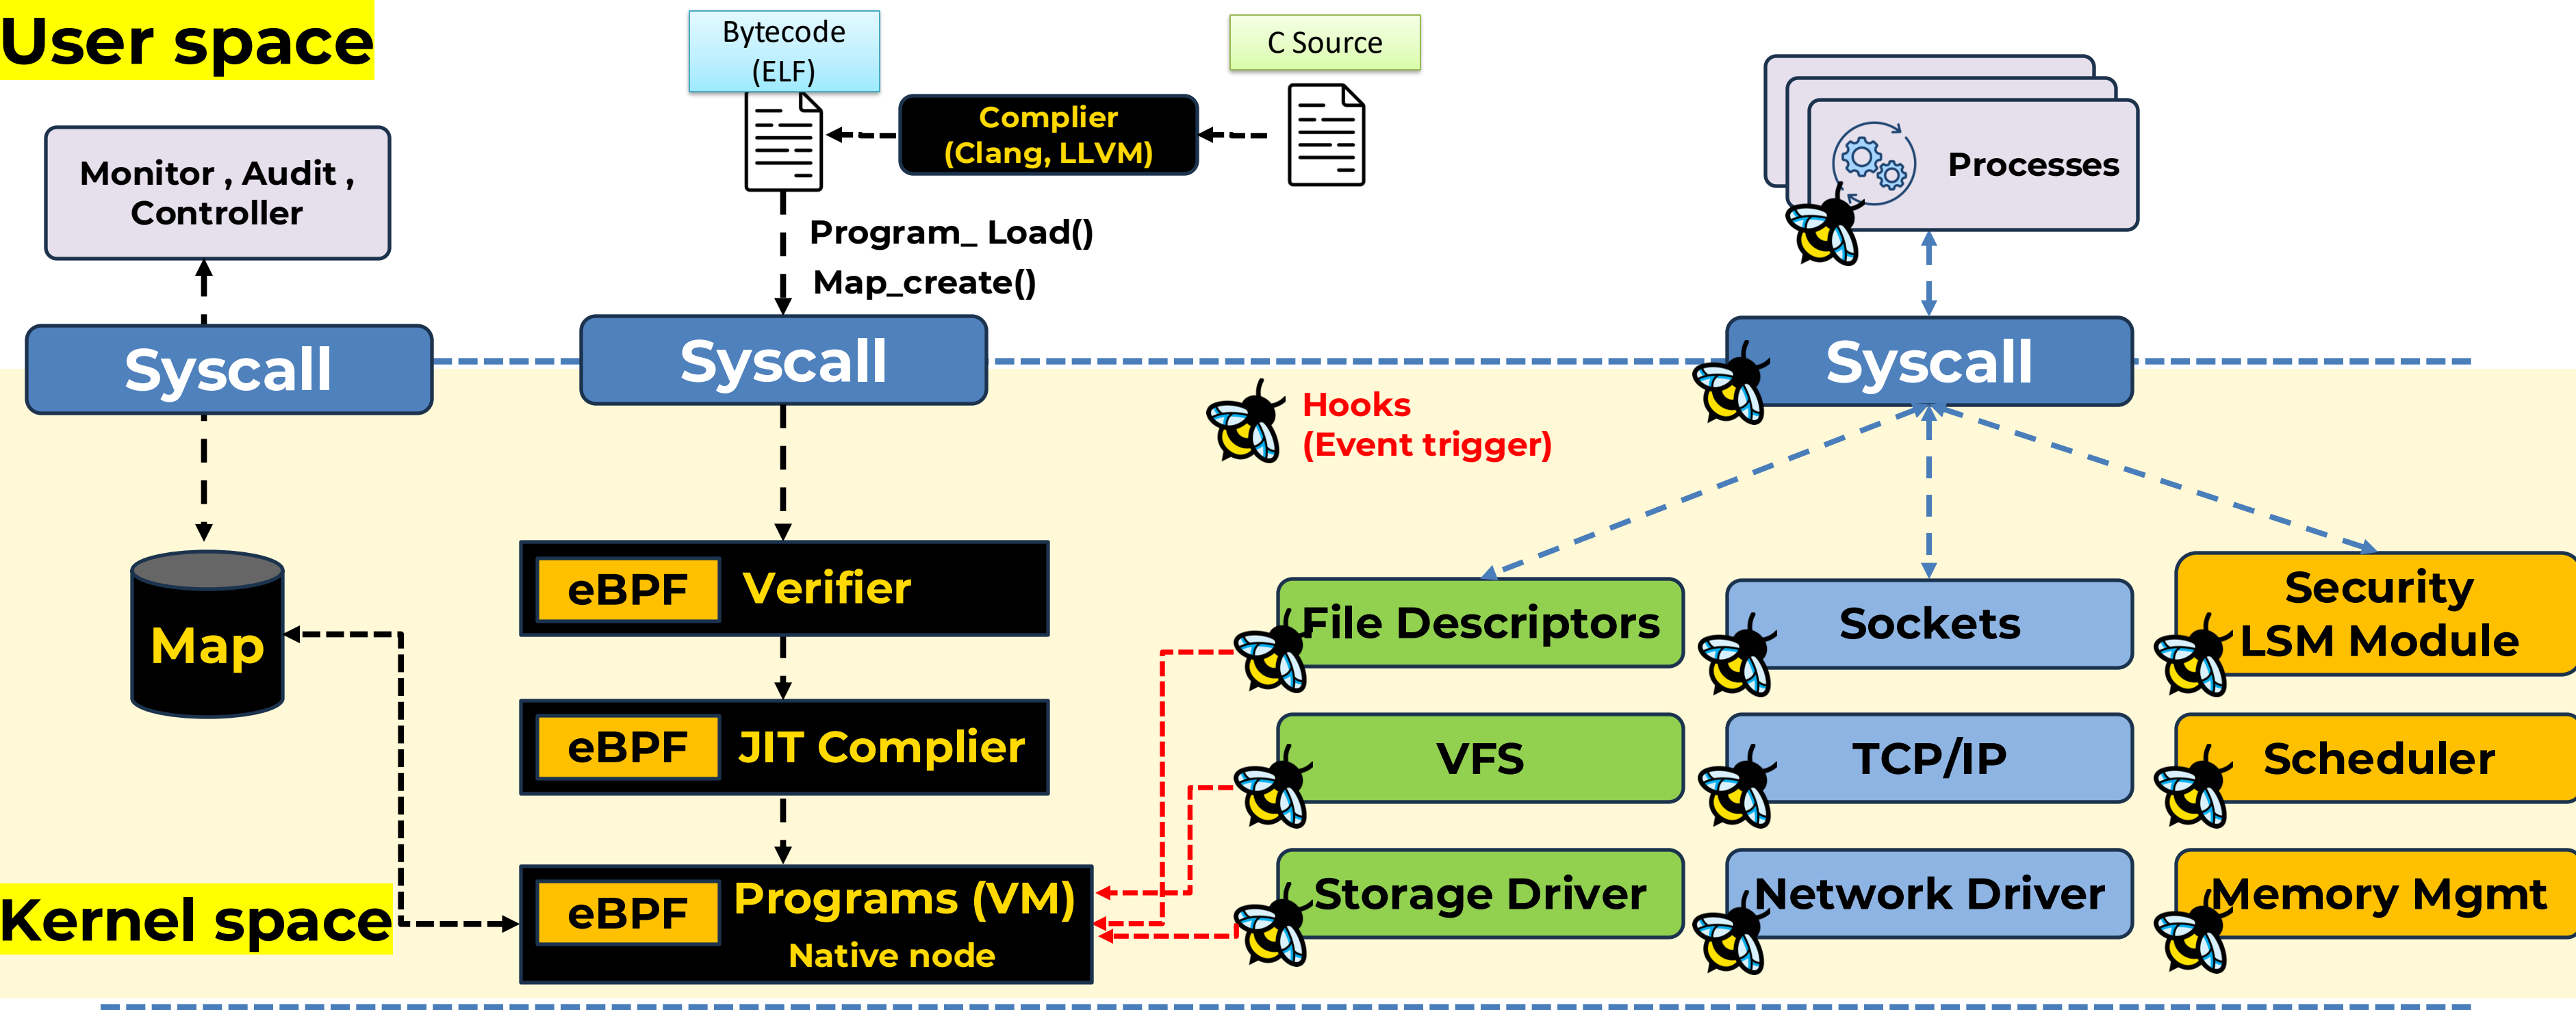
Accurate Definition

 **eBPF** is a **programming language** & **runtime** to extend **Operating Systems**

*eBPF refers to extended BPF
BPF - Berkeley Packet Filtering*

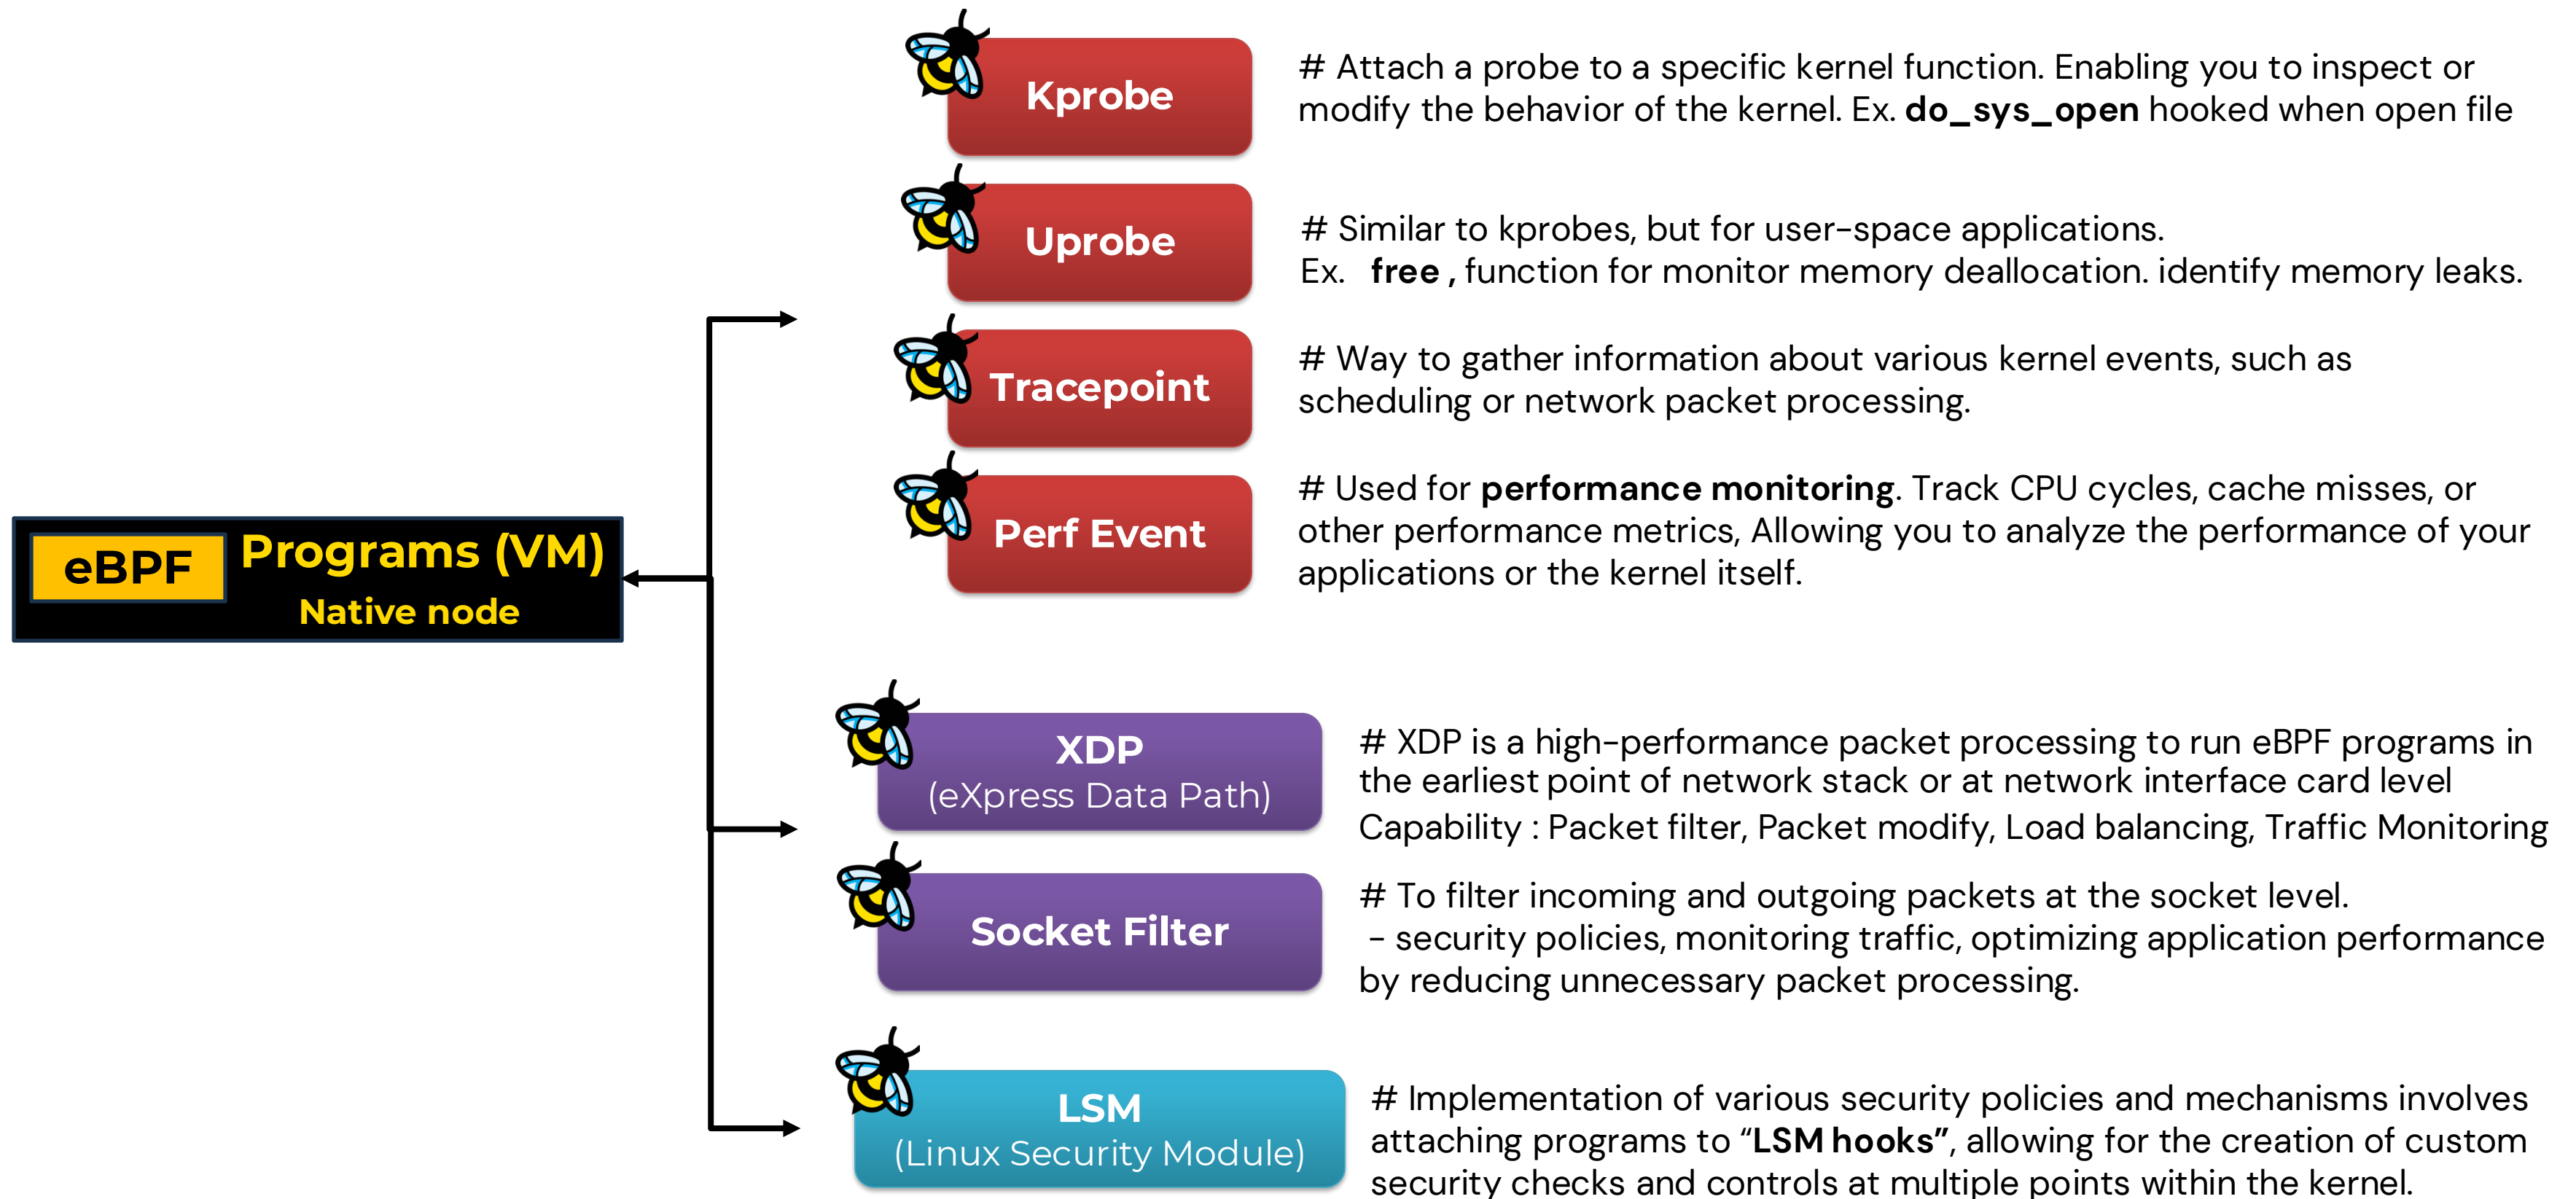
eBPF Diagram and Capability

User space




Hardware

Popular Attach point or Event Hook of eBPF




eBPF Use Cases

Security



Networking activities at the packet and socket levels. And enforce security policies by blocking or modifying certain system events based on predefined rules.

Observability




Trace network packets, system calls, function calls, and other events in real-time, enabling deep visibility into system behavior.

Networking



Firewalls, load balancers, and network monitoring tools. The low overhead and dynamic loading capabilities of eBPF

Tracing & Data Profiling



Profiling and tracing the runtime behavior of systems and applications

Using eBPF to Secure Linux Systems

Visibility and monitoring

Detect unusual behavior, identify potential threats, and analyze the attack surface.

Network Security

implementation of custom security traffic policies.

Container Security

Monitor and enforce security policies within containerized by attaching programs to **cgroup** hooks.



Intrusion Detection & Prevention

Monitor file system activity, restricting unauthorized access or modification of sensitive data.

Policy enforcement

Enforce access controls, regulate resource usage, and protect sensitive system components

Incident response & forensics

Collect detailed information about system activity and state during an incident

Kubernetes Security with eBPF

Security capabilities related to eBPF in kubernetes

Network Policy
Enforcement

Security Observability

Runtime Security
Monitoring

Intrusion Detection
Systems (IDP)

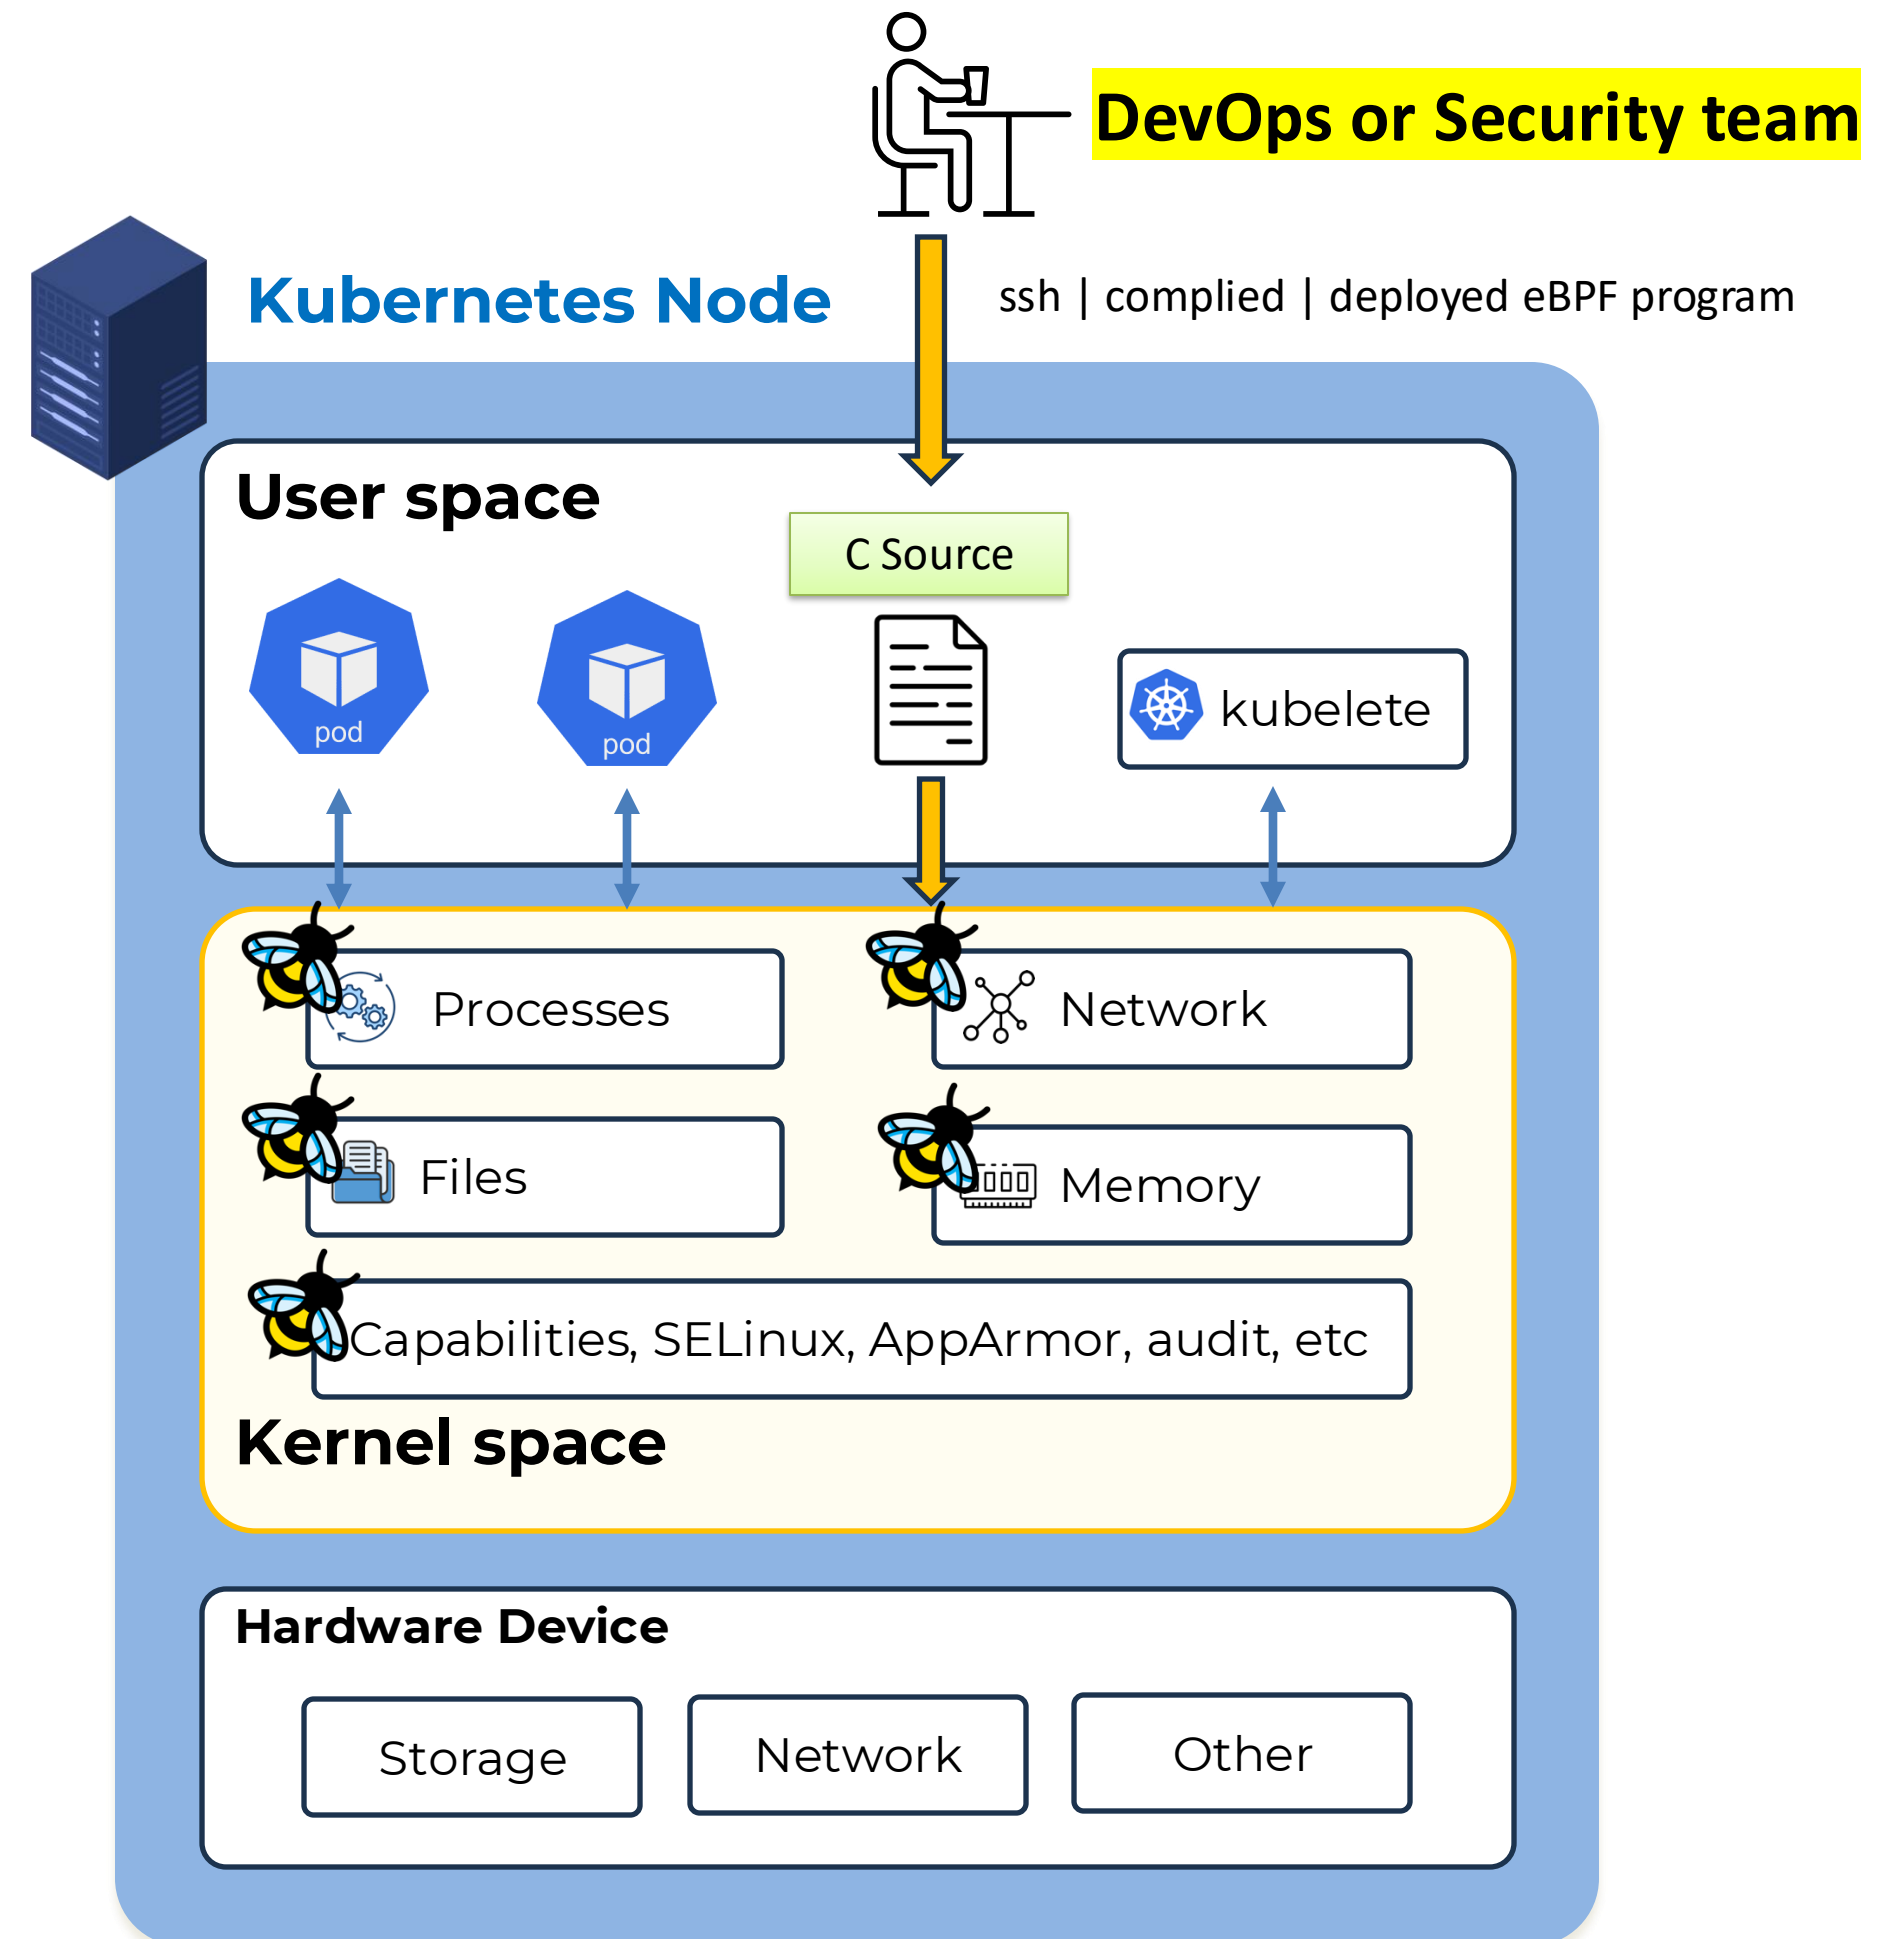
Cgroup & Namespace
Isolation

Policy Enforcement

Service Mesh



How to working with eBPF in Kubernetes





Kubernetes Node



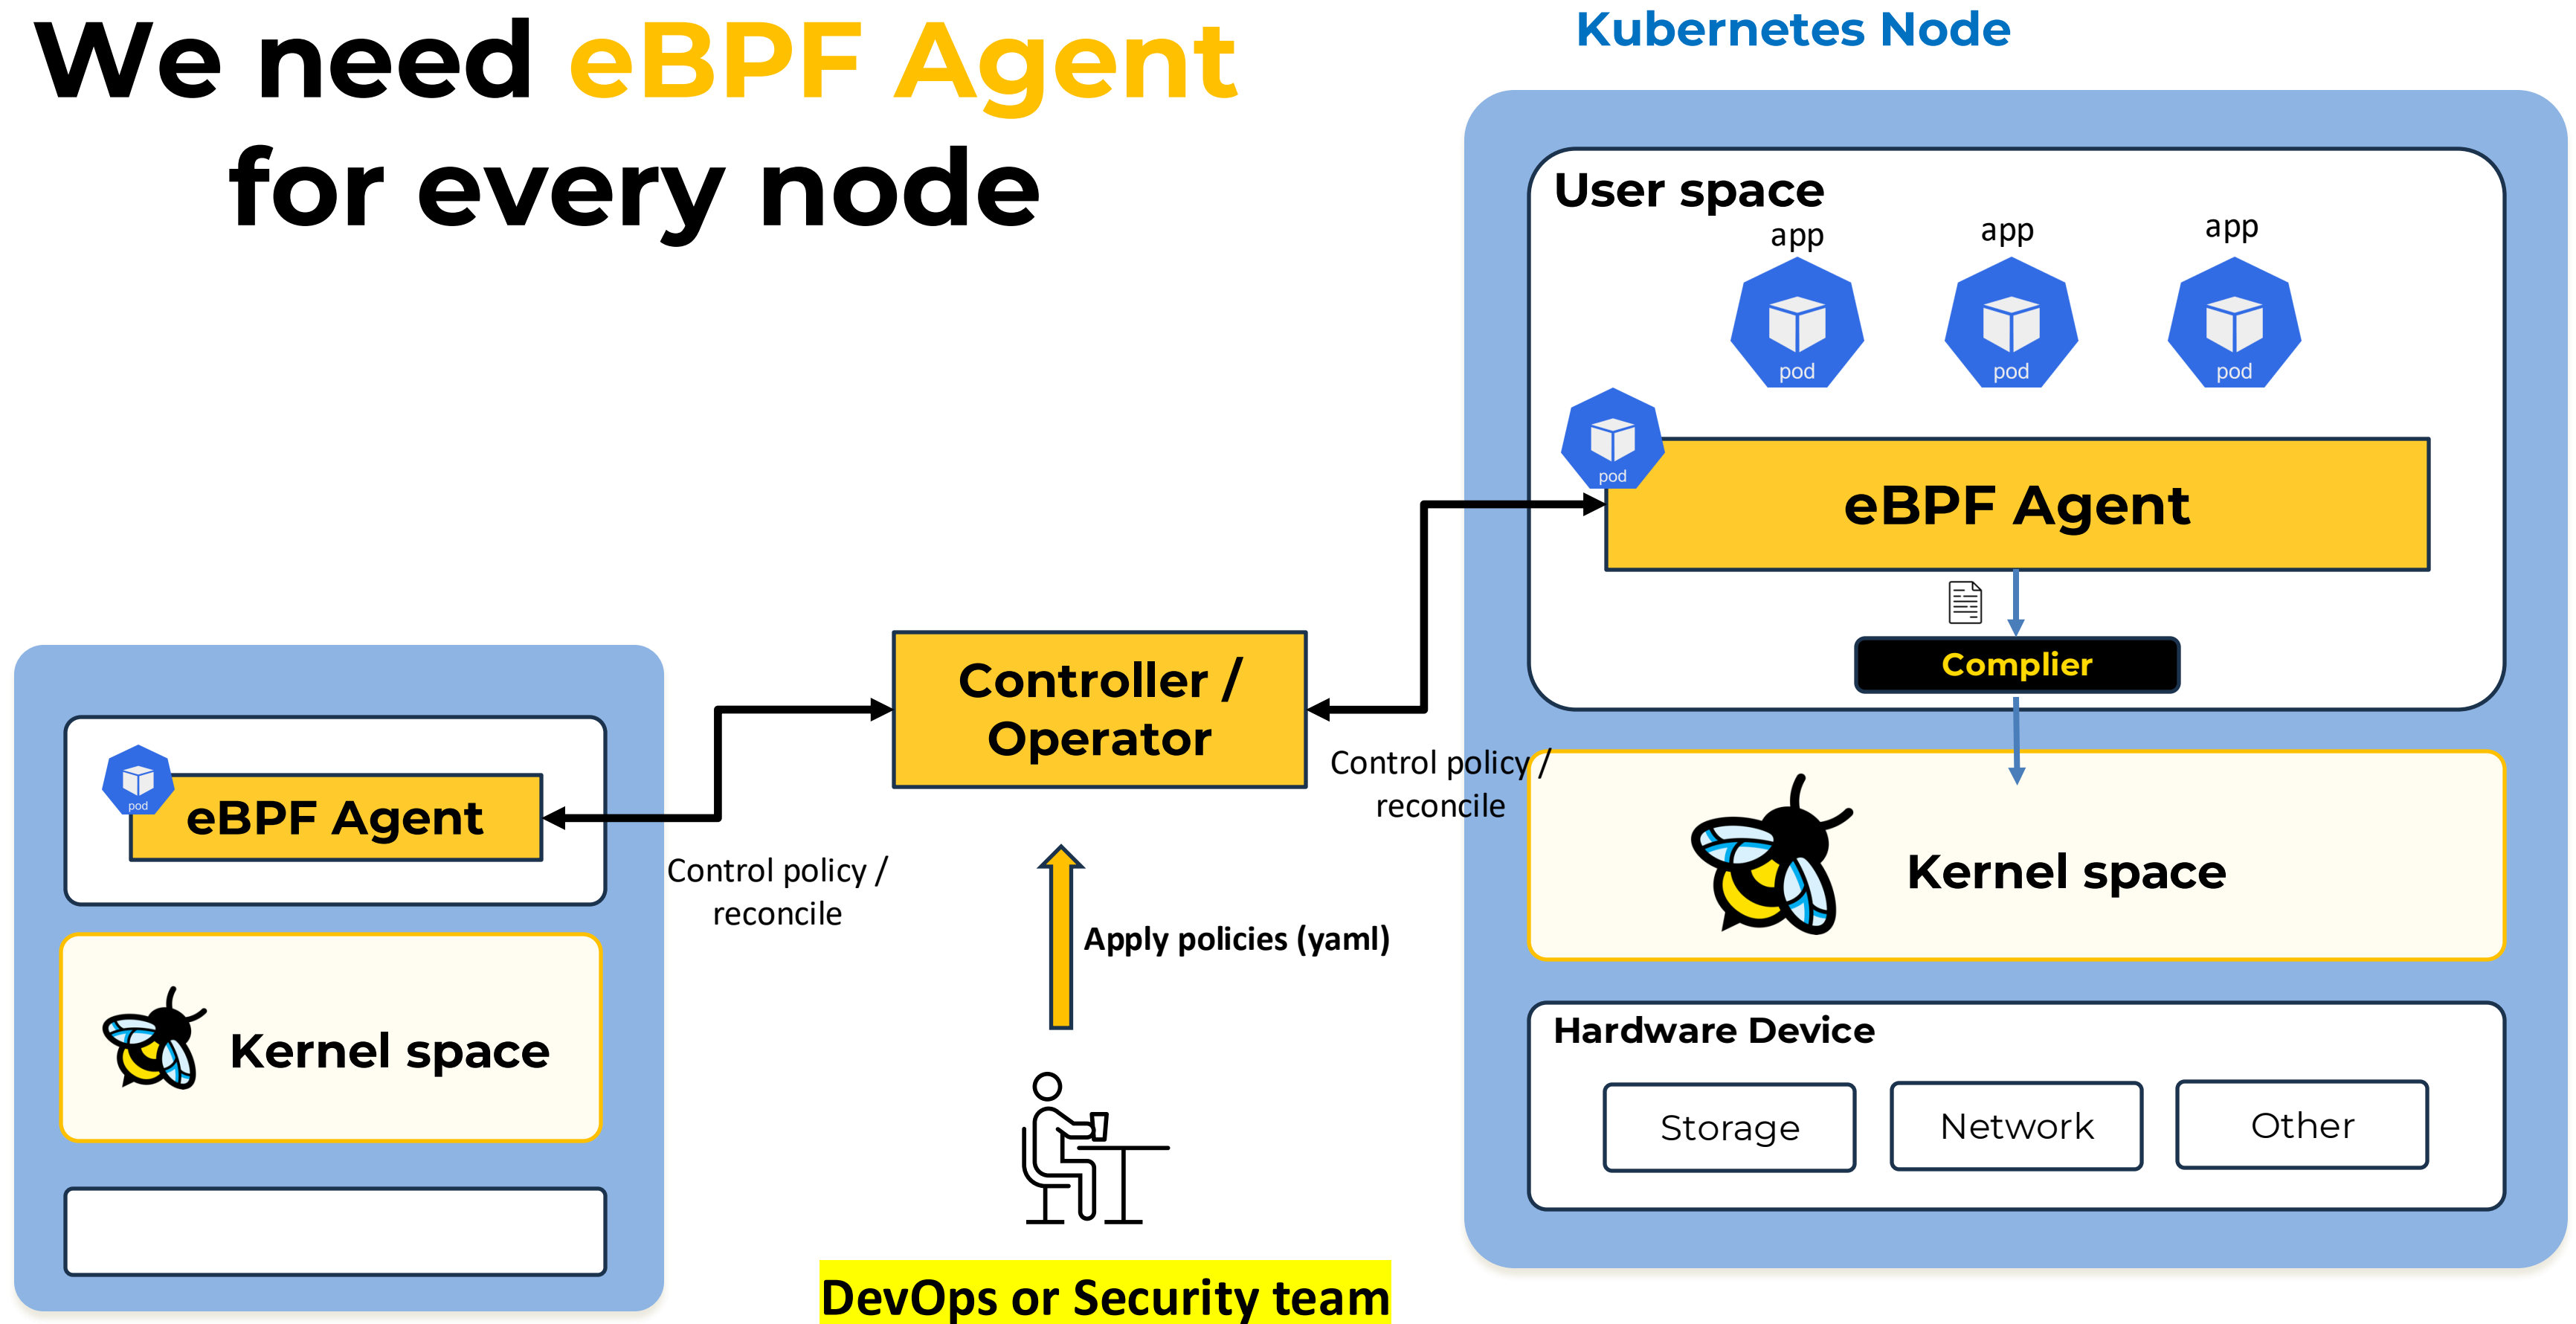
Just 1 Node

Kubernetes Node

If 10-100 Nodes?

A diagram illustrating a Kubernetes cluster. It features a central horizontal grey bar with the text "If 10-100 Nodes?" in red. Surrounding this bar are several blue squares of varying sizes, each representing a node. Each node has a dark blue server icon on top. The nodes are distributed across the image, with some being larger and others smaller, suggesting a heterogeneous or scalable node pool. The overall layout is symmetrical and clean, emphasizing the concept of a multi-node environment.

We need **eBPF Agent** for every node



Most Popular Tools for Security

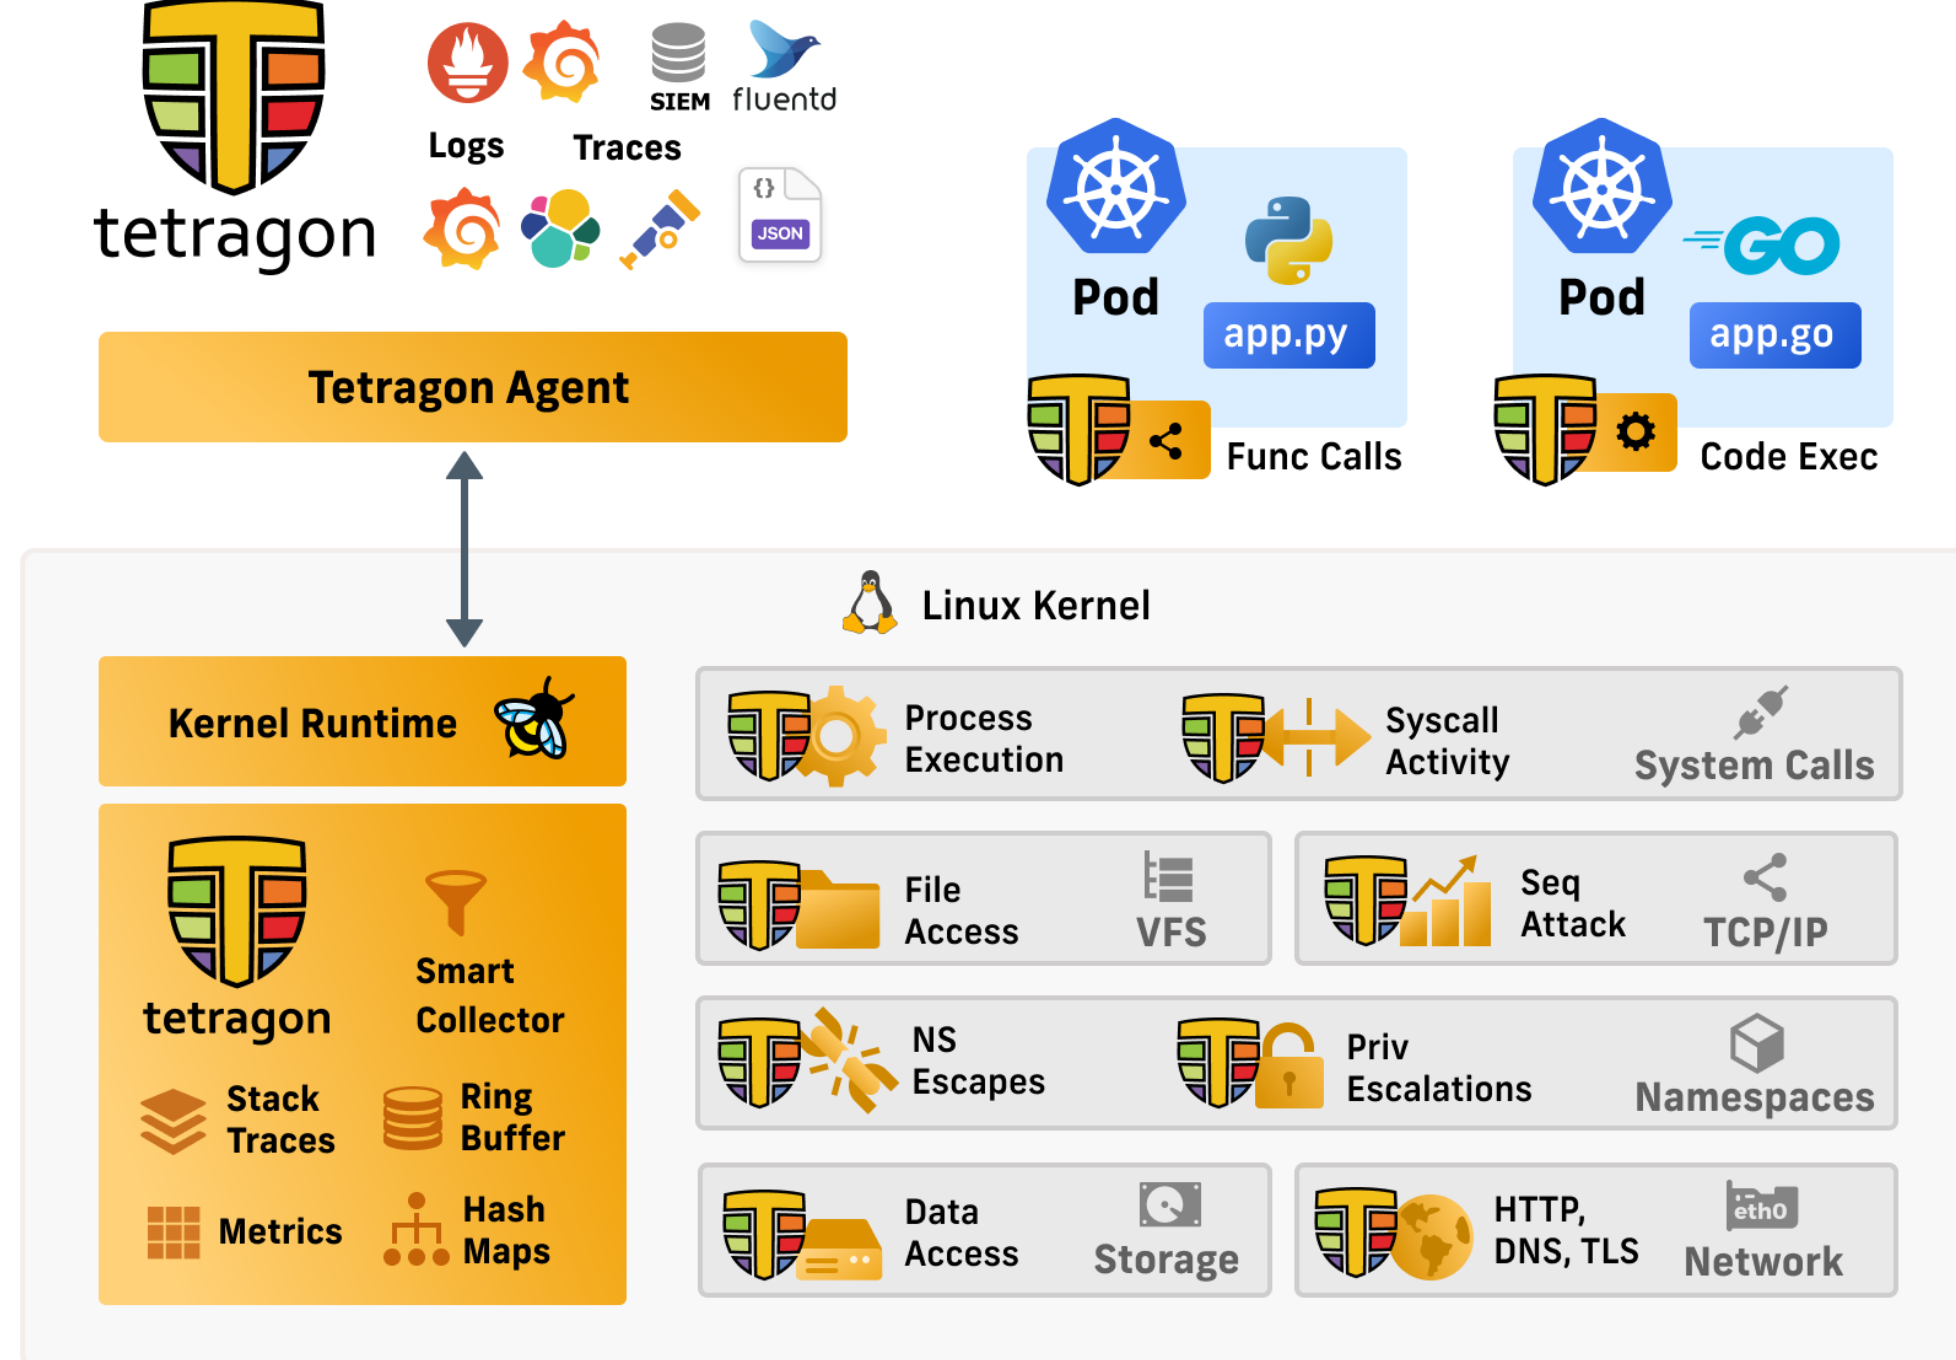


Network Policy, Network Security, Service Mesh

**Runtime Security, Runtime Enforcement,
Real-time threat detection,**

Security Observability

ebpf.io/applications



Use Cases

Process lifecycle

Tetragon observes by default the process lifecycle via exec and exit

Filename access

Monitor filename access using kprobe hooks

Network observability

Monitor TCP connect using kprobe hooks

Linux process credentials

Monitor Linux process credentials

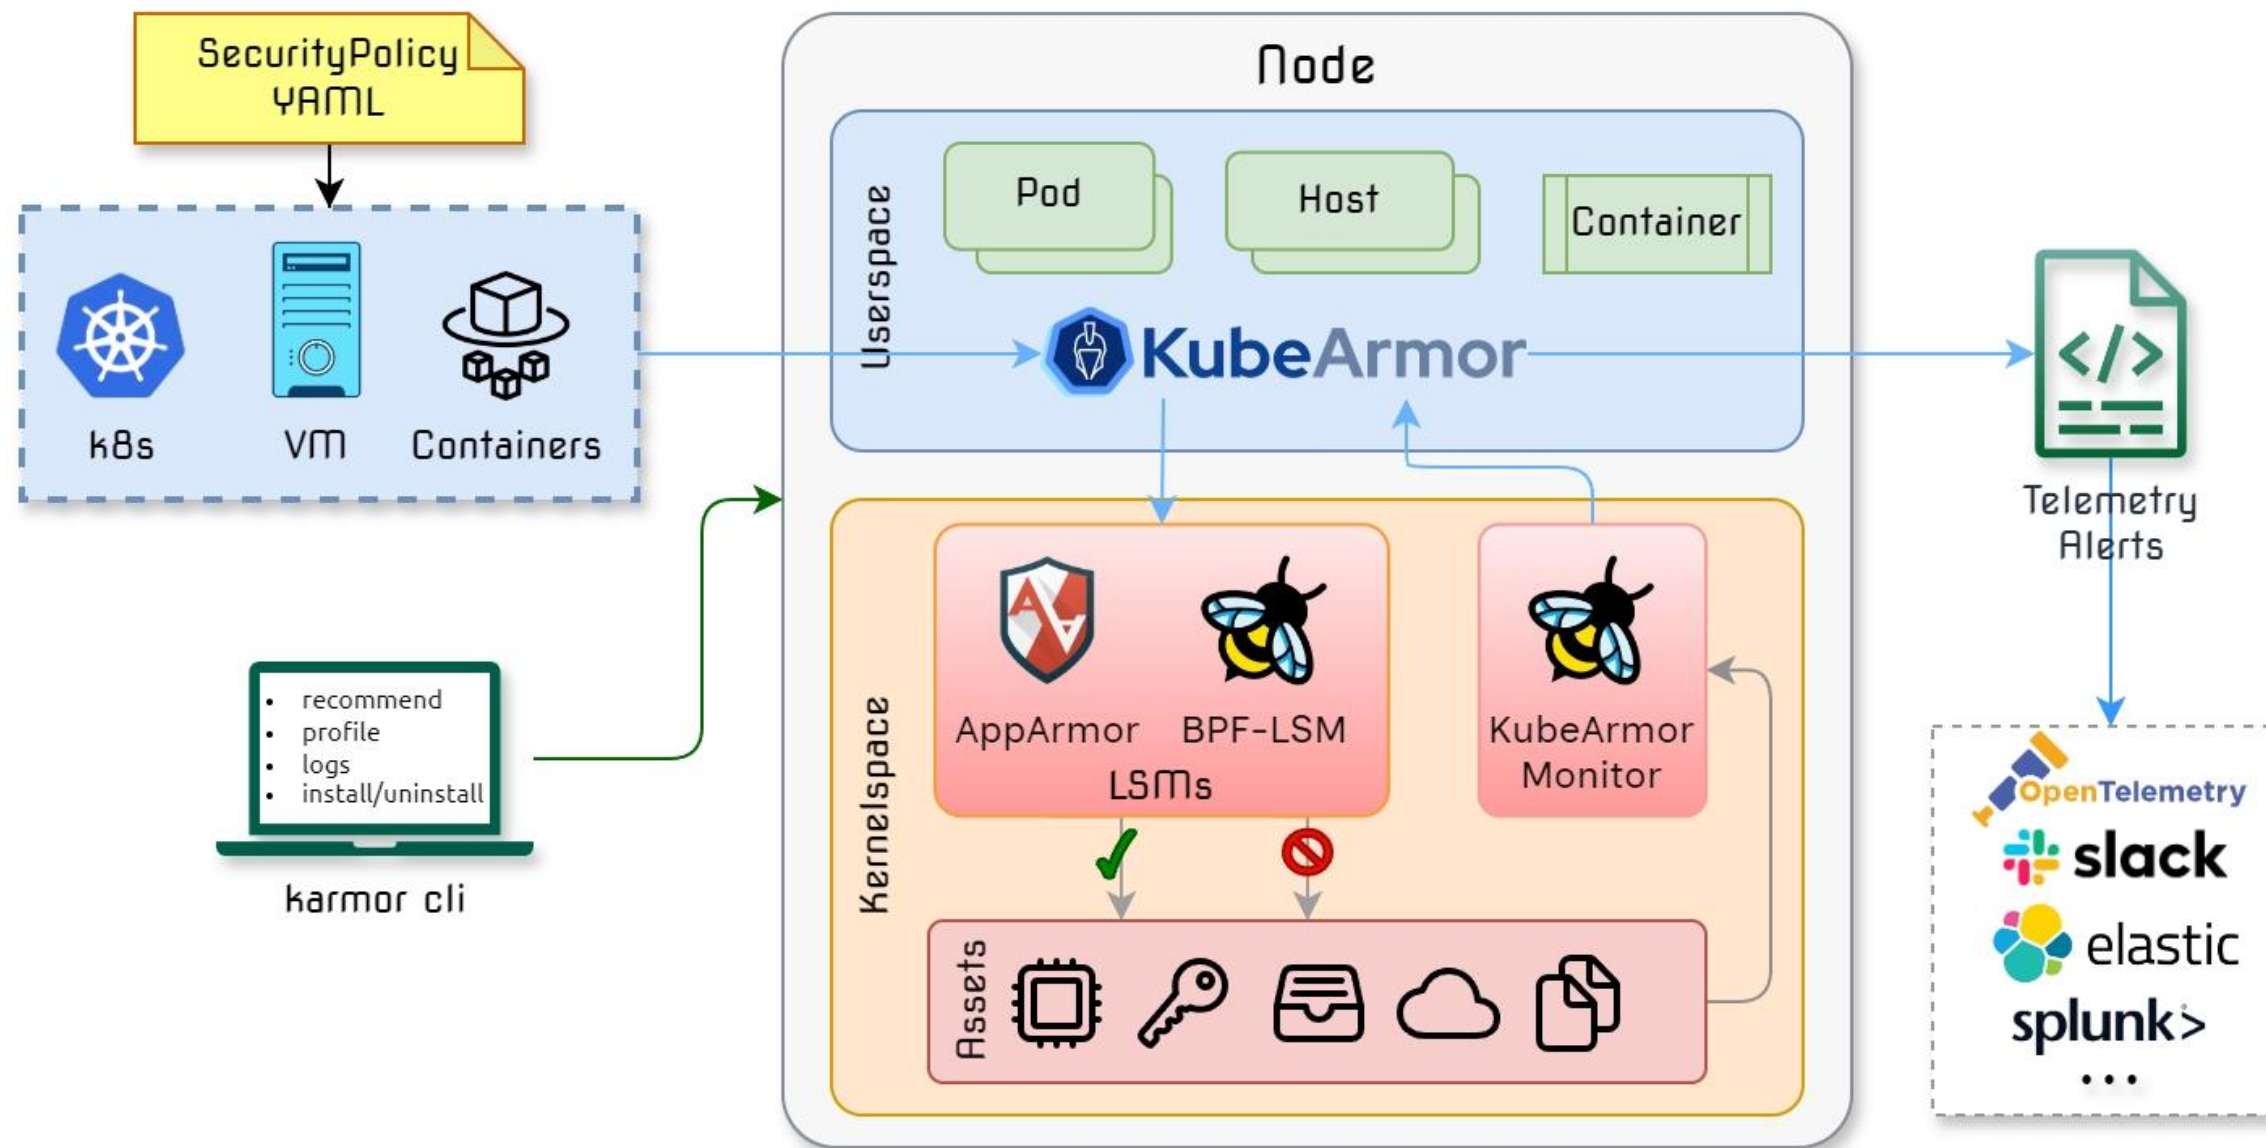
Host System Changes

Monitor Host System changes

Security Profiles

Observe and record security events




tetragon.io






kubearmor.io

Use Cases




Harden Infrastructure

-  Protect critical paths such as cert bundles
-  MITRE, STIGs, CIS based rules
-  Restrict access to raw DB table

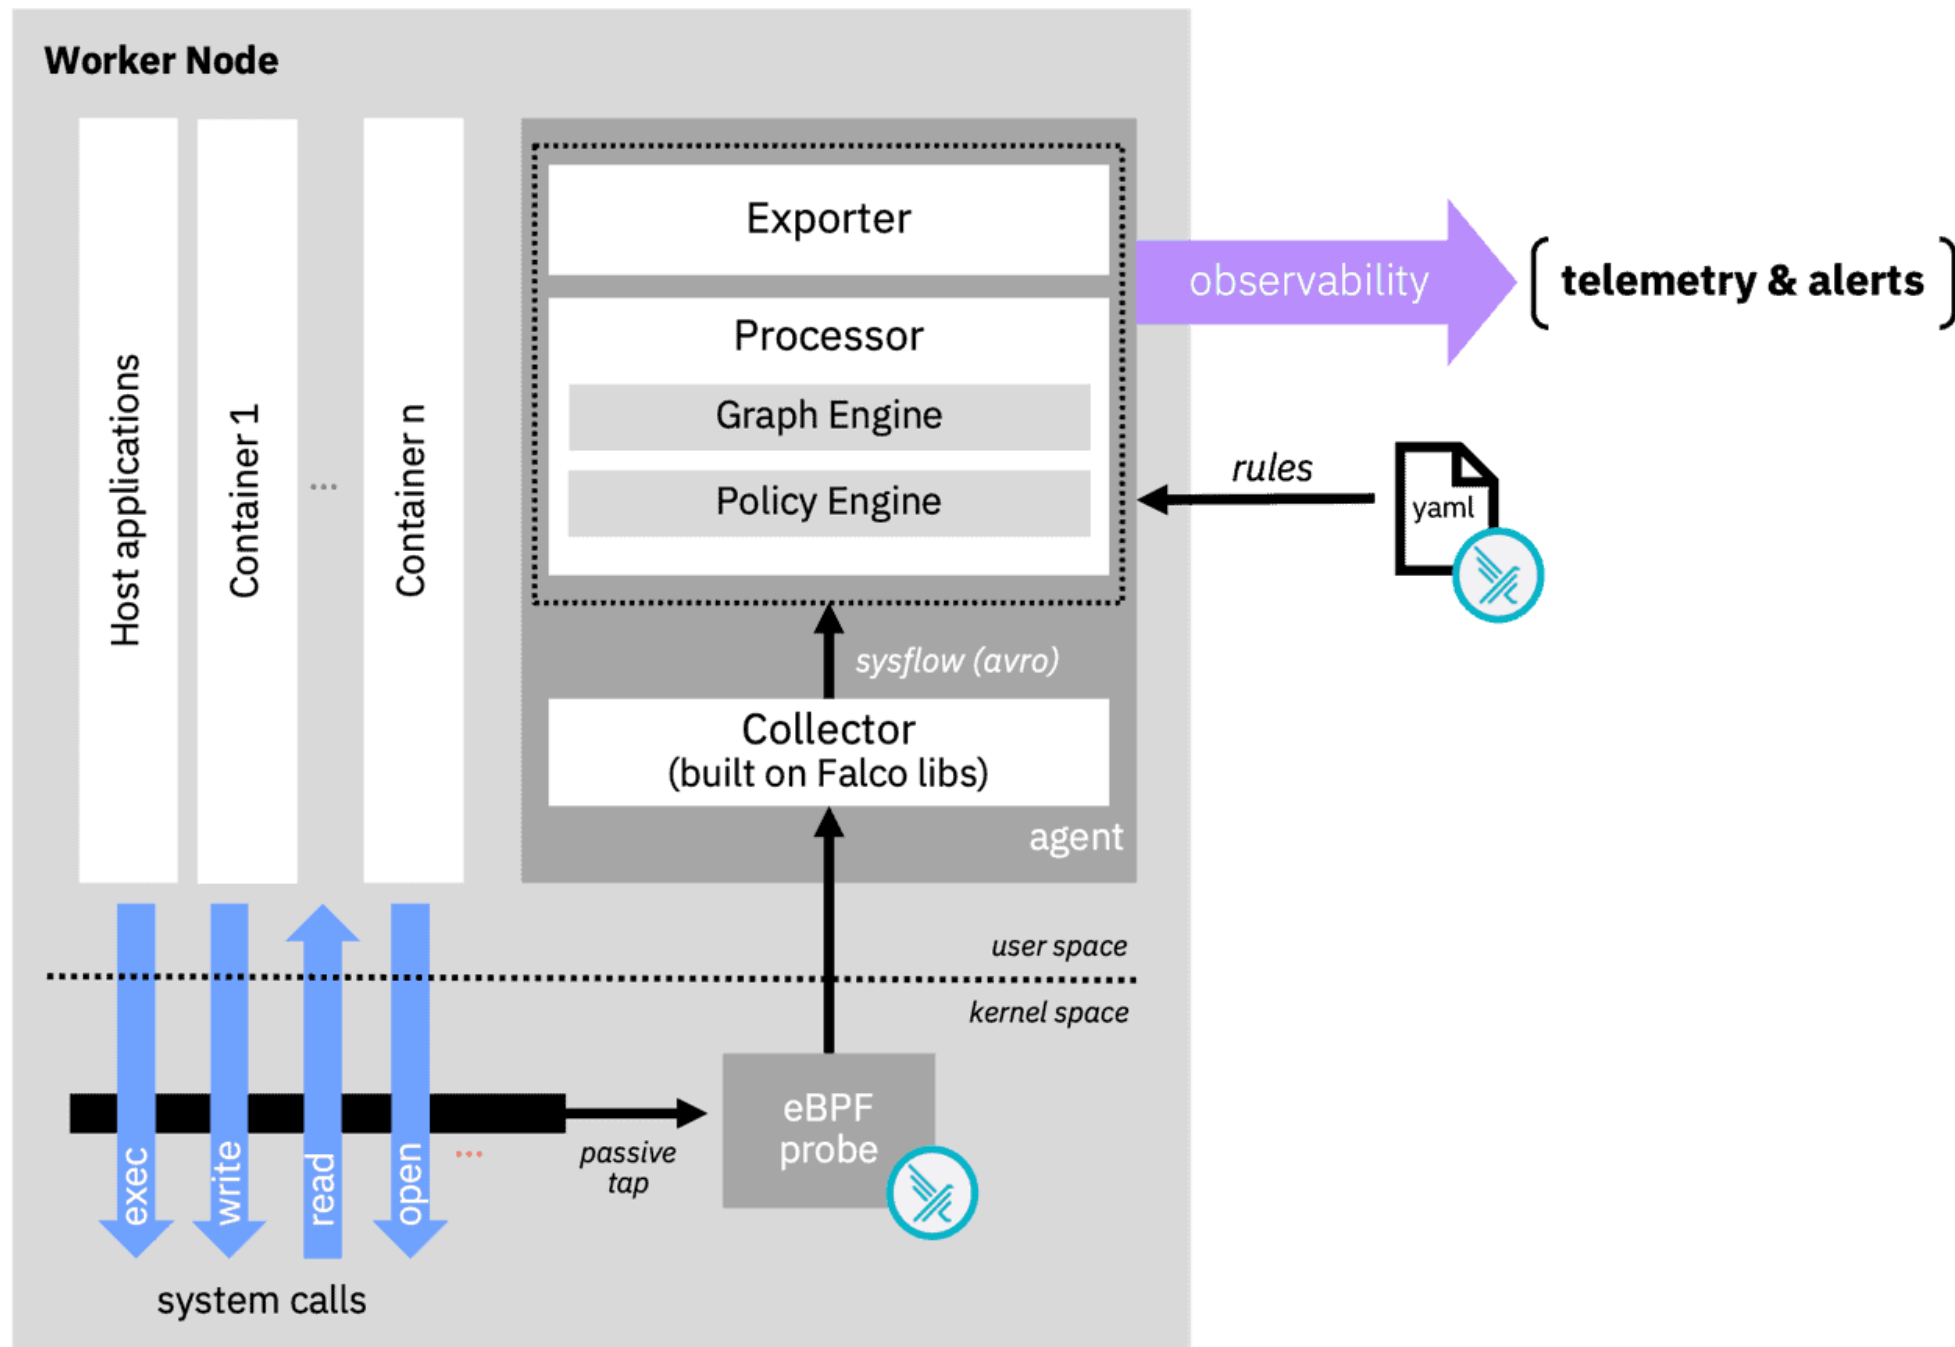
Least Permissive Access

-  Process Whitelisting
-  Network Whitelisting
-  Control access to sensitive assets

Application Behavior

-  Process execs, File System accesses
-  Service binds, Ingress, Egress connections
-  Sensitive system call profiling





Network Connection outside Local Subnet	incubating	WARNING	enabled
Mount Launched in Privileged Container	incubating	WARNING	enabled
Launch Ingress Remote File Copy Tools in Container	incubating	NOTICE	enabled
Read environment variable from /proc files	incubating	WARNING	enabled
Exfiltrating Artifacts via Kubernetes Control Plane	incubating	NOTICE	enabled
Adding ssh keys to authorized_keys	incubating	WARNING	enabled
Potential Local Privilege Escalation via Environment Variables Misuse	incubating	NOTICE	enabled
Directory traversal monitored file read	stable	WARNING	enabled
Read sensitive file trusted after startup	stable	WARNING	enabled
Read sensitive file untrusted	stable	WARNING	enabled
Run shell untrusted	stable	NOTICE	enabled
System user interactive	stable	INFO	enabled
Terminal shell in container	stable	NOTICE	enabled
Contact K8S API Server From Container	stable	NOTICE	enabled
Netcat Remote Code Execution in Container	stable	WARNING	enabled
Search Private Keys or Passwords	stable	WARNING	enabled
Clear Log Activities	stable	WARNING	enabled

Example Falco Default Rules (more than 90 rules)

Bonus

Bonus

Interesting eBPF project



[bpftrace](#)

High-level tracing language for Linux eBPF



[bpftop](#)

Real-time eBPF Program Monitoring and Performance Statistics



[Pixie](#)

Scriptable observability for Kubernetes



[Hubble](#)

Network, Service & Security Observability for Kubernetes using eBPF




[kubectl trace](#)

Schedule bpftrace programs on your Kubernetes cluster



[Caretta](#)

eBPF based Kubernetes service map



Active eBPF Landscape project by year

Source: github

Bonus

Interesting eBPF project



[bpftrace](#)

High-level tracing language for Linux eBPF



[Pixie](#)

Scriptable observability for Kubernetes



[kubectl trace](#)

Schedule bpftrace programs on your Kubernetes cluster



[bpftop](#)

Real-time eBPF Program Monitoring and Performance Statistics




[Hubble](#)

Network, Service & Security Observability for Kubernetes using eBPF



[Caretta](#)

eBPF based Kubernetes service map



Active eBPF Landscape project by year

Source: github

eBPF Case Studies



Cloudflare uses eBPF for network security, performance monitoring, and network observability



Red Hat uses eBPF at scale for load balancing and tracing in their private cloud

NETFLIX

Netflix uses eBPF at scale for network insights



Datadog uses eBPF for networking and security in their SaaS product



Meta uses eBPF to process and load balance every packet coming into their data centers



Google uses eBPF for security auditing, packet processing, and performance monitoring



Alibaba uses eBPF through Cilium to provide networking in their cloud



Line Corporation uses eBPF at scale for load balancing and tracing in their private cloud



“eBPF will be for everyone, but we don’t expect everyone to know about it”



Thank You

Mongkol Thongkraikaew



: Mongkol Thongkraikaew



: @mongkol.ttm



: Mongkol Thongkraikaew

